

THE ABI CHALLENGE

feat. inline namespaces

`arvid@libtorrent.org`

WHO AM I?

- writing C++ since 1997
- co-author of luabind (inspired by boost.python)
- author of libtorrent (bittorrent library)
- work at  **Blockstream**

WHY TALK ABOUT ABI?

- A lot of people run into ABI issues, sometimes without knowing it

WHY TALK ABOUT ABI?

- A lot of people run into ABI issues, sometimes without knowing it
- Causes memory corruption or (at best) link errors

WHY TALK ABOUT ABI?

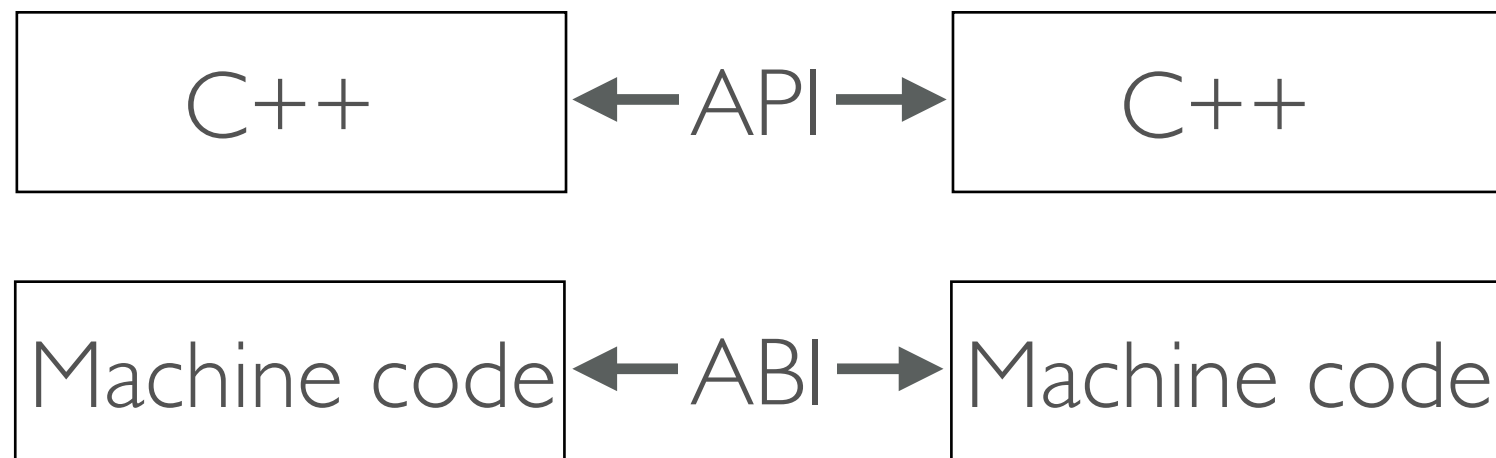
- A lot of people run into ABI issues, sometimes without knowing it
- Causes memory corruption or (at best) link errors
- Few people seem to appreciate the problem

AGENDA

- what is an ABI?
- linking symbols
- name mangling
- ABI errors in practice
- build systems
- inline namespaces
- forward declarations


ABI

- Application Binary Interface. Like API but for machine code



ABI

- compatible API → recompilation
- compatible ABI → linking with an existing binary



ABI

- **calling convention**
 - which registers to pass arguments in
 - pass two 32bit arguments in 64 bit registers
 - split class and pass fields in registers
 - pass floats in special registers
 - pass arguments in SIMD registers
 - return value optimization [C++17]
 - how are exceptions implemented?

ABI

- **class layouts**

- vtable layout (virtual inheritance)
- where/how do we pad fields, alignment
- empty base class optimization
[[no_unique_address]]
- std::pair (compressed_pair)
- std::string (small string optimization)

ABI

- across library boundaries
 - C++ version
 - may affect layout, calling convention, name mangling
 - defines
 - may affect layout (e.g. `_GLIBCXX_USE_CXX11_ABI`)
- any compiler flag that alters class layout or calling conventions

ABI

- across library boundaries
 - **C++ version**
may affect layout, calling convention, name mangling
 - **defines**
may affect layout (e.g. `_GLIBCXX_USE_CXX11_ABI`)
- any compiler flag that alters class layout or calling conventions

LINKING

- Define **$f()$** in one translation unit, call it from another

LINKING

f.cpp

```
int f(int x) {  
    return x*x;  
}
```

main.cpp


```
int f(int);  
  
int main() {  
    std::cout  
        << f(10)  
        << "\n";  
}
```

LINKING

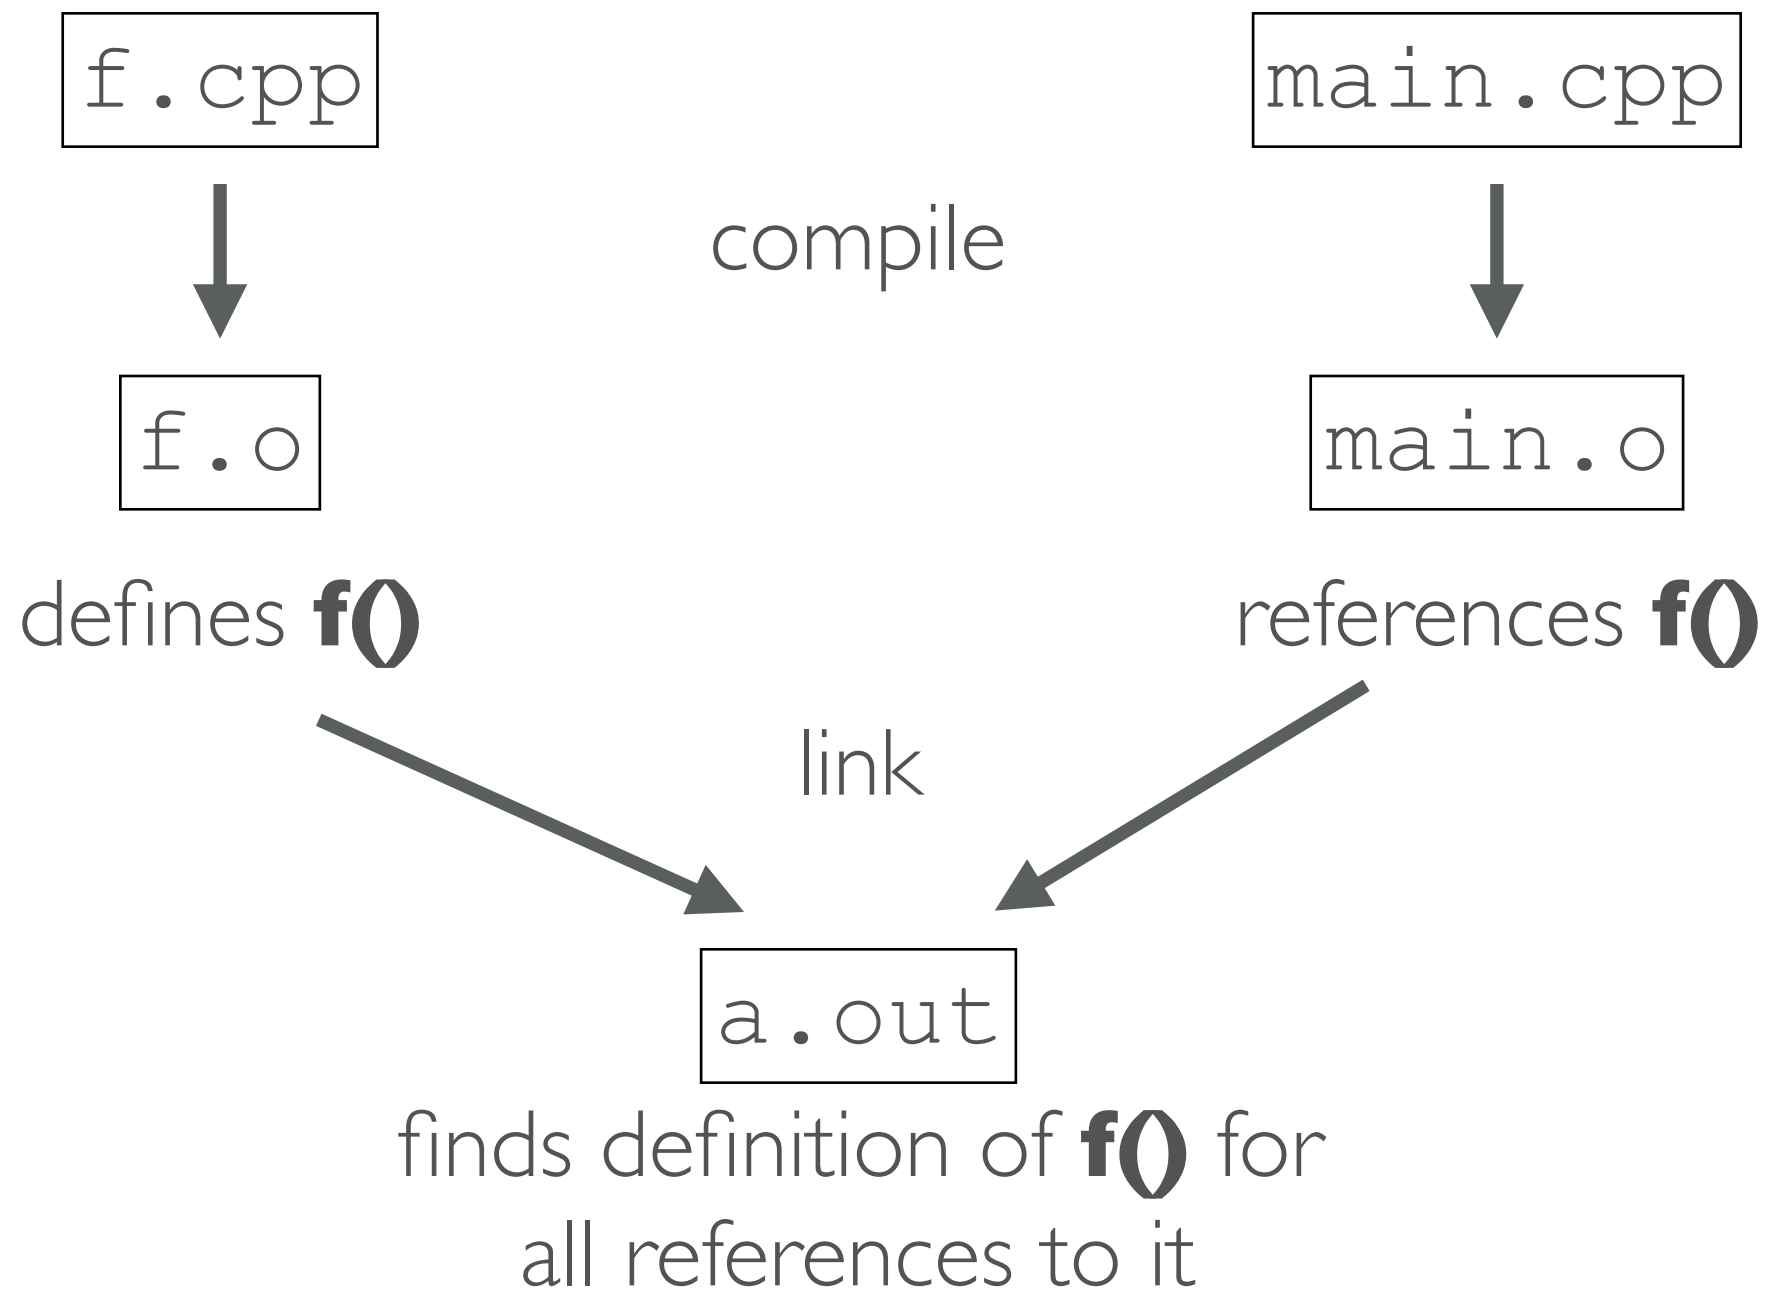
f.cpp

main.cpp

LINKING



LINKING



NAME MANGLING (C)

f.c

```
int f(int x)
{
    return x*x;
}
```

main.c

```
int f(int);
int main()
{
    printf("%d\n", f(10));
}
```

NAME MANGLING (C)

f.c

```
int f(int x)
{
    return x*x;
}
```

main.c

```
int f(int);
int main()
{
    printf("%d\n", f(10));
}
```

```
$ nm f.o
T _f
```

```
$ nm main.o
U _f
```

NAME MANGLING (C)

f.c

```
int f(int x)
{
    return x*x;
}
```

Defined symbol

main.c

```
int f(int);
int main()
{
    printf("%d\n", f(10));
}
```

External symbol

\$ nm f.o

T _f

\$ nm main.o

U _f

NAME MANGLING (C)

f.c

```
float f(float x)
{
    return x*x;
}
```

main.c

```
int f(int);
int main()
{
    printf("%d\n", f(10));
}
```

```
$ nm f.o
T _f
```

```
$ nm main.o
U _f
```

NAME MANGLING (C)

f.c

```
float f(float x)
{
    return x*x;
}
```

main.c

```
int f(int);
int main()
{
    printf("%d\n", f(10));
}
```

linked together. memory
corruption at runtime!

```
$ nm f.o
T _f
```

```
$ nm main.o
U _f
```

NAME MANGLING (C++)

f.cpp

```
float f(float x)
{
    return x*x;
}
```

main.cpp

```
int f(int);
int main()
{
    printf("%d\n", f(10));
}
```

NAME MANGLING (C++)

f.cpp

```
float f(float x)
{
    return x*x;
}
```

```
$ nm f.o
T __Z1ff
```

main.cpp

```
int f(int);
int main()
{
    printf("%d\n", f(10));
}
```

```
$ nm main.o
U __Z1fi
```


NAME MANGLING (C++)

f.cpp

```
float f(float x)
{
    return x*x;
}
```

main.cpp

```
int f(int);
int main()
{
    printf("%d\n", f(10));
}
```

symbol name mismatch
link error!

```
$ nm f.o
T __Z1ff
```

```
$ nm main.o
U __Z1fi
```

NAME MANGLING

- C++ encode signature in symbol names
 - to support overloading
 - increases type safety

NAME MANGLING

- C++ encode signature in symbol names
 - to support overloading
 - increases type safety
- **BUT**
 - encodes user types by name
 - return type (normally) not included

NAME MANGLING

```
struct foobar;  
void f(foobar const& x)  
{  
    // ...  
}
```

```
$ nm f.o  
T __Z1fRK6foobar
```

NAME MANGLING

```
struct foobar;  
void f(foobar const& x)  
{  
    // ...  
}
```

"foobar" encoded in symbol

```
$ nm f.o  
T __Z1fRK6foobar
```



NAME MANGLING

- Considering the following definition

```
struct foobar {  
    int x;  
    #ifndef NDEBUG  
    int debug_state;  
    #endif  
};
```

NAME MANGLING

- One-definition rule (ODR) violation
[basic.def.odr]

Every program shall contain exactly one definition of every non-inline function or variable that is odr-used in that program outside of a discarded statement; no diagnostic required.

NAME MANGLING

- One-definition rule (ODR) violation
[basic.def.odr]

Every program shall contain exactly one definition of every non-inline function or variable that is odr-used in that program outside of a discarded statement;
no diagnostic required.

ABI

- It's hard to maintain ABI compatibility across C++ versions

ABI

- It's hard to maintain ABI compatibility across C++ versions

"At some point, we'll need to acknowledge the problem which results from building Boost with one -std and user code with another"

— Peter Dimov ([boost mailing list](#))

ABI

"The thing that should work and the only thing that can be guaranteed to work is when your **whole code base is compiled with the same compiler with the same standard**"

— degski ([boost mailing list](#))

(emphasis mine)

ABI

- What happens in the wild when using shared libraries
 - library in release-mode, client in debug mode
 - library in C++11, client in C++14
 - library headers may be newer than library binary
the system building the client may have a newer version of the library than the system running it

EXAMPLE 1

Peter Dimov's example from boost mailing list:

```
struct X;
struct Y {
    void f ( X const& x );
    #if !defined(BOOST_NO_CXX11_RVALUE_REFERENCES)
        void f ( X&& x );
    #endif
};
```

EXAMPLE 2

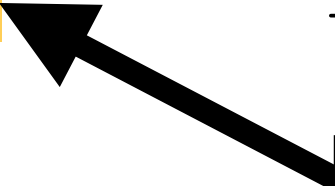
std::string may be different types depending on language version

```
std::string error_code::message() const {  
    // ...  
    return msg;  
};
```

EXAMPLE 2

std::string may be different types depending on language version

```
std::string error_code::message() const {  
    // ...  
    return msg;  
};
```



std::string may depend on language version. Not encoded in mangled name

EXAMPLE 3

```
template <class T>  
using map_string  
    = std::map<std::string, T, strview_less>;
```


EXAMPLE 3

```
#if (__cplusplus > 201103)

template <class T>
using map_string
    = std::map<std::string, T, strview_less>;

#else

template <class T>
struct map_string : std::map<std::string, T> {
    // ...
};

#endif
```

EXAMPLE 3

```
#if (__cplusplus > 201103)
```

```
template <class T>
```

```
using map_string
```

```
= std::map<std::string, T, strview_less>;
```

```
#else
```

```
template <class T>
```

```
struct map_string : std::map<std::string, T> {
```

```
// ...
```

```
};
```

```
#endif
```

these are different types,
not compatible



SOLUTION

SOLUTION

- better build systems
don't link incompatible objects

SOLUTION

- better build systems
don't link incompatible objects
- better libraries
turn memory corruption into link errors

SOLUTION

- better build systems
don't link incompatible objects
- better libraries
turn memory corruption into link errors
- increased awareness

BUILD SYSTEMS

- Build all your dependencies from source


BUILD SYSTEMS

- Build all your dependencies from source
- Use a build system that ensures link-compatibility
(or at least doesn't encourage you to create link incompatibilities)

BUILD SYSTEMS

Example 1


color indicates ABI



BUILD SYSTEMS

Example 1


color indicates ABI



build configurations propagate down


BUILD SYSTEMS

Example 2



BUILD SYSTEMS


Example 2



build restrictions propagate "up"

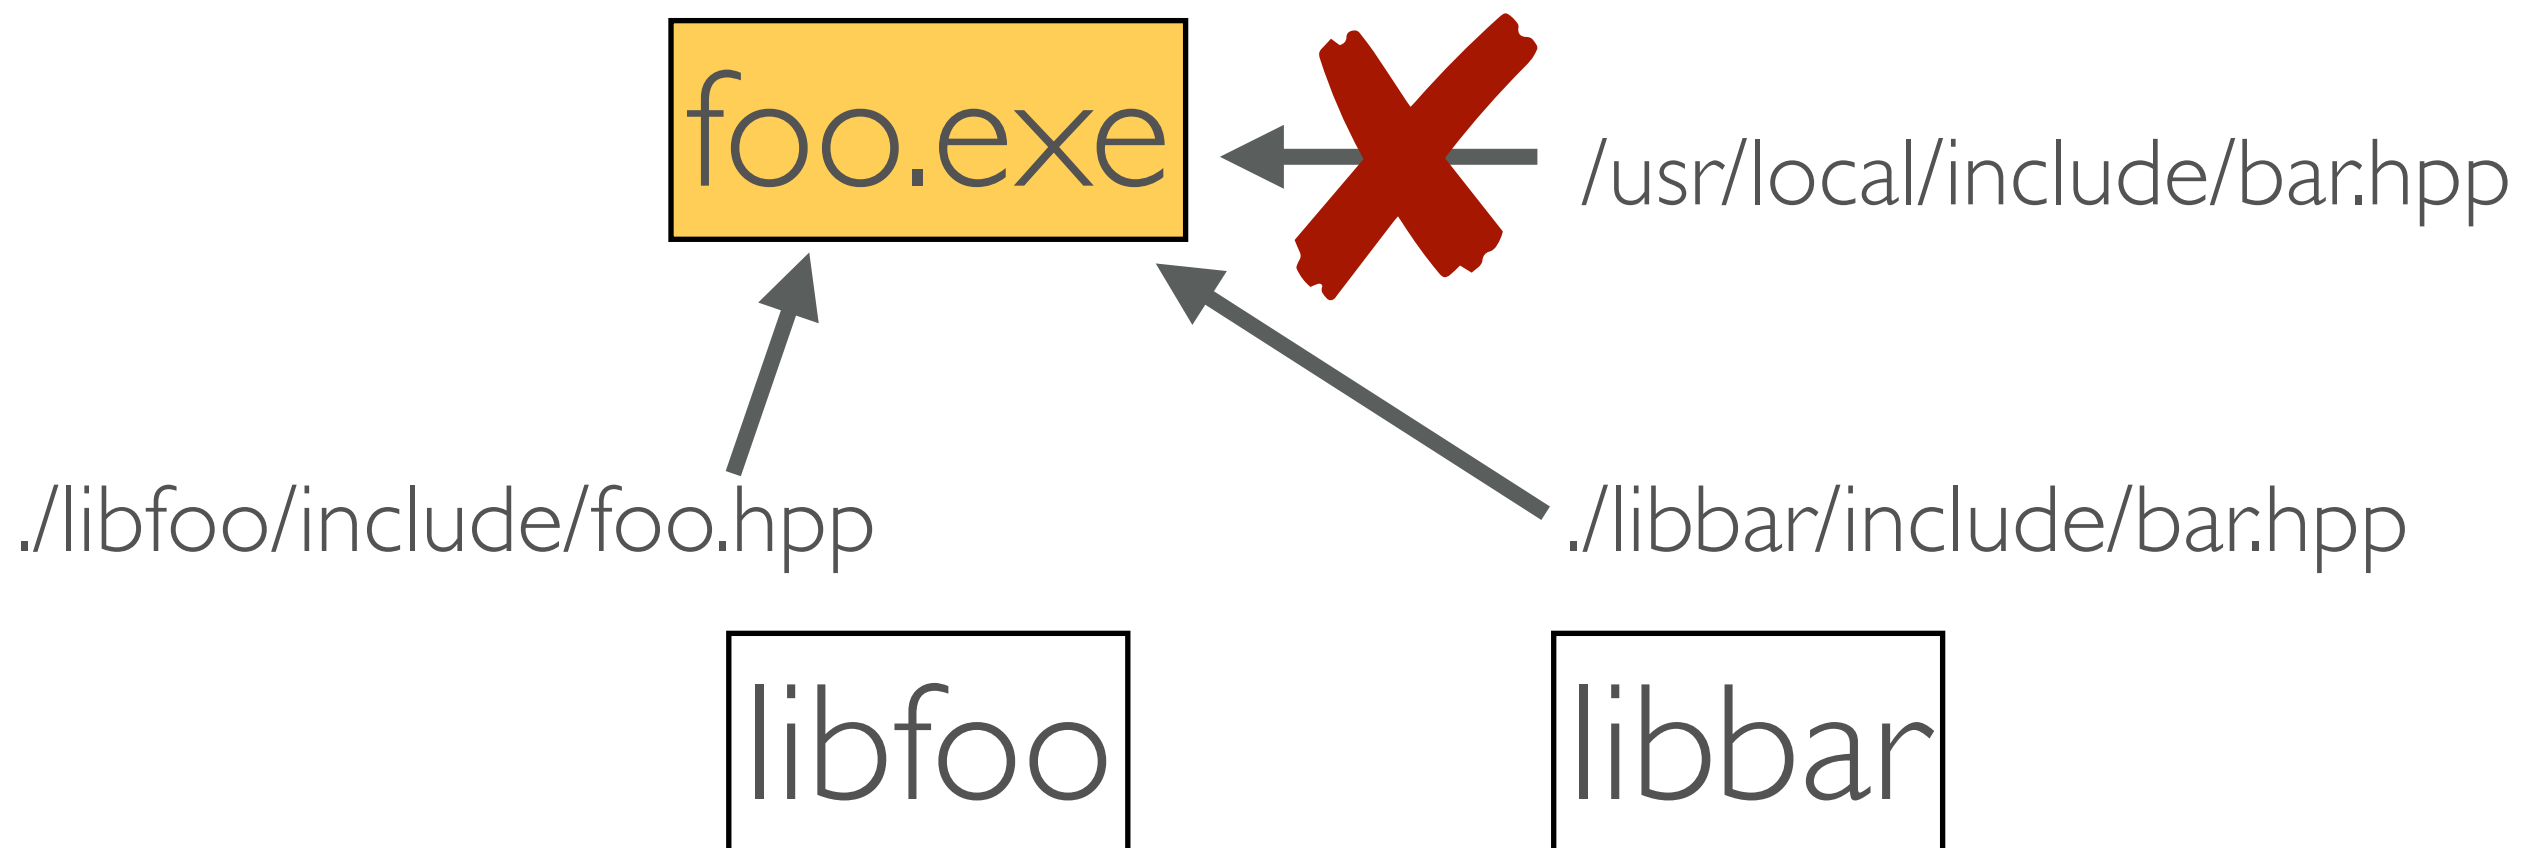
BUILD SYSTEMS

Example 3




BUILD SYSTEMS

Example 3




BUILD SYSTEMS

Example 4



BUILD SYSTEMS

Example 4



avoid `-L` linker option

BUILD SYSTEMS

- boost-build gets close to this

BUILD SYSTEMS

- boost-build gets close to this
- 3rd party dependencies is still a challenge.
example: openssl

BUILD SYSTEMS

Your library may be distributed as a binary, whether you like it or not.

what to do?

INLINE NAMESPACES

- inline namespaces let you inject information into the mangled name (**ABI**), without altering the **API**

INLINE NAMESPACES

- inline namespaces let you inject information into the mangled name (**ABI**), without altering the **API**
- inline namespaces affect the *linker names* of symbols, while preserving the API

INLINE NAMESPACES

- inline namespaces let you inject information into the mangled name (**ABI**), without altering the **API**
- inline namespaces affect the *linker names* of symbols, while preserving the API
- e.g. the two **foobar** types could be given different names

NAME MANGLING

- foobar is defined as:

```
struct foobar {  
    int x;  
    #ifndef NDEBUG  
        int debug_state;  
    #endif  
};
```

NAME MANGLING

- foobar is defined as:

```
#ifdef NDEBUG  
inline namespace rel {  
#else  
inline namespace dbg {  
#endif  
  
struct foobar {  
    int x;  
#ifndef NDEBUG  
    int debug_state;  
#endif  
};  
} // inline namespace
```


INLINE NAMESPACES

When built in debug mode

```
$ nm f.o  
T __Z1fRKN3dbg6foobarE
```

When built in release mode

```
$ nm f.o  
T __Z1fRKN3rel6foobarE
```

NAME MANGLING

```
Undefined symbols for architecture x86_64:  
  "f(dgb::foobar const&)", referenced from:  
      _main in main.o  
ld: symbol(s) not found for architecture  
x86_64  
clang-7: error: linker command failed with  
exit code 1 (use -v to see invocation)
```

INLINE NAMESPACES

- inline namespaces make all its declarations available in its surrounding scope

INLINE NAMESPACES

```
namespace my_library {
```

```
inline namespace v1 {
```

```
int f(int) { ... };
```

```
} // v1
```

```
} // my_library
```

library code

```
...
```

```
my_library::f(10);
```

```
...
```

```
using my_library::f;
```

```
f(10)
```

client code

INLINE NAMESPACES

```
namespace my_library {  
  
    inline namespace v1 {  
  
        int f(int) { ... };  
  
    } // v1  
} // my_library
```

namespace **v1** not
necessary in user code

```
...  
my_library::f(10);  
...  
using my_library::f;  
f(10)
```

client code

INLINE NAMESPACES

- you can specialise templates

INLINE NAMESPACES

```
namespace my_library {
```

```
inline namespace v1 {
```

```
    template <typename T>
```

```
    struct foobar {};
```

```
} // v1
```

```
} // my_library
```

library code

```
namespace my_library {
```

```
    template <>
```

```
    struct foobar<int> { ... };
```

```
} // my_library
```

client code

INLINE NAMESPACES

```
namespace my_library {  
  
    inline namespace v1 {  
  
        template <typename T>  
        struct foobar {};  
  
    } // v1  
} // my_library
```

defined in **my_library::v1**
specialized in **my_library**



```
namespace my_library {  
  
    template <>  
    struct foobar<int> { ... };  
  
} // my_library
```

client code

INLINE NAMESPACES

- From the user's point of view, as if the inline namespace isn't there
- Very useful for versioning of functions and types in libraries

INLINE NAMESPACE

```
namespace std {
```

```
namespace cxx11 {
    template <...> class basic_string {...}; // SSO
}
```

```
namespace cxx98 {
    template <...> class basic_string {...}; // COW
}

// std
```

INLINE NAMESPACES

```
namespace std {  
  
    #if defined _GLIBCXX_USE_CXX11_ABI  
    inline namespace cxx11 {}  
    #else  
    inline namespace cxx98 {}  
    #endif  
  
    namespace cxx11 {  
        template <...> class basic_string {...}; // SSO  
    }  
  
    namespace cxx98 {  
        template <...> class basic_string {...}; // COW  
    }  
  
} // std
```

INLINE NAMESPACES

```
namespace std {
```

```
#if defined _GLIBCXX_USE_CXX11_ABI
```

```
inline namespace cxx11 {}
```

```
#else
```

```
inline namespace cxx98 {}
```

```
#endif
```

```
namespace cxx11 {
```

```
    template <...> class basic_string {...}; // SSO
```

```
}
```

```
namespace cxx98 {
```

```
    template <...> class basic_string {...}; // COW
```

```
}
```

```
} // std
```

inlineness depends on *first*
time namespace is opened



INLINE NAMESPACES

Inline namespaces to provide
backwards ABI compatibility

INLINE NAMESPACES

```
namespace library {  
  
    inline namespace v1 {}  
  
    namespace v1 {  
        void f(std::string const&);  
    }  
  
} // library
```

header

```
namespace library {  
  
    namespace v1 {  
        void f(std::string const& s) {  
            // do some work  
        }  
    }  
  
} // library
```

C++ file

INLINE NAMESPACES

```
namespace library {  
  
#ifdef USE_OLD_API  
inline namespace v1 {}  
#else  
inline namespace v2 {}  
#endif  
  
namespace v1 {  
    void f(std::string const&);  
}  
  
namespace v2 {  
    void f(std::string_view);  
}  
  
} // library
```

header

```
namespace library {  
  
namespace v1 {  
    void f(std::string const& s) {  
        v2::f(s);  
    }  
}  
  
namespace v2 {  
    void f(std::string_view) {  
        // do some work  
    }  
}  
  
} // library
```

C++ file

FORWARD DECLARATIONS

- A symbol's actual namespace is an *implementation detail*
- clients may not forward declare 3rd party symbols

FORWARD DECLARATIONS

```
namespace third_party {  
  
    struct foo;  
  
} // third_party  
  
struct bar  
{  
    ...  
    third_party::foo* the_foo_;  
};
```

client code

FORWARD DECLARATIONS

```
namespace third_party {
```

```
    struct foo;
```

```
} // third_party
```

```
struct bar
```

```
{
```

```
...
```

```
third_party::foo* the_foo_;
```

```
};
```

foo does not need to be a complete type



client code

FORWARD DECLARATIONS

```
namespace third_party {
```

```
    struct foo;
```

```
} // third_party
```

```
struct bar
```

```
{
```

```
    ...
```

```
    third_party::foo* the_foo_;
```

```
};
```



may be tempting to forward declare it, despite coming from a 3rd party

client code

FORWARD DECLARATIONS

```
namespace third_party {  
    struct foo;  
} // third_party  
  
struct bar  
{  
    ...  
    third_party::foo* the_fo  
};
```

```
namespace third_party {  
inline namespace v2 {  
    struct foo;  
} // v2  
} // third_party
```

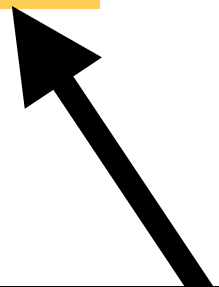
**library
code**

client code

FORWARD DECLARATIONS

```
namespace third_party {  
    struct foo;  
} // third_party  
  
struct bar  
{  
    ...  
    third_party::foo* the_fo  
};
```

```
namespace third_party {  
    inline namespace v2 {  
        struct foo;  
    } // v2  
} // third_party
```



namespace v2 is an
implementation detail

code

client code

FORWARD DECLARATIONS

```
#include "third_party/fwd.hpp"
```

```
struct bar  
{  
    ...  
    third_party::foo* the_foo_;  
};
```

FORWARD DECLARATIONS

```
#include "third_party/fwd.hpp"
```

```
struct bar  
{  
    ...  
    third_party::foo* the_foo_;  
};
```



forward declaration
controlled by library

FORWARD DECLARATIONS

non-trivial libraries should provide
forward declarations headers

FORWARD DECLARATIONS

- `<iosfwd>`
- `<netfwd>`
- `<boost/numeric/ublas/fwd.hpp>`
- `<boost/math/distributions/fwd.hpp>`

FORWARD DECLARATIONS

surprisingly few libraries
provide fwd headers

SUMMARY

- Build your dependencies from source
- Use a good build system
- Use **inline namespace** for
 - backward compatible upgrades to your ABI
 - ABI-safe build configurations
- Library authors, provide forward declaration headers!

THANK YOU



`github.com/arvidn`



`arvid@libtorrent.org`