

# Trivially Relocatable

Arthur O'Dwyer  
2019-05-09

## PART 1:

- What is relocation (move+destroy)?
  - Applications for reliable detection of trivial relocatability [5–44]
- Prior art (Folly, EASTL, BSL) [45–55]
  - And why it's fragile and error-prone [56–68]
- P1144 definition [69–90]

## PART 2:

- Versus past and current proposals [91–98, 99–103, 104–109, 110–114]
- Versus “persistent memory” [115–127]
- Open questions [128–129]

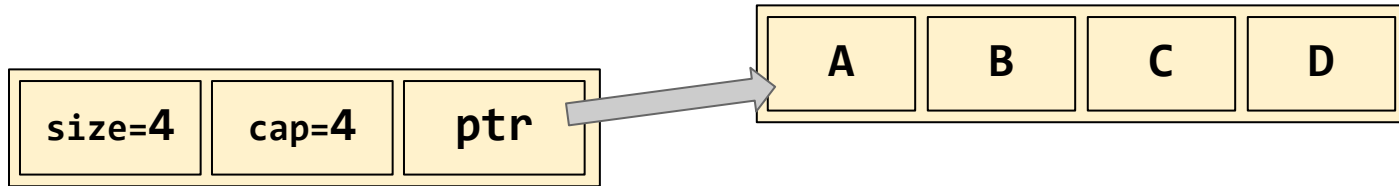
## BONUS SLIDES [130–156]

Hey look!  
Slide numbers!

# Motivating “relocation”

Consider what happens when we resize a `std::vector<T>`.

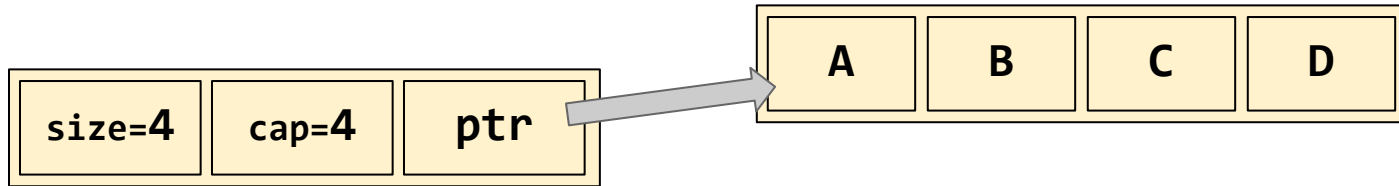
```
std::vector<T> vec { A, B, C, D };
```



# Motivating “relocation”

Consider what happens when we resize a `std::vector<T>`.

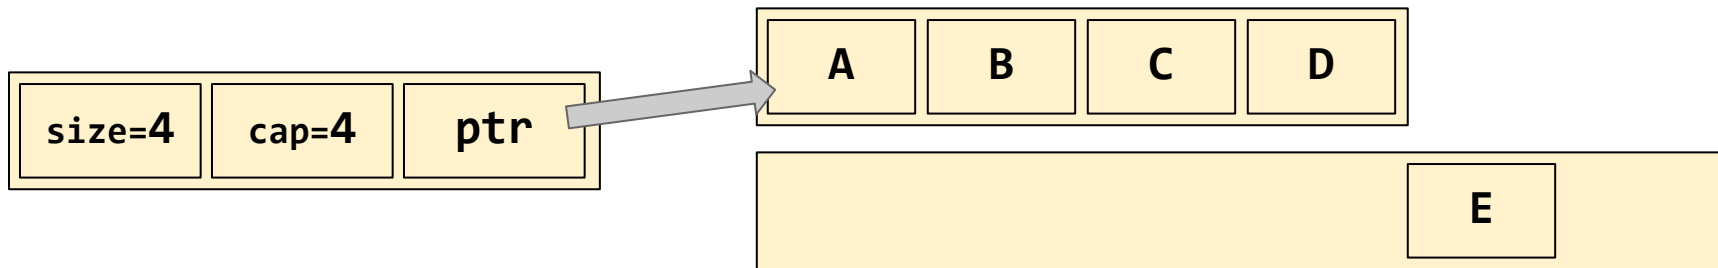
```
std::vector<T> vec { A, B, C, D };  
vec.push_back(E);
```



# Motivating “relocation”

Consider what happens when we resize a `std::vector<T>`.

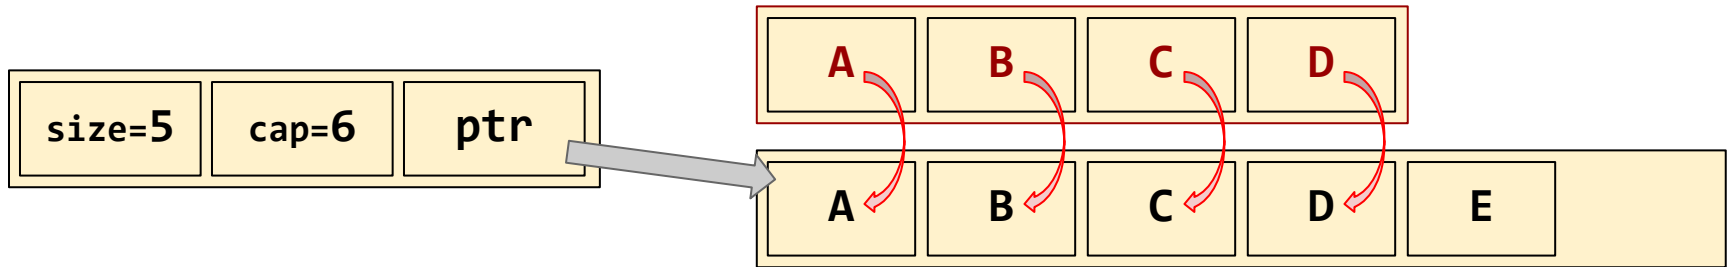
```
std::vector<T> vec { A, B, C, D };  
vec.push_back(E);
```



# Motivating “relocation”

Consider what happens when we resize a `std::vector<T>`.

```
std::vector<T> vec { A, B, C, D };  
vec.push_back(E);
```




**The “relocation” of objects A, B, C, D involves 4 calls to the move-constructor, followed by 4 calls to the destructor.**

# Relocating trivially copyable `int*`

C++ source #1

Save/Load Add new... C++

```
1 #include <vector>
2 using std::vector;
3
4 void foo(vector<int*>& dest)
5 {
6     dest.reserve(100);
7 }
```



x86-64 gcc (trunk) (Editor #1, Compiler #1) C++

x86-64 gcc (trunk) -O3 -fomit-frame-pointe

A- 11010 .LX0: .text // \s+ Intel Demangle

Libraries Add new...

```
26 cmpq %r14, %rbp
27 je .L3
28 movq %r13, %rdx
29 movq %rbp, %rsi
30 movq %rax, %rdi
31 call memcpy
32 .L3:
33 movq (%rbx), %rdi
34 testq %rdi, %rdi
35 je .L4
36 call operator delete(void*)
37 .L4:
38 movq %r12, (%rbx)
39 addq %r12, %r13
40 addq $800, %r12
41 movq %r13, 8(%rbx)
42 movq %r12, 16(%rbx)
43 popq %rbx
44 popq %rbp
45 popq %r12
46 popq %r13
47 popq %r14
48 ret
```

48 lines of assembly

x86-64 clang (trunk) (Editor #1, Compiler #2) C++

x86-64 clang (trunk) -O3 -fomit-frame-pointe

A- 11010 .LX0: .text // \s+ Intel Demangle

Libraries Add new...

```
22 movq %r12, %rsi
23 movq %r11, %rdx
24 callq memcpy
25 .LBB0_5:
26 sarg $3, %r14
27 movq (%rbx), %rdi
28 testq %rdi, %rdi
29 je .LBB0_5
30 callq operator delete
31 .LBB0_5:
32 movq %r15, (%rbx)
33 leaq (%r15,%r14,8), %r12
34 movq %rax, 8(%rbx)
35 addq $800, %r15
36 movq %r15, 16(%rbx)
37 .LBB0_6:
38 addq $8, %rsp
39 popq %rbx
40 popq %r12
41 popq %r14
42 popq %r15
43 retq
```

43 lines of assembly


Output (0/0) g++ (GCC-Explorer-Build) 9.0.0 20180924

Output (0/0) clang version 8.0.0 (trunk 342934) - cached

# Relocating non-trivial shared\_ptr

A- Save/Load + Add new... C++

```
1 #include <memory>
2 #include <vector>
3 using std::shared_ptr;
4 using std::vector;
5
6 using P = shared_ptr<int>;
7
8 void foo(vector<P>& dest)
9 {
10     dest.reserve(100);
11 }
```



x86-64 gcc (trunk) -O3 -fomit-frame-pointe

A- 11010 .LX0: .text // \s+ Intel Demangle

Libraries+ Add new...

```
83 jne .L34
84 .L16:
85 testq %r15, %r15
86 je .L17
87 movl $-1, %eax
88 lock xaddl %eax, 12(%rbx)
89 .L18:
90 cmpl $1, %eax
91 jne .L12
92 movq (%rbx), %rax
93 movq %rcx, 8(%rsp)
94 movq %rbx, %rdi
95 movq 24(%rax), %rdx
96 cmpq $std::_Sp_counted_base<
97 jne .L19
98 call *(%rax)
99 movq 8(%rsp), %rcx
100 addq $16, %rbp
101 cmpq %rbp, %rcx
102 jne .L8
103 .L32:
104 movq 0(%r13), %rbp
```

138 lines  
of assembly

x86-64 clang (trunk) -O3 -fomit-frame-pointe

A- 11010 .LX0: .text // \s+ Intel Demangle

Libraries+ Add new...

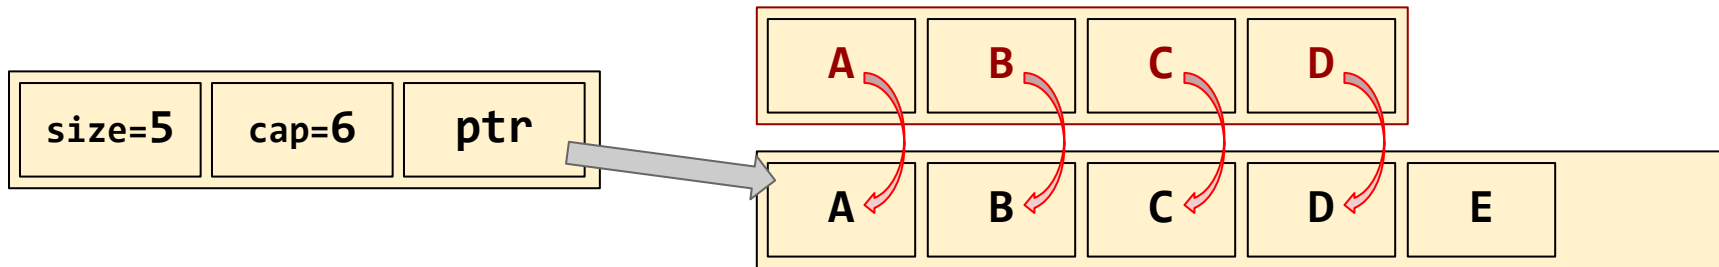
```
84 je .LBB1_24
85 movq %r13, 8(%rsp)
86 movq %r15, 16(%rsp)
87 movl $_pthread_key_
88 .LBB1_12:
89 movq 8(%rbx), %r13
90 testq %r13, %r13
91 je .LBB1_22
92 testq %r15, %r15
93 je .LBB1_15
94 movl $-1, %eax
95 lock xaddl
96 cmpl $1, %eax
97 je .LBB1_17
98 jmp .LBB1_22
99 .LBB1_15:
100 movl 8(%r13), %eax
101 leal -1(%rax), %ecx
102 movl %ecx, 8(%r13)
103 cmpl $1, %eax
104 jne .LBB1_22
105 .LBB1_17:
```

159 lines  
of assembly



# Relocating non-trivial types

In principle, we **can** implement the “relocation” of objects A, B, C, D here with a simple memcpy. `shared_ptr`’s move constructor is non-trivial, and its destructor is also non-trivial, but if we always call them together, the **result** is tantamount to memcpy.



The operation of “calling the move-constructor and the destructor together in pairs” is known as **relocation**.

A type whose relocation operation is tantamount to memcpy is **trivially relocatable**.

# Not all types are trivially relocatable

A type whose relocation operation is ***not*** tantamount to memcpy is called ***non-trivially relocatable***.

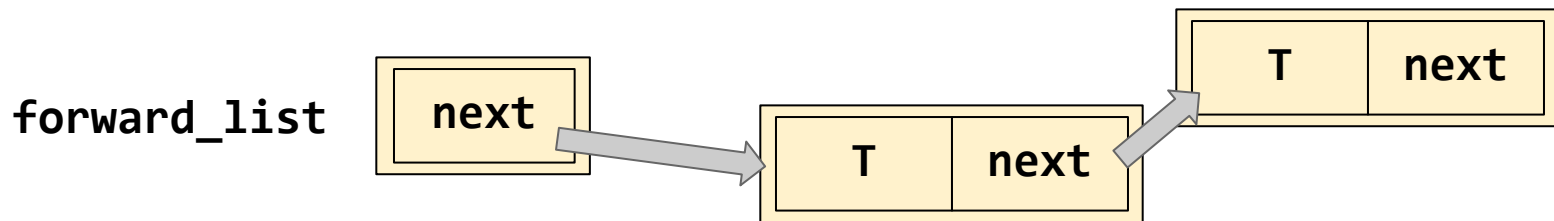
Although many everyday types are trivially relocatable, there do exist non-trivially relocatable types.

Some of them are pretty common. For example, libc++'s `std::list`.

Let's compare libc++'s `std::forward_list<int>` and `std::list<int>`.

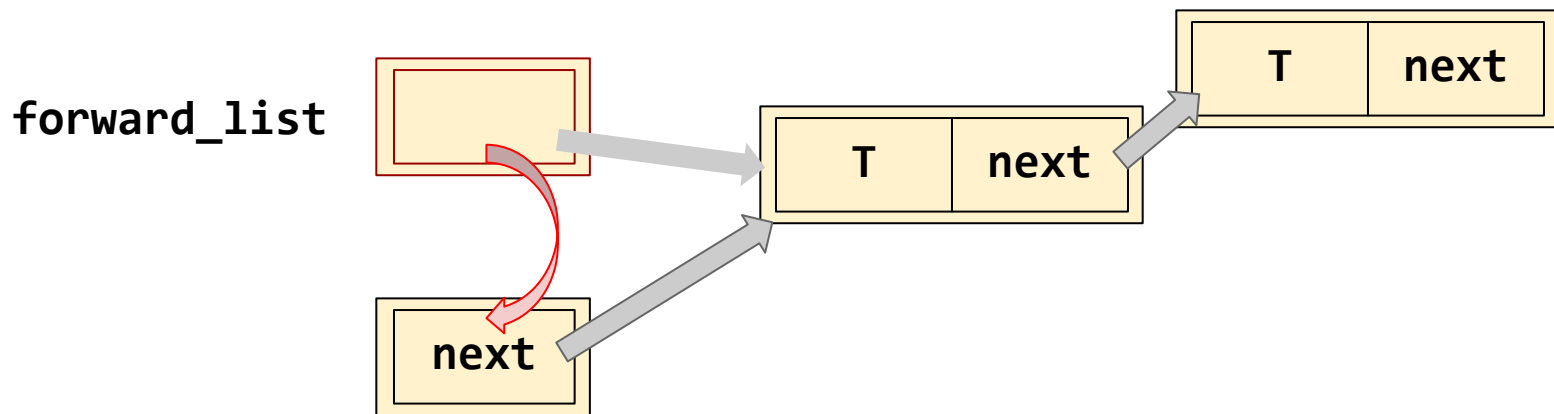
# Not all types are trivially relocatable

libc++'s `std::forward_list<T>` (with `std::allocator`) is trivially relocatable.



# Not all types are trivially relocatable

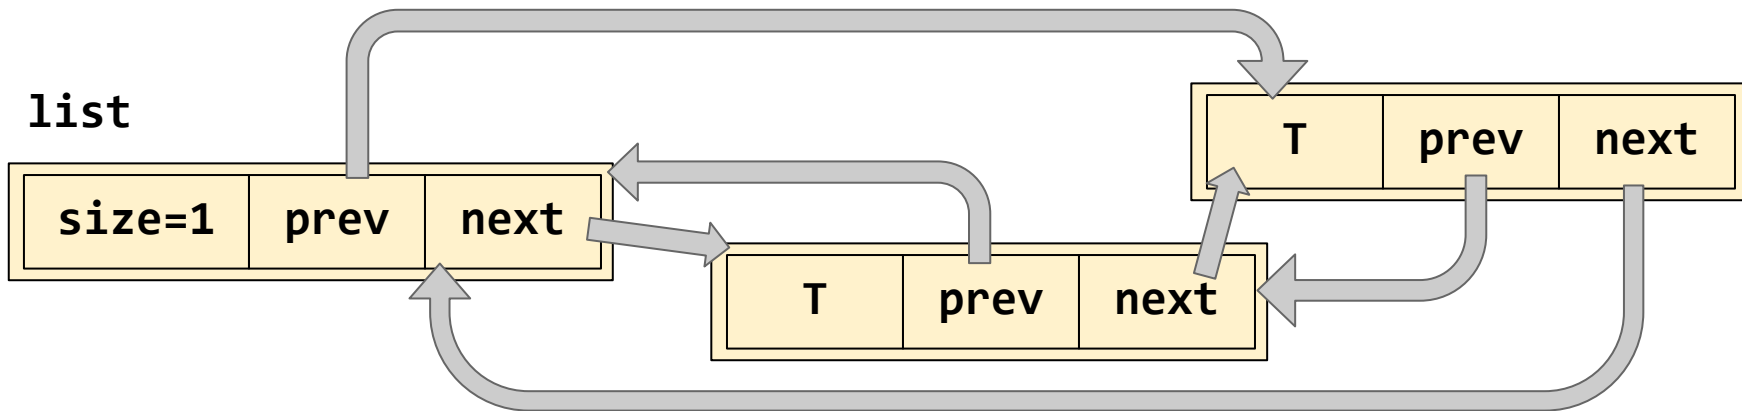
libc++'s `std::forward_list<T>` (with `std::allocator`) is trivially relocatable.



➡: The memory may still hold that bit-pattern, but the C++ object's lifetime is already over.

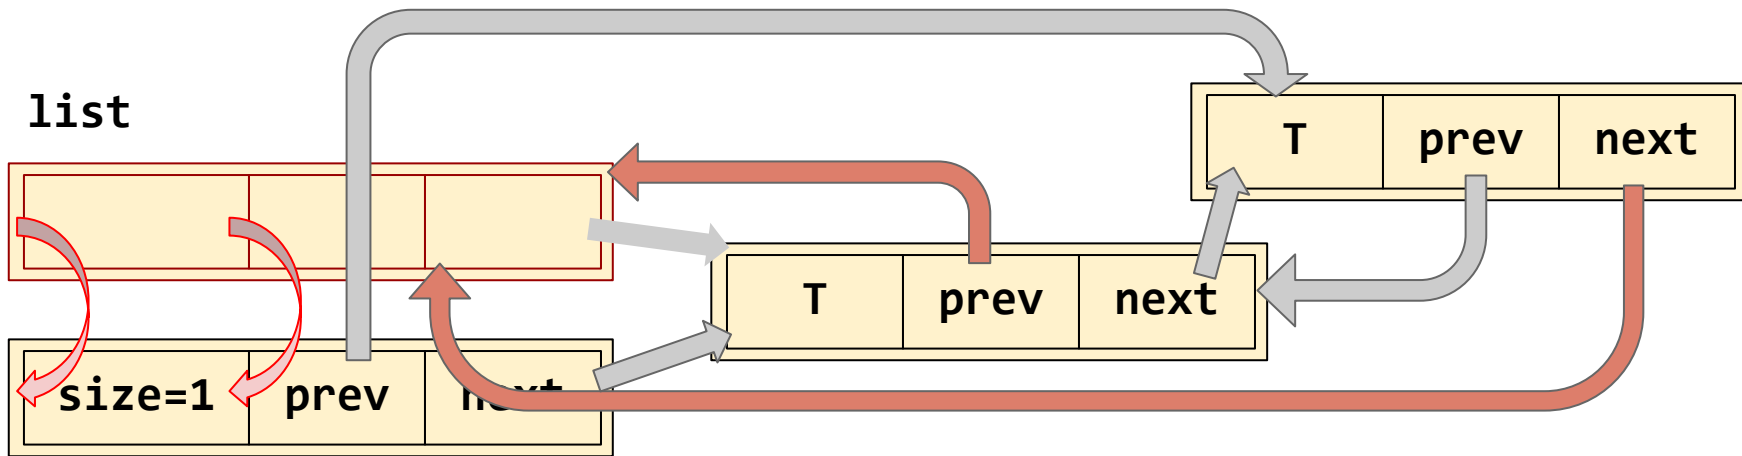
# Not all types are trivially relocatable

libc++'s `std::list<T>` is non-trivially relocatable.



# Not all types are trivially relocatable

libc++'s `std::list<T>` is non-trivially relocatable.



Relocating the `list` object requires some “fixup” beyond just a simple memcopy.

# Not all types are trivially relocatable

If we can reliably distinguish trivially relocatable types from non-trivially relocatable types, then we can use trivial `memcpy` for the former and fall back to non-trivial `move+destroy` for the latter.

Reallocating a vector of `std::forward_list<T>`? Use `memcpy`!

Reallocating a vector of `std::list<T>`? Use `move+destroy` in a loop.

Let's teach `libc++` that `shared_ptr` is trivially relocatable, and see what improvement we get on our “vector of `shared_ptr`” example.

# Trivially relocatable shared\_ptr

Compiler: x86-64 clang (experimental P1144) -O3 -fomit-frame-pointers

```
1 #include <memory>
2 #include <vector>
3 using std::shared_ptr;
4 using std::vector;
5
6 using P = shared_ptr<int>;
7
8 void foo(vector<P>& dest)
9 {
10     dest.reserve(100);
11 }
```

Assembly (x86-64 clang (trunk) -O3 -fomit-frame-pointers):

```
84     je     .LBB1_24
85     movq   %r13, 8(%rsp)
86     movq   %r15, 16(%rsp)
87     movl   $__pthread_key_0, %eax
88 .LBB1_12:
89     movq   8(%rbx), %r13
90     testq  %r13, %r13
91     je     .LBB1_22
92     testq  %r15, %r15
93     je     .LBB1_15
94     movl   $-1, %eax
95     lock   xaddl %eax, 8(%rbx)
96     cmpl   $1, %eax
97     je     .LBB1_17
98     jmp    .LBB1_22
99 .LBB1_15:
100     movl   8(%r13), %eax
101     leal   -1(%rax), %ecx
102     movl   %ecx, 8(%r13)
103     movl   %eax, %ecx
104     jne    .LBB1_22
105 .LBB1_17:
```

Assembly (x86-64 clang (experimental P1144) -O3 -fomit-frame-pointers):

```
19     callq  operator new(unsigned long, void*)
20     movq   %rax, %r15
21     leaq   (%rax,%r12), %rbp
22     addq   $1600, %r15 # imm = 0x6400
23     movq   %rbp, %r13
24     testq  %r12, %r12
25     jle    .LBB0_3
26     subq   %r12, %r13
27     movq   %r13, %rdi
28     movq   %r14, %rsi
29     movq   %r12, %rcx
30     callq  memcpy@PLT
31 .LBB0_3:
32     movq   %r13, (%rbx)
33     movq   %rbp, 8(%rbx)
34     movq   %r15, 16(%rbx)
35     testq  %r14, %r14
36     je     .LBB0_4
```

Output (0/0) clang version 8.0.0  
(<https://github.com/Quuxplusone/clang>  
fe01be88b1a4cd75fc6467eeb001a99b83035b3a)

54 lines of assembly

159 lines of assembly



**BUT WAIT!**

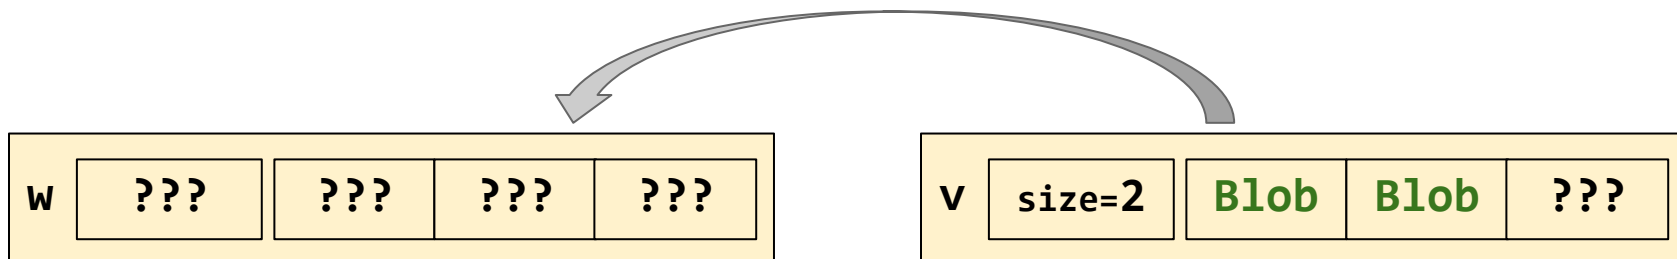


**THERE'S MORE!**

# Application #2 for relocatability

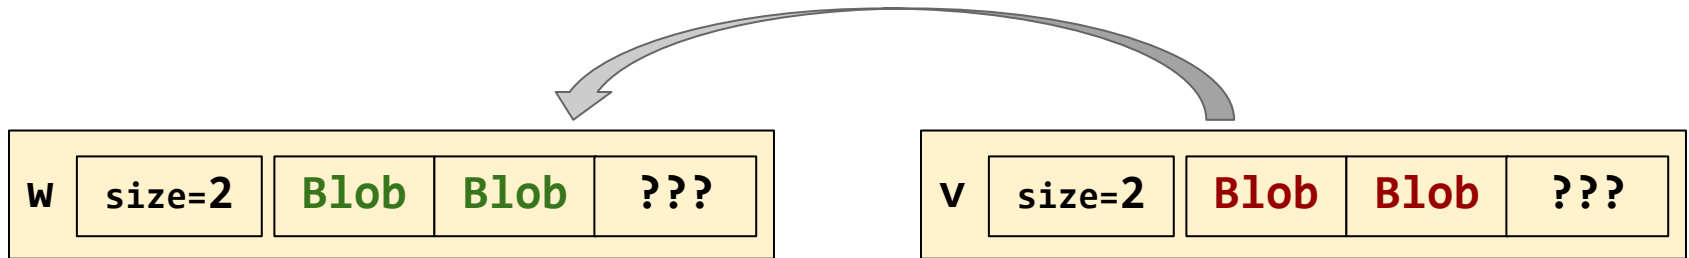
A reliable way of detecting "trivial relocatability" permits optimizing the move operations of `fixed_capacity_vector`:

```
fixed_capacity_vector<Blob, 3> v = { ... };  
auto w = std::move(v);
```



# Application #2 for relocatability

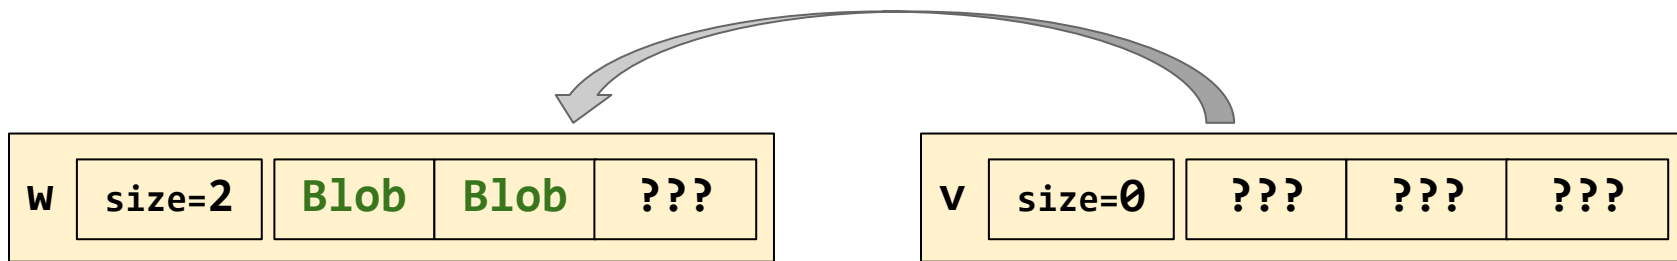
```
fixed_capacity_vector(fixed_capacity_vector&& rhs) {  
    uninitialized_move(rhs.begin(), rhs.end(), begin());  
    size_ = rhs.size_;  
}
```



**MOVING IS  
INEFFICIENT AND BAD**

# Application #2 for relocatability

```
fixed_capacity_vector(fixed_capacity_vector&& rhs) {  
    uninitialized_relocate(rhs.begin(), rhs.end(), begin());  
    size_ = std::exchange(rhs.size_, 0);  
}
```



**RELOCATING IS  
EFFICIENT AND GOOD**

# Moving fixed\_capacity\_vector

```
Save/Load + Add new... CppInsights C++  
  
int size() const { return size_; }  
bool empty() const { return size() == 0; }  
fixed_capacity_vector(fixed_capacity_vec  
#if SLOW  
    std::uninitialized_move(rhs.begin(),  
    size_ = rhs.size_;  
#else  
    std::uninitialized_relocate(rhs.beg  
    size_ = std::exchange(rhs.size_, 0);  
#endif  
}  
~fixed_capacity_vector() {  
    std::destroy(begin(), end());  
}  
};  
  
#include <future>  
using T = std::promise<int>;  
  
fixed_capacity_vector<T, 100>  
test(fixed_capacity_vector<T, 100> vec) {  
    vec.emplace_back();  
    return vec;  
}
```



x86-64 clang (experimental) P1144

A

11010 .LX0: lib.f: .text // s+ Int

Libraries + Add new... Add tool...

```
20 addq $1, %rbx  
21 movl %ebx, (%r15)  
22 movl $0, (%r14)  
23 testl %ebx, %ebx  
24 je .LBB0_2  
25 movq %rbx, %rdx  
26 shlq $3, %rdx  
27 leaq 8(%r15), %rsi  
28 leaq 8(%r14), %rdi  
29 callq memcpy  
30 .LBB0_2:  
31 movl $0, (%r15)  
32 movl %ebx, (%r14)  
33 movq %r14, %rax  
34 popq %rbx  
35 popq %r14  
36 popq %r15  
37 retq  
38 std::_1::__assoc_sub_state::~__as  
39 pushq %rbx  
40 movq %rdi, %rbx
```

87 lines  
of assembly

x86-64 clang (trunk) -O3 -

A

11010 .LX0: lib.f: .text //

Libraries + Add new... Add tool...

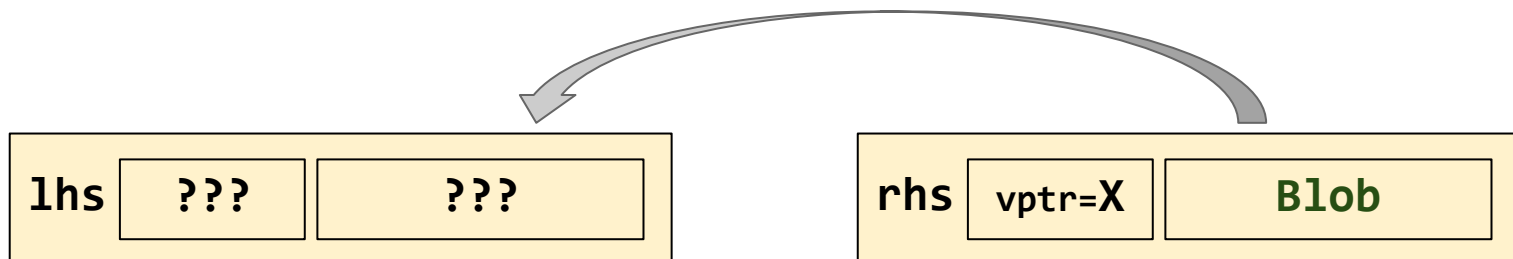
```
24 movl $0, (%r15)  
25 testl %r12d, %r12d  
26 je .LBB0_1  
27 shlq $3, %r12  
28 leaq 8(%r14), %r13  
29 xorl %ebx, %ebx  
30 .LBB0_3: # =>This Inner Loop  
31 leaq (%rbx,%r13), %rdi  
32 movq (%r13,%rbx), %rax  
33 movq %rax, 8(%r15,%rbx)  
34 movq $0, (%r13,%rbx)  
35 callq std::_1::promise<int  
36 addq $8, %rbx  
37 cmpq %rbx, %r12  
38 jne .LBB0_3  
39 movl (%r14), %eax  
40 jmp .LBB0_5  
41 .LBB0_1:  
42 xorl %eax, %eax  
43 .LBB0_5:  
44 movl $0, (%r14)  
45 movl %eax, (%r15)  
46 movq %r15, %rax
```

184 lines  
of assembly

# Application #3 for relocatability

Consider the move-constructor of `std::function`.

What options do we have, to get the data from point rhs to point lhs?

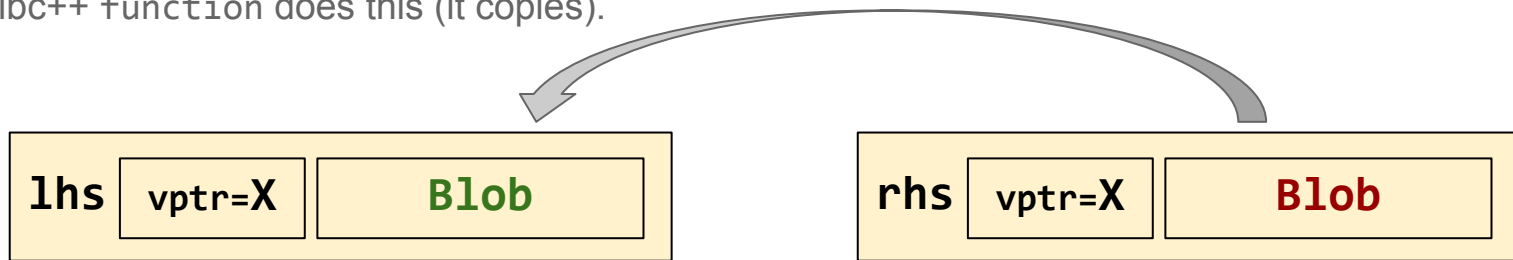


# Application #3 for relocatability

```
function(function&& rhs) {  
    rhs.vptr_->move(rhs_->storage_, this->storage_);  
}
```

We could just move (or copy) the data, leaving rhs engaged.

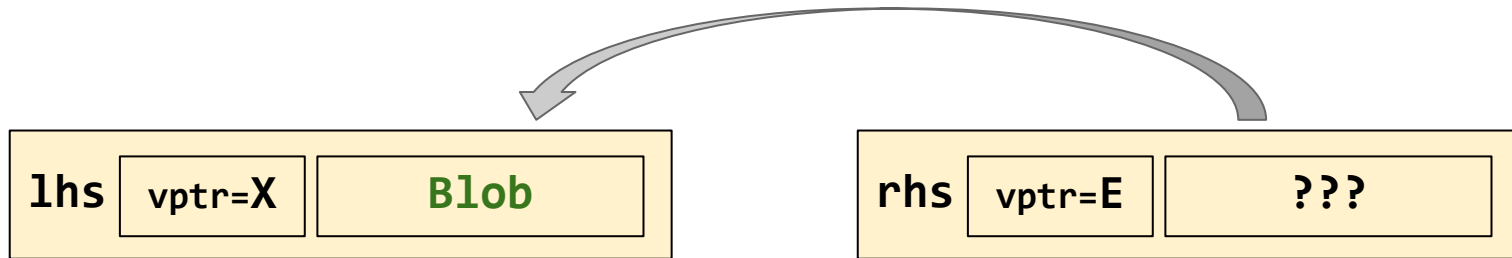
libc++ function does this (it copies).



**MOVING IS  
INEFFICIENT AND BAD**

# Application #3 for relocatability

```
function(function&& rhs) {  
    rhs.vptr_->move(rhs->storage_, this->storage_);  
    rhs.vptr_->destroy(rhs->storage_);  
    this->vptr_ = std::exchange(rhs.vptr_, &empty_vtable);  
}
```

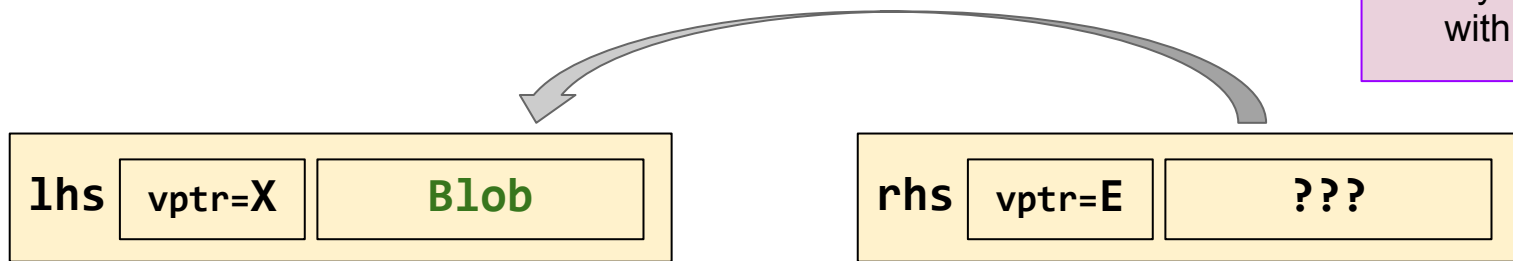


**MOVING IS STILL  
INEFFICIENT AND BAD**



# Application #3 for relocatability

```
function(function&& rhs) {  
    rhs.vptr_->relocate(rhs->storage_, this->storage_);  
    this->vptr_ = std::exchange(rhs.vptr_, &empty_vtable);  
}
```



**RELOCATING IS  
EFFICIENT AND GOOD**

# Now we can do something cool...

Look at this line:

```
rhs.vptr_->relocate(rhs->storage_, this->storage_);
```

When we “type-erase” T into a relocate function, we generally have to do

```
[](void *src, void *dest) {  
    T& from = *(T*)src;  
    ::new (dest) T(std::move(from));  
    from.~T();  
}
```

Different code for each different T. Lots of linker symbols.

# Now we can do something cool...

But a reliable way of detecting "**trivial** relocatability" permits deduplicating the codepaths for all trivially relocatable Ts of a given size!

```
rhs.vptr_->relocate(rhs->storage_, this->storage_);
```

When we “type-erase” a trivially relocatable T, we can just use

```
[](void *src, void *dest) {  
    memcpy(dest, src, sizeof(T));  
}
```

which means we can use the **same single piece of code** for double, int\*, std::unique\_ptr<int>, and so on. Less work for the linker!

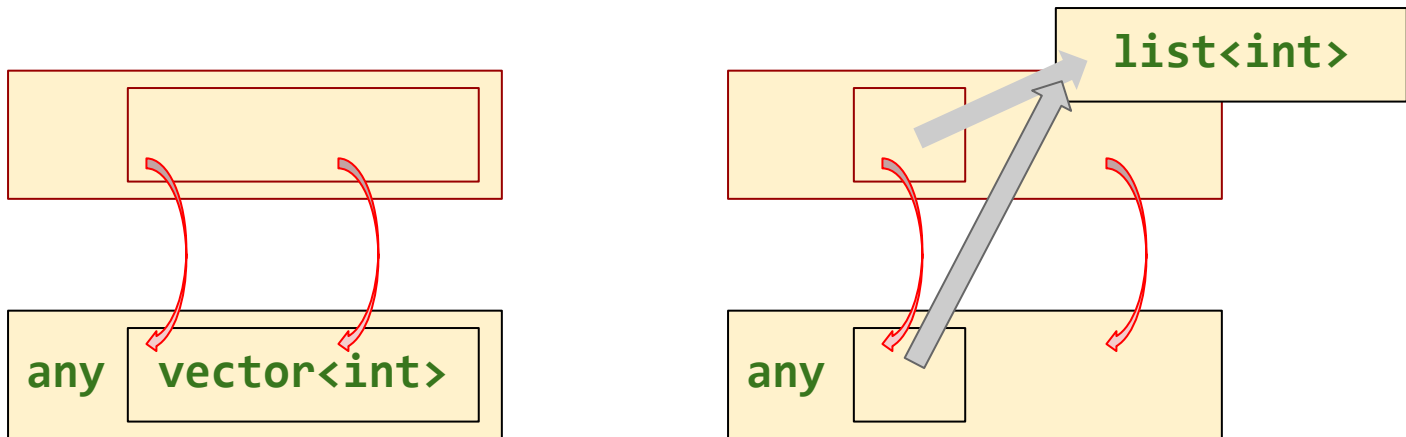
# Application #4 for relocatability

With a reliable way of detecting "trivial relocatability,"  
we can **avoid SBO for wrappees that are not trivially relocatable**.  
Then the wrapper itself becomes trivially relocatable!



# Application #4 for relocatability

With a reliable way of detecting "trivial relocatability,"  
we can **avoid using our SBO for wrappees that are not trivially relocatable**.  
Then the wrapper itself becomes trivially relocatable!



# Application #5 for relocatability

- Consider this naïve user-defined class type.  
It has several members, all of trivially relocatable types.
- I call it “naïve” because it follows the Rule of Zero and does not define a custom ADL swap overload.
- I wish all my code were naïve!

```
struct Pair {  
    std::string first;  
    std::vector<int> second;  
};
```

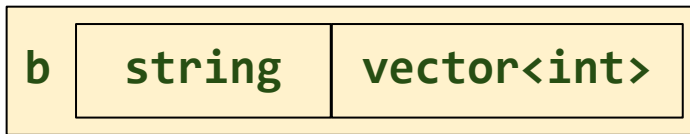
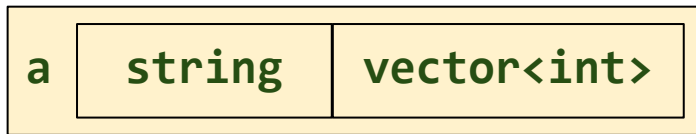
```
Pair a, b;  
std::swap(a, b);
```

# Application #5 for relocatability

```
struct Pair {  
    std::string first;  
    std::vector<int> second;  
};
```

```
Pair a, b;  
std::swap(a, b);
```

```
template<class T>  
void swap(T& a, T& b) {  
    T temp = std::move(a);  
    a = std::move(b);  
    b = std::move(temp);  
}
```

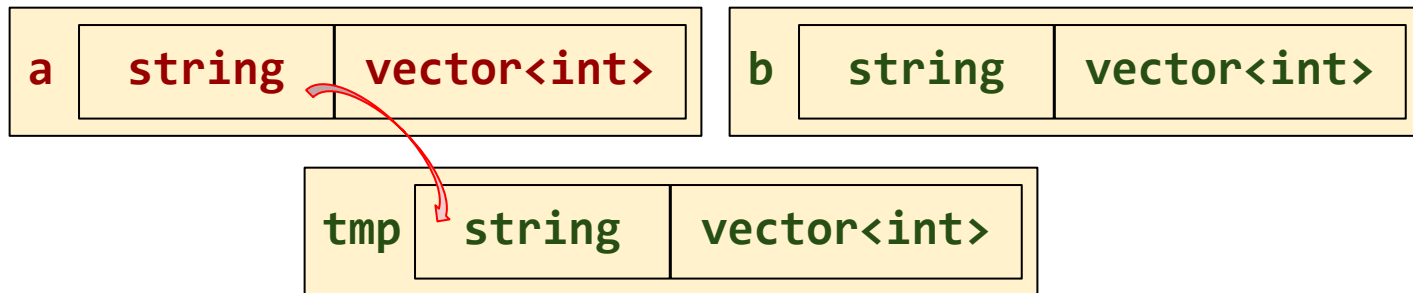


# Application #5 for relocatability

```
struct Pair {  
    std::string first;  
    std::vector<int> second;  
};
```

```
Pair a, b;  
std::swap(a, b);
```

```
template<class T>  
void swap(T& a, T& b) {  
    T temp = std::move(a);  
    a = std::move(b);  
    b = std::move(temp);  
}
```



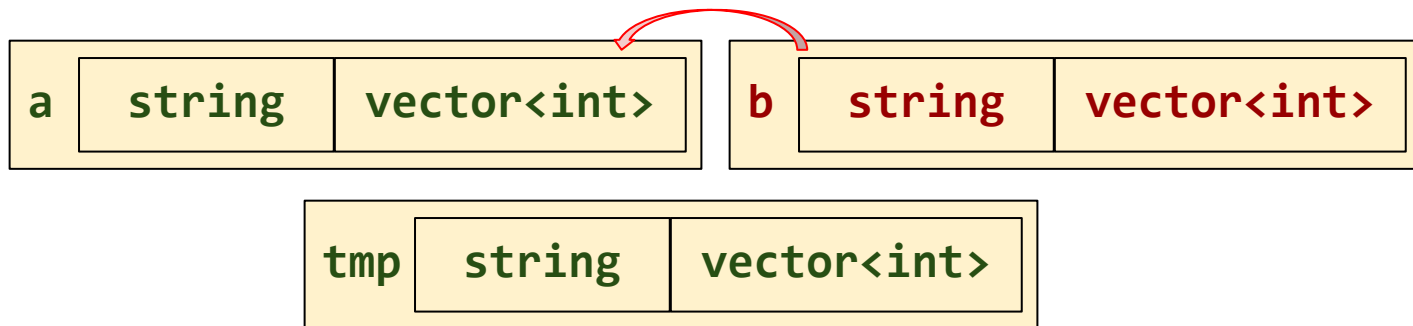


# Application #5 for relocatability

```
struct Pair {  
    std::string first;  
    std::vector<int> second;  
};
```

```
Pair a, b;  
std::swap(a, b);
```

```
template<class T>  
void swap(T& a, T& b) {  
    T temp = std::move(a);  
    a = std::move(b);  
    b = std::move(temp);  
}
```

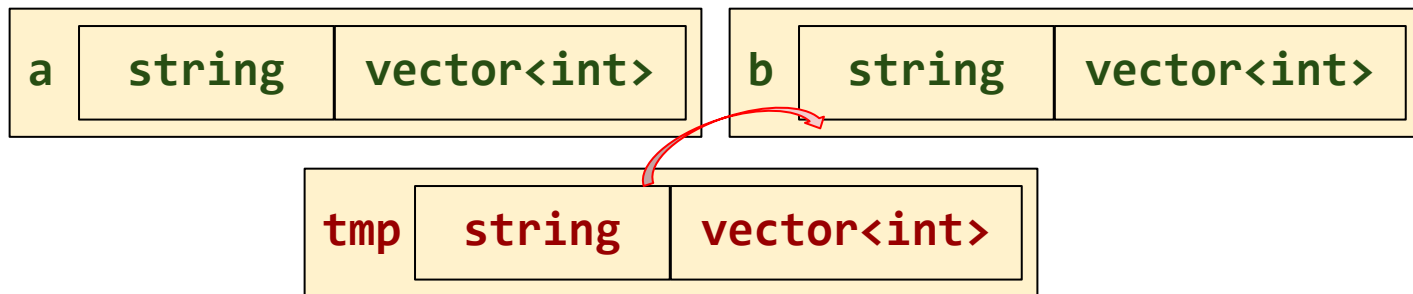


# Application #5 for relocatability

```
struct Pair {  
    std::string first;  
    std::vector<int> second;  
};
```

```
Pair a, b;  
std::swap(a, b);
```

```
template<class T>  
void swap(T& a, T& b) {  
    T temp = std::move(a);  
    a = std::move(b);  
    b = std::move(temp);  
}
```

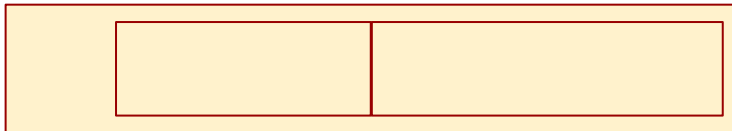
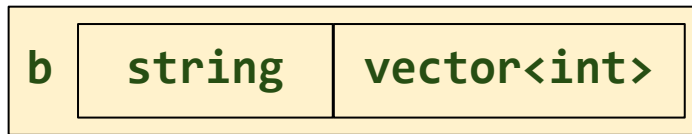
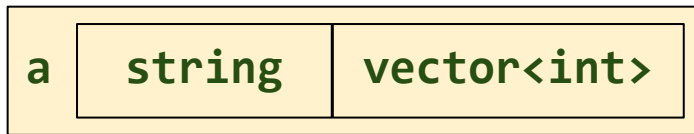


# Application #5 for relocatability

```
struct Pair {  
    std::string first;  
    std::vector<int> second;  
};
```

```
Pair a, b;  
std::swap(a, b);
```

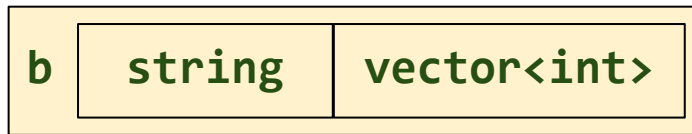
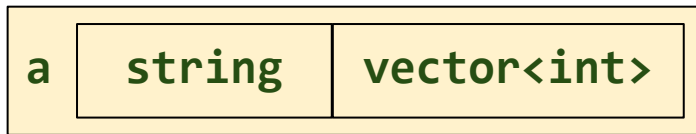
```
template<class T>  
void swap(T& a, T& b) {  
    T temp = std::move(a);  
    a = std::move(b);  
    b = std::move(temp);  
}
```



# But what *should* happen?

```
struct Pair {  
    std::string first;  
    std::vector<int> second;  
};
```

```
Pair a, b;  
std::swap(a, b);
```

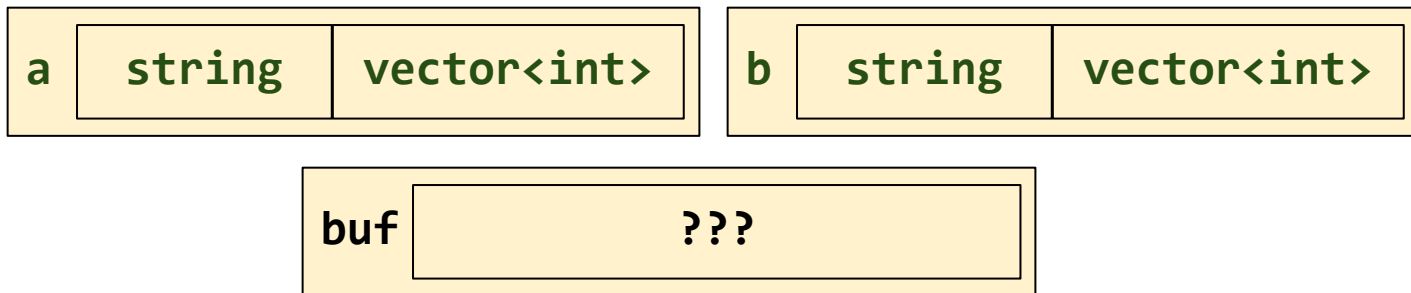


# But what *should* happen?

```
struct Pair {  
    std::string first;  
    std::vector<int> second;  
};
```

```
Pair a, b;  
std::swap(a, b);
```

```
template<class T>  
void swap(T& a, T& b) {  
    alignas(T) char buf[sizeof(T)];  
    std::relocate_at(&a, (T*)buf);  
    std::relocate_at(&b, &a);  
    std::relocate_at((T*)buf, &b);  
}
```

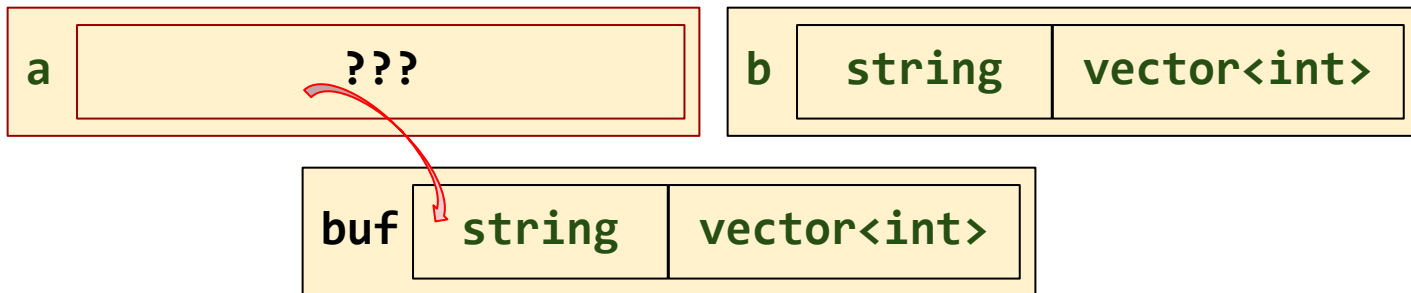


# But what *should* happen?

```
struct Pair {  
    std::string first;  
    std::vector<int> second;  
};
```

```
Pair a, b;  
std::swap(a, b);
```

```
template<class T>  
void swap(T& a, T& b) {  
    alignas(T) char buf[sizeof(T)];  
    std::relocate_at(&a, (T*)buf);  
    std::relocate_at(&b, &a);  
    std::relocate_at((T*)buf, &b);  
}
```

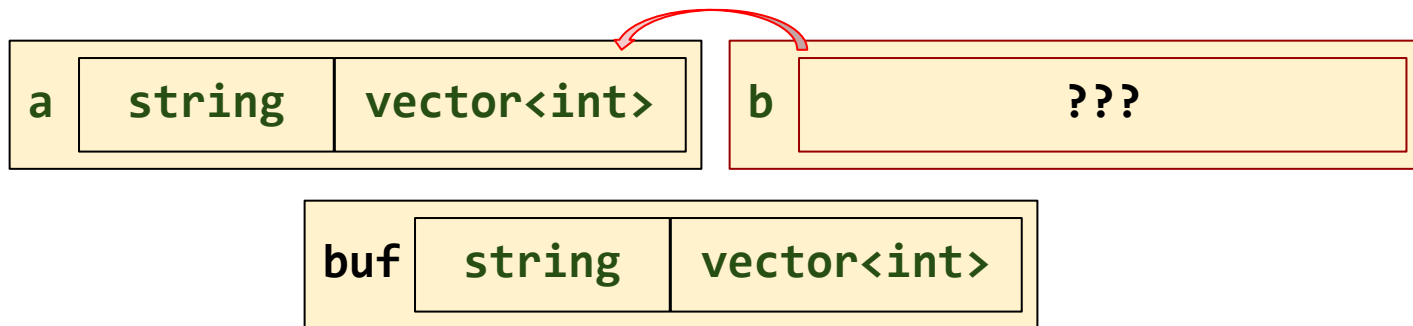


# But what *should* happen?

```
struct Pair {  
    std::string first;  
    std::vector<int> second;  
};
```

```
Pair a, b;  
std::swap(a, b);
```

```
template<class T>  
void swap(T& a, T& b) {  
    alignas(T) char buf[sizeof(T)];  
    std::relocate_at(&a, (T*)buf);  
    std::relocate_at(&b, &a);  
    std::relocate_at((T*)buf, &b);  
}
```

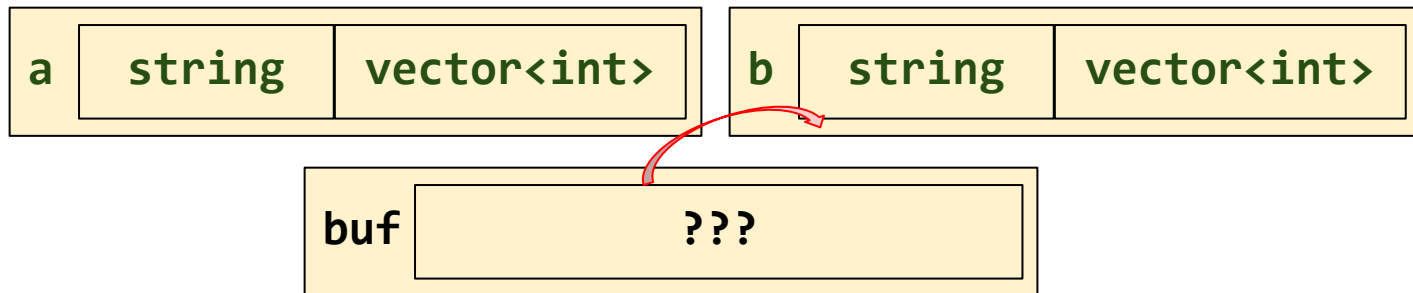


# But what *should* happen?

```
struct Pair {  
    std::string first;  
    std::vector<int> second;  
};
```

```
Pair a, b;  
std::swap(a, b);
```

```
template<class T>  
void swap(T& a, T& b) {  
    alignas(T) char buf[sizeof(T)];  
    std::relocate_at(&a, (T*)buf);  
    std::relocate_at(&b, &a);  
    std::relocate_at((T*)buf, &b);  
}
```





# Comparison of swap approaches

```
template<class T>
void swap(T& a, T& b) {
    T temp = std::move(a);
    a = std::move(b);
    b = std::move(temp);
}
```

- 4 invocations of 3 different operations
- Usually, **none** of the 3 are trivial

```
template<class T>
void swap(T& a, T& b) {
    alignas(T) char buf[sizeof(T)];
    std::relocate_at(&a, (T*)buf);
    std::relocate_at(&b, &a);
    std::relocate_at((T*)buf, &b);
}
```

- Just 3 invocations of 1 operation
- ...which ***is usually trivial***

# Benchmarking `std::rotate`

```
1 #include <algorithm>
2 #include <string>
3 #include <vector>
4
5 struct Pair {
6     int key;
7     std::string first;
8     std::vector<int> second;
9 };
10
11 void test(Pair *first, int n, int k) {
12     std::rotate(first, first + n, first + k)
13 }
```



Rewriting swap for an enormous class of types affects codegen for all swap-based algorithms. My libc++ doesn't change a single line of `std::rotate`!

x86-64 clang (experimental P1144) -O3 -std=c++11

Libraries + Add new... Add tool...

```
111 movups %xmm2, 32(%rdi)
112 movups %xmm1, 16(%rdi)
113 movups %xmm0, (%rdi)
114 movq -8(%rsp), %rax
115 movq %rax, 48(%rcx)
116 movaps -56(%rsp), %xmm0
117 movaps -40(%rsp), %xmm1
118 movaps -24(%rsp), %xmm2
119 movups %xmm2, 32(%rcx)
120 movups %xmm1, 16(%rcx)
121 movups %xmm0, (%rcx)
122 leaq 56(%rdi), %rsi
123 addq $56, %rcx
124 cmpq %r8, %rcx
125 jne .LBB0_13
126 movq %r11, %rcx
127 movq %rsi, %rdi
128 cmpq %rsi, %r11
129 jne .LBB0_10
130 .LBB0_12:
131 retq
```

131 lines of assembly

x86-64 clang (trunk) -O3 -std=c++11

Libraries + Add new... Add tool...

```
154 movq %r15, 24(%r14)
155 movq $0, 15(%rsp)
156 movq $0, 8(%rsp)
157 movq 32(%r14), %rdi
158 testq %rdi, %rdi
159 je .LBB1_8
160 movq %rdi, 16(%r14)
161 callq operator delete(void*)
162 movaps %xmm0, %xmm2
163 movups %xmm0, (%rbp)
164 movq $0, 16(%rbp)
165 .LBB1_8:
166 movaps 32(%rsp), %xmm0 # 16
167 movups %xmm0, 32(%r14)
168 movq %r12, 48(%r14)
169 addq $56, %rsp
170 popq %rbx
171 popq %r12
172 popq %r13
173 popq %r14
174 popq %r15
175 popq %rbp
176 retq
```

176 lines of assembly

**So by this point you're  
probably thinking...**



**GIVE IT TO ME**

# Prior Art

Folly, BSL, EASTL

# Prior art in Electronic Arts EASTL

[https://github.com/electronicarts/EASTL/blob/master/doc/  
EASTL%20Best%20Practices.html](https://github.com/electronicarts/EASTL/blob/master/doc/EASTL%20Best%20Practices.html)

If the user has a class that is relocatable (i.e. can safely use `memcpy` to copy values), the user can use the `EASTL_DECLARE_TRIVIAL_RELOCATE` declaration to tell the compiler ... This will automatically significantly speed up some containers and algorithms that use that class.

```
EASTL_DECLARE_TRIVIAL_RELOCATE(Widget);
```

```
vector<Widget> wVector;
```

```
wVector.erase(wVector.begin()); // This operation will use memcpy.
```

`has_trivial_relocate` is ... very useful in allowing for the generation of optimized object moving operations. It is similar to the `is_pod` type trait, but goes further and allows non-POD classes to be categorized as relocatable. Such categorization is something that no compiler can do, as only the user can know if it is such. Thus `EASTL_DECLARE_TRIVIAL_RELOCATE` is provided to allow the user to give the compiler a hint.

# Prior art in Bloomberg BSL

`bslmf::IsBitwiseMoveable` ... allows generic code to determine whether `TYPE` can be destructively moved using `memcpy`. Given a pointer `p1` to an object of type `TYPE`, and a pointer `p2` of the same type pointing to allocated but uninitialized storage, a destructive move from `p1` to `p2` comprises the following pair of operations:

```
new ((void*) p2) TYPE(*p1); // OR new ((void*) p2) TYPE(std::move(*p1));  
p1->~TYPE();
```

An object of type `TYPE` is ***bitwise moveable*** if the above operation can be replaced by the following operation without affecting correctness:

```
std::memcpy(p2, p1, sizeof(TYPE));
```

# Prior art in Bloomberg BSL

```
class Pair {  
    bsl::string first;  
    bsl::vector<int> second;  
public:  
    BSLMF_NESTED_TRAIT_DECLARATION(Pair,  
                                    bslmf::IsBitwiseMoveable);  
};  
  
static_assert(bslmf::IsBitwiseMoveable<Pair>::value);
```



**Folly has by far  
the most quotable  
explanation of  
relocation semantics.**

# Prior art in Facebook Folly

<https://github.com/facebook/folly/blob/master/folly/docs/FBVector.md>  
[#object-relocation](#)

C++'s assumption of non-relocatable values hurts everybody for the benefit of a few questionable designs. The issue is that moving a C++ object "by the book" entails (a) creating a new copy from the existing value; (b) destroying the old value. This is quite vexing and violates common sense. Consider this hypothetical conversation between Captain Picard and an incredulous alien:

Incredulous Alien: "So, this teleporter, how does it work?"

Picard: "It beams people and arbitrary matter from one place to another."

Incredulous Alien: "Hmmm... is it safe?"

Picard: "Yes, but earlier models were a hassle. They'd clone the person to another location. Then the teleporting chief would have to shoot the original. Ask O'Brien, he was an intern during those times. A bloody mess, that's what it was."

# Prior art in Facebook Folly

- I'm going to pick on Folly here. I'm going to show you its bugs.
- This is **not** because Folly is poor-quality code!
- It's because Folly is highly **comprehensible** code.

There are two ways of constructing a software design: One way is to make it so simple that there are *obviously no deficiencies*, and the other way is to make it so complicated that there are *no obvious deficiencies*. The first method is far more difficult.

—C.A.R. Hoare

# Prior art in Facebook Folly

- I'm going to pick on Folly here. I'm going to show you its bugs.
- This is **not** because Folly is poor-quality code!
- It's because Folly is highly **comprehensible** code.

There are two ways of constructing a software design: One way is to make it so simple that *its deficiencies are obvious*, and the other way is to make it so complicated that *its deficiencies are not obvious*. The first method is far more difficult.

—with apologies to C.A.R. Hoare

# Prior art in Facebook Folly

<https://github.com/facebook/folly/blob/master/folly/docs/FBVector.md>  
#object-relocation

`folly::fbvector<Widget>` will not compile unless you provide a “warrant” that `Widget` is relocatable.

Here’s what a warrant looks like, according to the Folly docs:

```
namespace folly {  
    struct IsRelocatable<Widget> : std::true_type {};  
}
```

# Two downsides of explicit warrants

With explicit warrants, you need to know ***how*** to write one.

In each of our three case studies, this requires

- using a library-specific macro, and/or
- knowing the C++ syntax for explicit template specialization.

Did you notice that the Folly docs actually get the syntax wrong?

With explicit warrants, you need to know ***when*** to write one.

- You have to know whether your type actually *is* trivially relocatable.

Even the developers get it wrong!

# Get the warrant wrong

<https://github.com/facebook/folly/issues/889>

Folly provides these warrants in their "Traits.h" header:

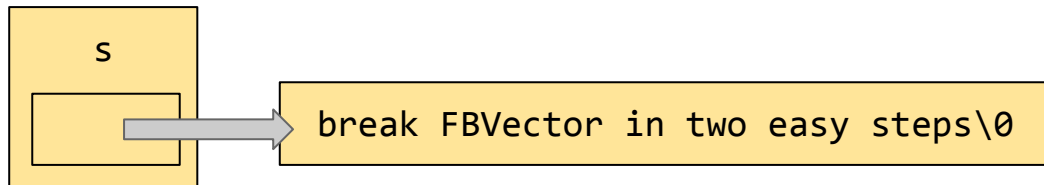
```
#ifndef _MSC_VER
FOLLY_ASSUME_FBVECTOR_COMPATIBLE_2(std::vector)
FOLLY_ASSUME_FBVECTOR_COMPATIBLE_2(std::deque)
FOLLY_ASSUME_FBVECTOR_COMPATIBLE_2(std::unique_ptr)
FOLLY_ASSUME_FBVECTOR_COMPATIBLE_1(std::shared_ptr)
FOLLY_ASSUME_FBVECTOR_COMPATIBLE_1(std::function)
#endif
```

Unfortunately, if using libc++, the warrant for `std::function` is simply a lie; & the warrant for `std::unique_ptr<T, D>` can cause problems for some D.

# How to break FBVector

```
#include <folly/FBVector.h>
```

```
int main() {  
    std::string s = "break FBVector in two easy steps";  
    folly::fbvector<std::function<int()>> v;  
    v.push_back([s]() { std::cout << s << std::endl; });  
    v.reserve(v.capacity() + 1);  
    v[0]();  
}
```

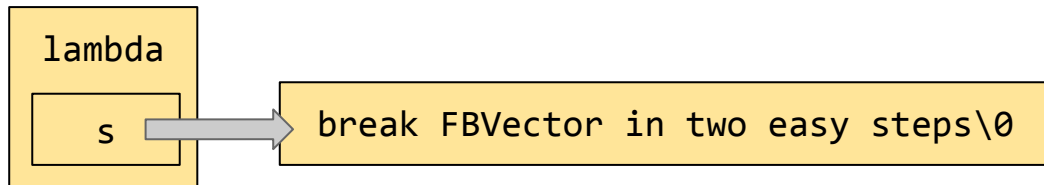




# How to break FBVector

```
#include <folly/FBVector.h>
```

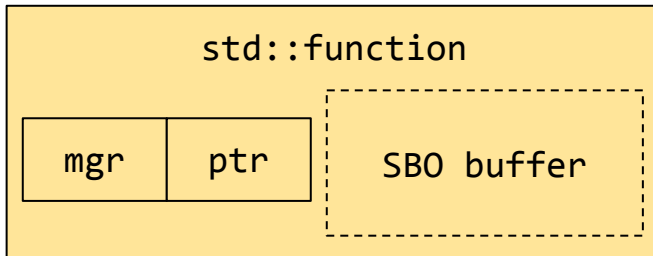
```
int main() {  
    std::string s = "break FBVector in two easy steps";  
    folly::fbvector<std::function<int()>> v;  
    v.push_back([s]() { std::cout << s << std::endl; });  
    v.reserve(v.capacity() + 1);  
    v[0]();  
}
```



# How to break FBVector

```
#include <folly/FBVector.h>
```

```
int main() {  
    std::string s = "break FBVector in two easy steps";  
    folly::fbvector<std::function<int()>> v;  
    v.push_back([s]() { std::cout << s << std::endl; });  
    v.reserve(v.capacity() + 1);  
    v[0]();  
}
```

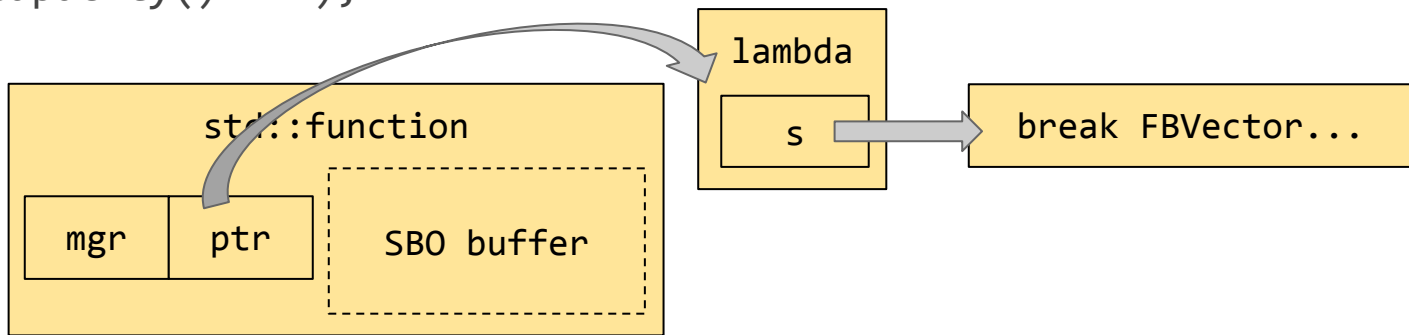


libc++'s `std::function` has a pointer to a “manager” function (basically a manual vtable), a pointer to the wrapped data, and also a buffer for the Small Buffer Optimization.

# How to break FBVector

```
#include <folly/FBVector.h>
```

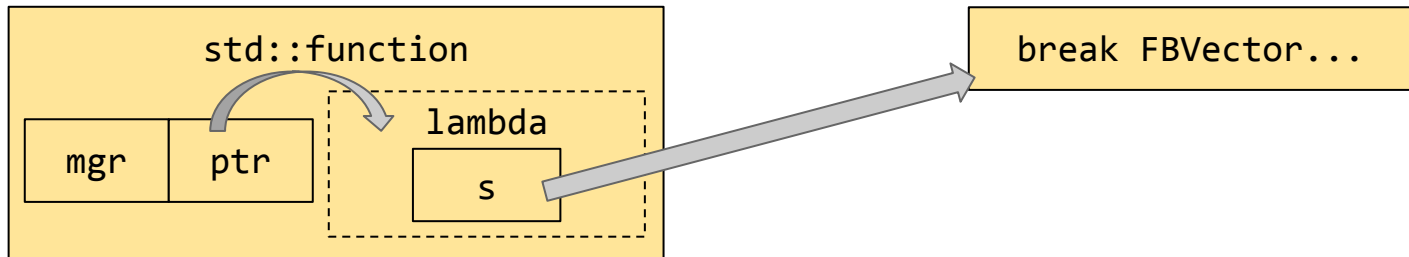
```
int main() {  
    std::string s = "break FBVector in two easy steps";  
    folly::fbvector<std::function<int()>> v;  
    v.push_back([s]() { std::cout << s << std::endl; });  
    v.reserve(v.capacity() + 1);  
    v[0]();  
}
```



# How to break FBVector

```
#include <folly/FBVector.h>
```

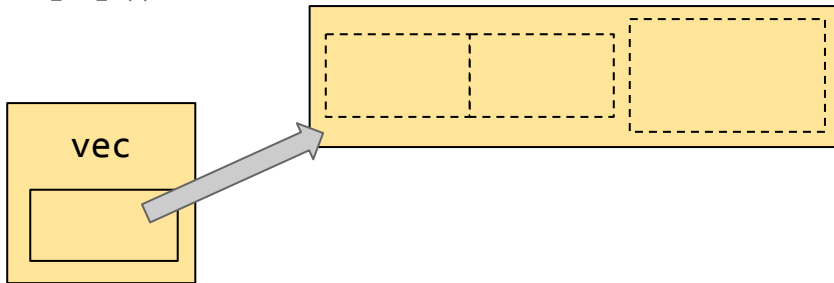
```
int main() {  
    std::string s = "break FBVector in two easy steps";  
    folly::fbvector<std::function<int()>> v;  
    v.push_back([s]() { std::cout << s << std::endl; });  
    v.reserve(v.capacity() + 1);  
    v[0]();  
}
```



# How to break FBVector

```
#include <folly/FBVector.h>
```

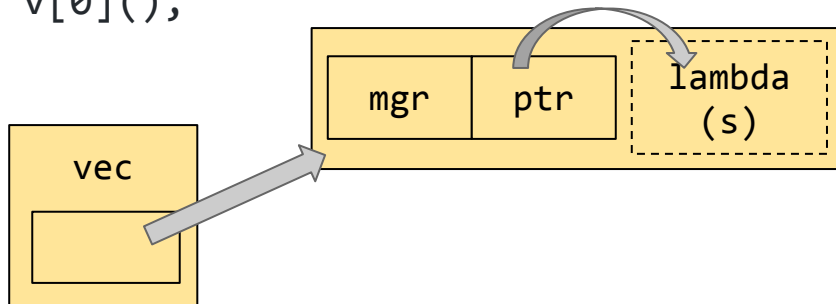
```
int main() {  
    std::string s = "break FBVector in two easy steps";  
    folly::fbvector<std::function<int()>> v;  
    v.push_back([s]() { std::cout << s << std::endl; });  
    v.reserve(v.capacity() + 1);  
    v[0]();  
}
```



# How to break FBVector

```
#include <folly/FBVector.h>
```

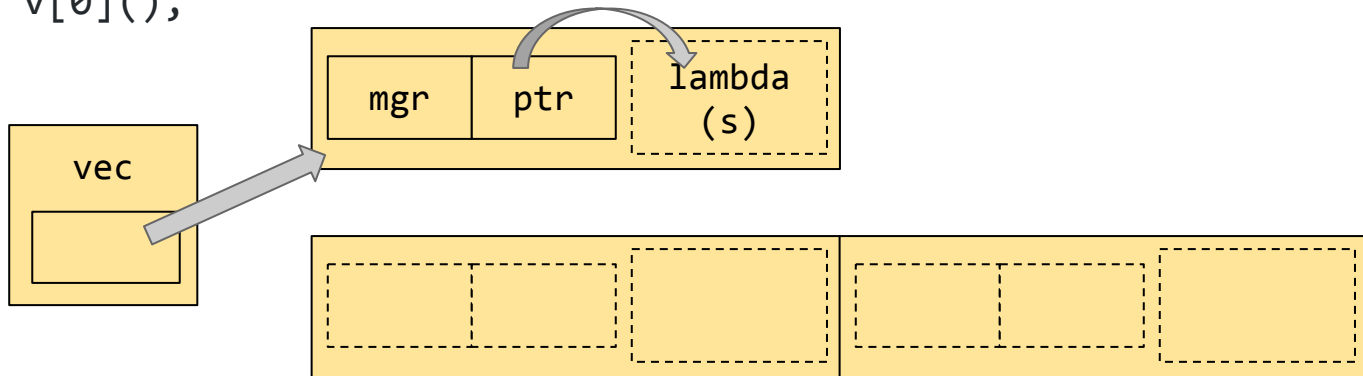
```
int main() {  
    std::string s = "break FBVector in two easy steps";  
    folly::fbvector<std::function<int()>> v;  
    v.push_back([s]() { std::cout << s << std::endl; });  
    v.reserve(v.capacity() + 1);  
    v[0]();  
}
```



# How to break FBVector

```
#include <folly/FBVector.h>
```

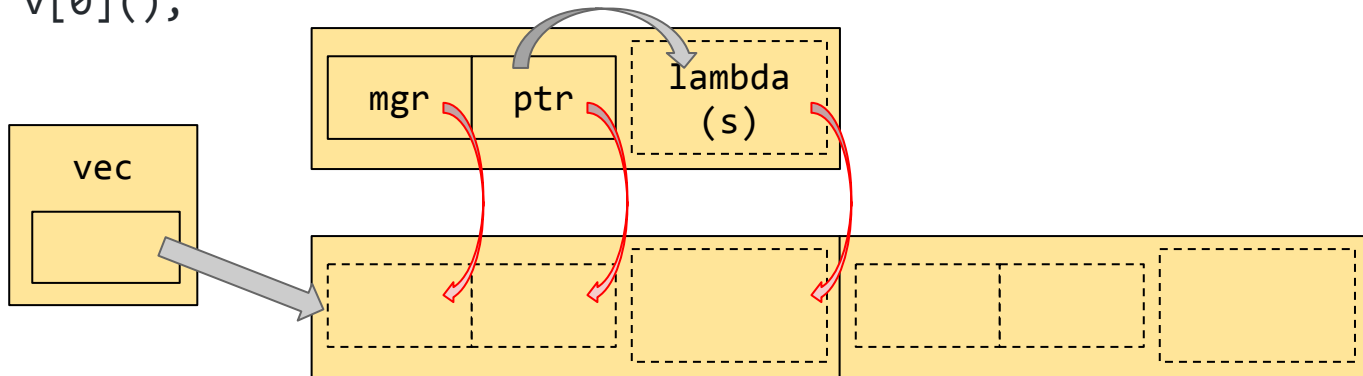
```
int main() {  
    std::string s = "break FBVector in two easy steps";  
    folly::fbvector<std::function<int()>> v;  
    v.push_back([s]() { std::cout << s << std::endl; });  
    v.reserve(v.capacity() + 1);  
    v[0]();  
}
```



# How to break FBVector

```
#include <folly/FBVector.h>
```

```
int main() {  
    std::string s = "break FBVector in two easy steps";  
    folly::fbvector<std::function<int()>> v;  
    v.push_back([s]() { std::cout << s << std::endl; });  
    v.reserve(v.capacity() + 1);  
    v[0]();  
}
```

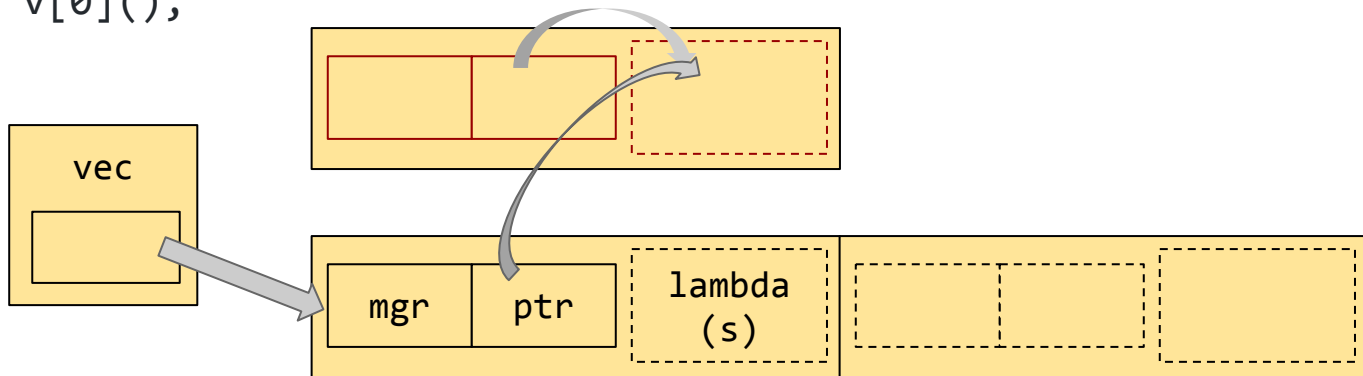




# How to break FBVector

```
#include <folly/FBVector.h>
```

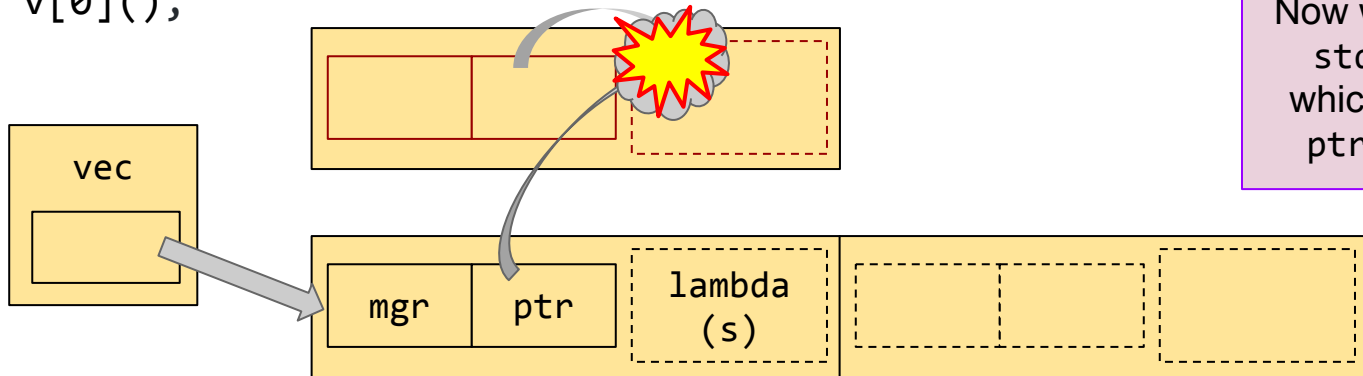
```
int main() {  
    std::string s = "break FBVector in two easy steps";  
    folly::fbvector<std::function<int()>> v;  
    v.push_back([s]() { std::cout << s << std::endl; });  
    v.reserve(v.capacity() + 1);  
    v[0]();  
}
```



# How to break FBVector

```
#include <folly/FBVector.h>
```

```
int main() {  
    std::string s = "break FBVector in two easy steps";  
    folly::fbvector<std::function<int()>> v;  
    v.push_back([s]() { std::cout << s << std::endl; });  
    v.reserve(v.capacity() + 1);  
    v[0]();  
}
```



Now we try to call the `std::function`, which dereferences `ptr`, and... boom.

# Hand-written warrants are bug-prone

- `folly::fbvector` relies on hand-written warrants
  - written by the Folly developers (who get it wrong)
  - written by the client developer (who will get it wrong)
- Hand-written warrants are usually wrong, because they require that the ***developer*** know internal details of the ***library vendor***'s implementation.
- And they don't scale.

# Hand-written warrants don't scale

```
#include <folly/FBVector.h>
```

```
struct Pair {  
    std::string first;  
    std::vector<int> second;  
};
```

```
folly::fbvector<Pair> v;  // ERROR: no warrant
```

Folly “helpfully” prevents `fbvector<Pair>` from compiling. EASTL and BSL let it compile and just don’t do the optimization — so this code is silently pessimized.

Adding warrants is ***too much work*** for the developer.

**Enter P1144:  
“Object Relocation  
in terms of  
Move plus Destroy”**



# P1144 proposal in a nutshell

- Targets C++2b (“C++23”)
- Combination of core-language feature, type traits, library algorithms
- By design, ***preserves correctness*** of all C++17 code
- By design, ***preserves conformance*** of all C++17 library implementations (!!)

P1144 requires vendors to implement a few simple algorithms and type-traits, but does not require an ABI break. For example, libstdc++’s `std::string` is not trivially relocatable. This is 100% okay according to P1144.

What matters is that we have a reliable way to ***detect*** its trivial relocatability.

# P1144 proposal in a nutshell

- “Trivially relocatable” becomes a term of art, similar to “trivially copyable,” “trivially destructible,” etc.
- Just like those properties, the compiler exposes a built-in which is wrapped up by the STL into a type trait: `std::is_trivially_relocatable<T>`.
- Just like those properties, the compiler automatically propagates trivial relocatability to every Rule-of-Zero class, according to the trivial relocatability of its bases and members.
- All non-Rule-of-Zero classes are assumed to be non-trivially relocatable. Expert users can explicitly warrant the property via a class attribute: `[[trivially_relocatable]]`.

# P1144 is implemented on Godbolt

- Use the C++ compiler “x86-64 clang (experimental P1144)”
- This is built from stock llvm...
- ...with a branch of the clang front-end that implements `[[trivially_relocatable]]` and `__is_trivially_relocatable(T)`...
- ...and a branch of libc++ that implements P1144’s new library features.





# P1144 is implemented on Godbolt

- Use the C++ compiler “x86-64 clang (experimental P1144)”
- This is built from stock llvm...
- ...with a branch of the clang front-end that implements `[[trivially_relocatable]]` and `__is_trivially_relocatable(T)`...
- ...and a branch of libc++ that implements P1144’s new library features.
- And trivially relocatable `std::any`.

# P1144 is implemented on Godbolt

- Use the C++ compiler “x86-64 clang (experimental P1144)”
- This is built from stock `llvm`...
- ...with a branch of the `clang` front-end that implements `[[trivially_relocatable]]` and `__is_trivially_relocatable(T)`...
- ...and a branch of `libc++` that implements P1144’s new library features.
- And trivially relocatable `std::any`.
- And trivially copyable `std::vector<bool>::iterator`.

# P1144 is implemented on Godbolt

- And the entire `<memory_resource>` header.
- And CTAD deduction guides for all associative and unordered containers.
- And `std::priority_queue::replace_top()`.
- And container adaptors which are conditionally trivially destructible.
- And `string_view` support in `std::regex` (thanks to Mark de Wever!)

All this stuff is maintained in the `trivially-relocatable` branch of [github.com/Quuxplusone/libcxx](https://github.com/Quuxplusone/libcxx).

# You already know the entire feature

All of the optimizations I've shown can be implemented invisibly to the user-programmer. The proof is the Godbolt links/screenshots you've already seen.

Okay, there are two more little wrinkles:

- `[[trivially_relocatable(bool)]]`
- `std::relocate_at` and `std::uninitialized_relocate`

# Conditional trivial relocatability

```
template<class T>
struct ZeroWrap {
    T t;
};

using R = std::vector<int>;
using NR = std::list<int>;

static_assert(std::is_trivially_relocatable_v<ZeroWrap<R>>);
static_assert(!std::is_trivially_relocatable_v<ZeroWrap<NR>>);
```

Follow the Rule of Zero and you get conditional trivial relocatability for free. Here's an example.

# Conditional trivial relocatability

```
template<class T>
struct ExpertWrap {
    std::pair<int*, T> p;
    friend void swap(ExpertWrap& a, ExpertWrap& b) noexcept { std::swap(a.p, b.p); }
    ExpertWrap() { p.first = new int; }
    ExpertWrap(ExpertWrap&& rhs) noexcept : p(rhs.p) { rhs.p.first = nullptr; }
    auto& operator=(ExpertWrap rhs) noexcept { swap(*this, rhs); }
    ~ExpertWrap() { delete p.first; }
};
```

```
using R = std::vector<int>;
using NR = std::list<int>;
```

```
static_assert(!std::is_trivially_relocatable_v<ExpertWrap<R>>);
static_assert(!std::is_trivially_relocatable_v<ExpertWrap<NR>>);
```

Break the Rule of Zero and the compiler will conservatively assume that your type is **not** trivially relocatable. Remember our design goals!

# Conditional trivial relocatability

```
template<class T>
struct [[trivially_relocatable]] ExpertWrap {
    std::pair<int*, T> p;
    friend void swap(ExpertWrap& a, ExpertWrap& b) noexcept { std::swap(a.p, b.p); }
    ExpertWrap() { p.first = new int; }
    ExpertWrap(ExpertWrap&& rhs) noexcept : p(rhs.p) { rhs.p.first = nullptr; }
    auto& operator=(ExpertWrap rhs) noexcept { swap(*this, rhs); }
    ~ExpertWrap() { delete p.first; }
};
```

```
using R = std::vector<int>;
using NR = std::list<int>;
```

```
static_assert(std::is_trivially_relocatable_v<ExpertWrap<R>>);
static_assert(std::is_trivially_relocatable_v<ExpertWrap<NR>>);
```

As with Folly, BSL, or EASTL, if you give the wrong warrant, your code will be wrong. P1144 reduces the *number* of times you use the chainsaw; it doesn't change the *danger level* of the chainsaw.

# Conditional trivial relocatability (1)

```
template<class T, bool TrivReloc = false>
struct ExpertImpl {
    std::pair<int*, T> p;
    friend void swap(ExpertImpl& a, ExpertImpl& b) noexcept { std::swap(a.p, b.p); }
    ExpertWrap() { p.first = new int; }
    ExpertWrap(ExpertImpl&& rhs) noexcept : p(rhs.p) { rhs.p.first = nullptr; }
    auto& operator=(ExpertImpl rhs) noexcept { swap(*this, rhs); }
    ~ExpertImpl() { delete p.first; }
};

template<class T>
struct [[trivially_relocatable]] ExpertImpl<T, true> {
    cut and paste exactly the same code as above
};

template<class T>
struct ExpertWrap : private ExpertImpl<T, std::is_trivially_relocatable_v<T>> {
    this top-level class now follows the Rule of Zero
};
```

P1144R2 supported only this metaprogramming approach.



# Conditional trivial relocatability (2)

```
template<class T>
struct [[trivially_relocatable(std::is_trivially_relocatable_v<T>))] ExpertWrap {
    std::pair<int*, T> p;
    friend void swap(ExpertWrap& a, ExpertWrap& b) noexcept { std::swap(a.p, b.p); }
    ExpertWrap() { p.first = new int; }
    ExpertWrap(ExpertWrap&& rhs) noexcept : p(rhs.p) { rhs.p.first = nullptr; }
    auto& operator=(ExpertWrap rhs) noexcept { swap(*this, rhs); }
    ~ExpertWrap() { delete p.first; }
};

using R = std::vector<int>;
using NR = std::list<int>;

static_assert(std::is_trivially_relocatable_v<ExpertWrap<R>>);
static_assert(!std::is_trivially_relocatable_v<ExpertWrap<NR>>);
```

In P1144R3, I've decided that this syntax removes enough library complexity to be worth the added language complexity.

# Conditional trivial relocatability (3)

```
template<class T>
struct [[maybe_trivially_relocatable]] ExpertWrap {
    std::pair<int*, T> p;
    friend void swap(ExpertWrap& a, ExpertWrap& b) noexcept { std::swap(a.p, b.p); }
    ExpertWrap() { p.first = new int; }
    ExpertWrap(ExpertWrap&& rhs) noexcept : p(rhs.p) { rhs.p.first = nullptr; }
    auto& operator=(ExpertWrap rhs) noexcept { swap(*this, rhs); }
    ~ExpertWrap() { delete p.first; }
};

using R = std::vector<int>;
using NR = std::list<int>;

static_assert(std::is_trivially_relocatable_v<ExpertWrap<R>>);
static_assert(!std::is_trivially_relocatable_v<ExpertWrap<NR>>);
```

John McCall suggested this version.  
Here ExpertWrap is trivially relocatable  
if and only if all its bases and members are.

# Isn't this all undefined behavior?

- It's undefined behavior to memcpy bits and bytes between objects of non-trivially copyable types.
- P1144R1 says let's change that. It is now ***legal*** to memcpy between objects of non-trivially copyable (but trivially relocatable) types. Problem solved!
- But doesn't that permit the user to “make a copy” of a `unique_ptr` via `memcpy`? Copying should be legal only if we can extract a promise from the user not to look at the source object's value afterward.

# Isn't this all undefined behavior?

- P1144R3 adds a library algorithm `std::relocate_at(from, to)`. This algorithm is documented to begin the lifetime of its destination, like `std::construct_at`, and end the lifetime of its source, like `std::destroy_at`.
- For non-trivially relocatable types, this is just move-construct plus destroy.
- For trivially relocatable types, this *can* be implemented as a `memcpy` surrounded by some compiler magic to begin and end lifetimes appropriately.
- In practice, that “compiler magic” is a no-op.


# Similar to std::bless

P0593 std::bless works in a similar way.

```
struct ListNode {  
    ListNode *next;  
    T t;  
};  
  
void example(void *raw_bytes) {  
    std::bless(raw_bytes, sizeof(ListNode));  
    ListNode *node = (ListNode *)raw_bytes;  
    ::new (&node->t) T(); // There's no T there...  
    node->next = nullptr; // ...but there's a next ptr?  
}
```


# Easy implementation of `std::bless`

```
void example(void *raw_bytes) {  
    std::bless(raw_bytes, sizeof(ListNode));  
    ListNode *node = (ListNode *)raw_bytes;  
    ::new (&node->t) T();  
    node->next = nullptr;  
}
```



When the compiler sees this line...

```
void bless(void *p, size_t n);
```




...it has no idea what the body of `bless` looks like, because that's an out-of-line library function.

# Easy implementation of `std::bless`

```
void example(void *raw_bytes) {  
    std::bless(raw_bytes, sizeof(ListNode));  
    ListNode *node = (ListNode *)raw_bytes;  
    ::new (&node->t) T();  
    node->next = nullptr;  
}
```

```
void bless(void *p, size_t n) {  
    ::new (p) ListNode{};  
    ((ListNode*)p)->t.~T();  
}
```




For all the compiler knows, the body of `bless` might look like ***this!*** So it must optimize accordingly.

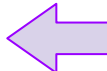
# Easy implementation of `relocate_at`

```
T *relocate_at(T *src, T *dest) {  
    if constexpr (std::is_trivially_relocatable_v<T>) {  
        ::new (dest) T(std::move(*src));  
        src->~T();  
    } else {  
        __triv_reloc_at(dest, src, sizeof(T));  
    }  
    return dest;  
}
```

```
void __triv_reloc_at(void*, void*, size_t);
```



When the compiler sees this line...



...it has no idea what the body of `__triv_reloc_at` looks like.



# Easy implementation of `relocate_at`

```
void __triv_reloc_at(void*, void*, size_t) {  
    if (you_are_relocating_a_string) {  
        ::new (dest) std::string(std::move(*src));  
        src->~basic_string();  
    } else if (you_are_relocating_a_vector) {  
        ::new (dest) std::vector<int>(std::move(*src));  
        src->~vector();  
    } else if (.....  
        .....  
}
```

The compiler must imagine that the body of `__triv_reloc_at` **might** look like this! The compiler is therefore not allowed to make any optimizations that would be incorrect in such a world.

Meanwhile, in our actual world, the linker symbol `__triv_reloc_at` is just an alias for `memcpy`.

**End of Part 1  
Time Check  
Questions?**

**P1144 relocation  
is not...**

# Not Pablo Halpern's N4158

Can I define my own “relocation” operation, to be used by `vector::resize` and so on, which is not quite `memcpy` but also more efficient than move-plus-destroy?

My P1144 says, definitively, ***no***.

`std::relocate_at(from, to)` is a standard library algorithm, like `std::destroy_at`.

Pablo's N4158 says, definitively, ***yes***.

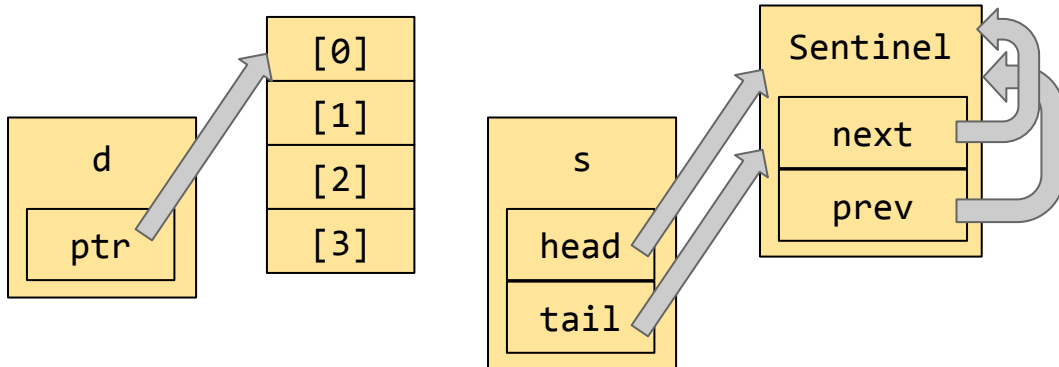
`uninitialized_destructive_move(from, to)` is an ADL customization point, like `swap`.

# Not Pablo Halpern's N4158

N4158 supports some use-cases that P1144 does not.

Consider MSVC's `std::list<T>` or GCC's `std::deque<T>`. They have a “never-completely-null” invariant:

```
std::deque<int> d;  
std::list<int> s;
```



They are trivially relocatable  
but **not** nothrow-move-constructible.

# Not Pablo Halpern's N4158

Vice versa, GCC's `std::list<T>` is nothrow-move-constructible but **not** trivially relocatable.

```
struct Pair {  
    std::deque<int> d;  
    std::list<int> s;  
};
```

*// Assuming libstdc++, then...*

```
static_assert(!std::is_trivially_relocatable_v<Pair>);  
static_assert(!std::is_nothrow_move_constructible_v<Pair>);
```

# Not Pablo Halpern's N4158

```
struct Pair {  
    std::deque<int> d;  
    std::list<int> s;  
};  
  
static_assert(!std::is_trivially_relocatable_v<Pair>);  
static_assert(!std::is_nothrow_move_constructible_v<Pair>);
```

P1144 says: game over, that's all you get.

N4158 says: You can customize the ADL customization point to relocate `d` (trivially) and then relocate `s` (non-trivially but still nothrow-ly). Thus, `Pair` can be made *nothrow-relocatable*.

# Not Pablo Halpern's N4158

```
struct Pair {  
    std::deque<int> d;  
    std::list<int> s;  
    friend void uninitialized_destructive_move(  
        Pair *from, Pair *to) noexcept {  
        using std::uninitialized_destructive_move;  
        uninitialized_destructive_move(&from->d, &to->d);  
        uninitialized_destructive_move(&from->s, &to->s);  
    }  
};
```

```
static_assert(!std::is_nothrow_move_constructible_v<Pair>);  
static_assert(std::is_nothrow_destructive_movable_v<Pair>);
```



# Not Pablo Halpern's N4158

```
struct Pair {  
    std::string first;  
    std::vector<int> second;  
};
```

*// N4158 says you must do this. Danger, Will Robinson!*

```
template<>  
struct std::is_trivially_destructive_movable<Pair>  
    : std::true_type {};  
  
static_assert(std::is_trivially_destructive_movable_v<Pair>);
```

# Not Pablo Halpern's N4158

TLDR:

- N4158 could achieve “nothrow-relocatability” in more cases than P1144
- Therefore it could avoid the vector pessimization in more cases
- But N4158 was ***not simple***
- Relied on an ADL customization point
- Also relied on reopening namespace std to provide explicit warrants
  - `std::is_trivially_destructive_movable`  
was a specialization point like `std::is_error_code_enum`,  
not a type-trait like `std::is_trivially_destructible`

**P1144 relocation  
is not...**

# Not Denis Bider's P0023

```
struct Example {  
    A a;  
    B b;  
  
    >>Example(Example& rhs) :  
        >>a(rhs.a), >>b(rhs.b) {} // or =default  
};  
  
void relocate(Example *from, Example *to) {  
    ::new (to) >>Example(*from);  
}
```

# Not Denis Bider's P0023

TLDR:

- P0023 adds a completely new operation: the “relocator.” (Like a constructor or a destructor, but not the same as either.)
- “Relocation” is a new core-language operation with its own core-language syntax. It is not “the same” as move plus destroy, except by convention.
- P0023 provides a higher-level `std::relocate_or_move(dest, src)` to dispatch between “relocate” (if possible) or “move+destroy,” just like the STL provides `std::move_if_noexcept` to dispatch between “move” (if possible) or “copy.”
- The compiler figures out trivial relocatability the same way it figures out trivial destructibility or trivial copy-constructibility.

# Not Denis Bider's P0023

```
struct Uniq {  
    int *p = nullptr;  
    Uniq() = default;  
    Uniq(Uniq&& rhs) : p(rhs.p) { rhs.p = nullptr; }  
    friend void swap(Uniq& a, Uniq& b) { std::swap(a.p, b.p); }  
    void operator=(Uniq rhs) { swap(*this, rhs); }  
    ~Uniq() { delete p; }  
};
```

```
static_assert(!std::is_relocatable_v<Uniq>);  
static_assert(!std::is_trivially_relocatable_v<Uniq>);
```

# Not Denis Bider's P0023

```
struct Uniq {  
    int *p = nullptr;  
    Uniq() = default;  
    Uniq(Uniq&& rhs) : p(rhs.p) { rhs.p = nullptr; }  
    friend void swap(Uniq& a, Uniq& b) { std::swap(a.p, b.p); }  
    void operator=(Uniq rhs) { swap(*this, rhs); }  
    ~Uniq() { delete p; }  
    >>Uniq(Uniq&) = default;  
};  
  
static_assert(std::is_relocatable_v<Uniq>);  
static_assert(std::is_trivially_relocatable_v<Uniq>);
```

**P1144 relocation  
is not...**



# Not Niall Douglas's P1029

```
struct Uniq {  
    int *p = nullptr;  
    Uniq() = default;  
    Uniq(Uniq&& rhs)  
        : p(rhs.p) { rhs.p = nullptr; }  
    friend void swap(Uniq& a, Uniq& b) { std::swap(a.p, b.p); }  
    void operator=(Uniq rhs) { swap(*this, rhs); }  
    ~Uniq() { delete p; }  
};
```

```
static_assert(!std::is_move_construction_relocating_v<Uniq>);
```

# Not Niall Douglas's P1029

```
struct Uniq {  
    int *p = nullptr;  
    Uniq() = default;  
    [[move_relocates]] Uniq(Uniq&& rhs)  
        : p(rhs.p) { rhs.p = nullptr; }  
    friend void swap(Uniq& a, Uniq& b) { std::swap(a.p, b.p); }  
    void operator=(Uniq rhs) { swap(*this, rhs); }  
    ~Uniq() { delete p; }  
};
```

```
static_assert(std::is_move_construction_relocating_v<Uniq>);
```

# Not Niall Douglas's P1029

My impression is that P1029 is extremely confused.

P1029 adds an attribute `[[move_relocates]]`. You put it on the move-constructor (only), in order to indicate a relationship between the move-constructor, the destructor, and the default constructor.

The relationship is:

“Move-constructing `d` from `s`” is tantamount to “memcpying `s` to `d`, and then setting `s`’s bit-pattern to the bit-pattern of a default-constructed `T`.”

Furthermore (although it’s not clearly explained in the paper):

“Destroying a `T` in the default-constructed state has no side effects.” &/or  
“Default-constructing a `T` has no side effects.”

# Not Niall Douglas's P1029

Furthermore, maybe (it's not clear in the paper):

“No operation on T cares where the T is physically located in memory.”

P1029 uses this last property to claim that any T which has “relocating move-construction” can therefore be passed in CPU registers because it is no longer important to match up the constructors and destructors.

```
struct Widget {  
    Widget() {  
        printf("A Widget constructed at %p...\n", (void*)this);  
    }  
    ~Widget() {  
        printf("...might validly be destroyed at %p.\n", (void*)this);  
    }  
};
```

# Not Niall Douglas's P1029

TLDR:

- P1029 proposes a property warranted by an attribute, just like P1144.
- P1029's property is very complex. It involves a “default-constructed state.” Types which aren't default-constructible cannot be “move-relocating.”
- P1029 deliberately, explicitly, has ABI implications. “Move-relocating” class types are supposed to be passed in registers. For example, a move-relocating `std::unique_ptr<int>` would be returned directly in `%rax`, instead of being returned on the stack.
- Thus, P1029 forces vendors to choose: status quo (no benefit for users), or ABI break (pass-by-register `unique_ptr`)? There's no way to indicate “My type is trivially relocatable but I don't want the trivial ABI.”

**P1144 relocation  
is not...**

# Not `[[clang::trivial_abi]]`

```
struct [[clang::trivial_abi]] Uniq {  
    int *p = nullptr;  
    Uniq() = default;  
    Uniq(Uniq&& rhs)  
        : p(rhs.p) { rhs.p = nullptr; }  
    friend void swap(Uniq& a, Uniq& b) { std::swap(a.p, b.p); }  
    void operator=(Uniq rhs) { swap(*this, rhs); }  
    ~Uniq() { delete p; }  
};
```

# Not `[[clang::trivial_abi]]`

```
struct Uniq {  
    int *p = nullptr;  
    et cetera  
};
```

```
Uniq test_function(Uniq x) {  
    *x.p = 42;  
    return x;  
}
```

Without `[[trivial_abi]]`:

```
movq %rdi, %rax  
movq (%rsi), %rcx  
movl $42, (%rcx)  
movq %rcx, (%rdi)  
movq $0, (%rsi)  
retq
```

With `[[trivial_abi]]`:

```
movq %rdi, %rax  
movl $42, (%rdi)  
retq
```



# I have a blog post on `[[trivial_abi]]`

- It's a hack inserted directly in the Itanium C++ ABI
- It affects the calling convention for parameters and return values of type `T`
- Since the parameter is passed in a register, the caller can't access it after the call. Responsibility for destroying that parameter object is transferred to the callee.
  - MSVC's convention is also "callee-destroy." The Itanium convention is "caller-destroy" by default, but `[[trivial_abi]]` changes that.
- `[[trivial_abi]]` does not elide or eliminate the observable side effects of move-constructors and destructors!

# I have a blog post on `[[trivial_abi]]`

```
struct [[clang::trivial_abi]] A { ~A(); };  
struct B { ~B(); };  
struct [[clang::trivial_abi]] C { ~C(); };
```

```
void f(A, B, C, B) { }
```

```
void test() {  
    f(A(), B(), C(), B());  
}
```

Parameters of types A and C are callee-destroy; parameters of type B remain caller-destroy. So the params are constructed in the order A,B,C,B and destroyed in the order C,A,B,B.

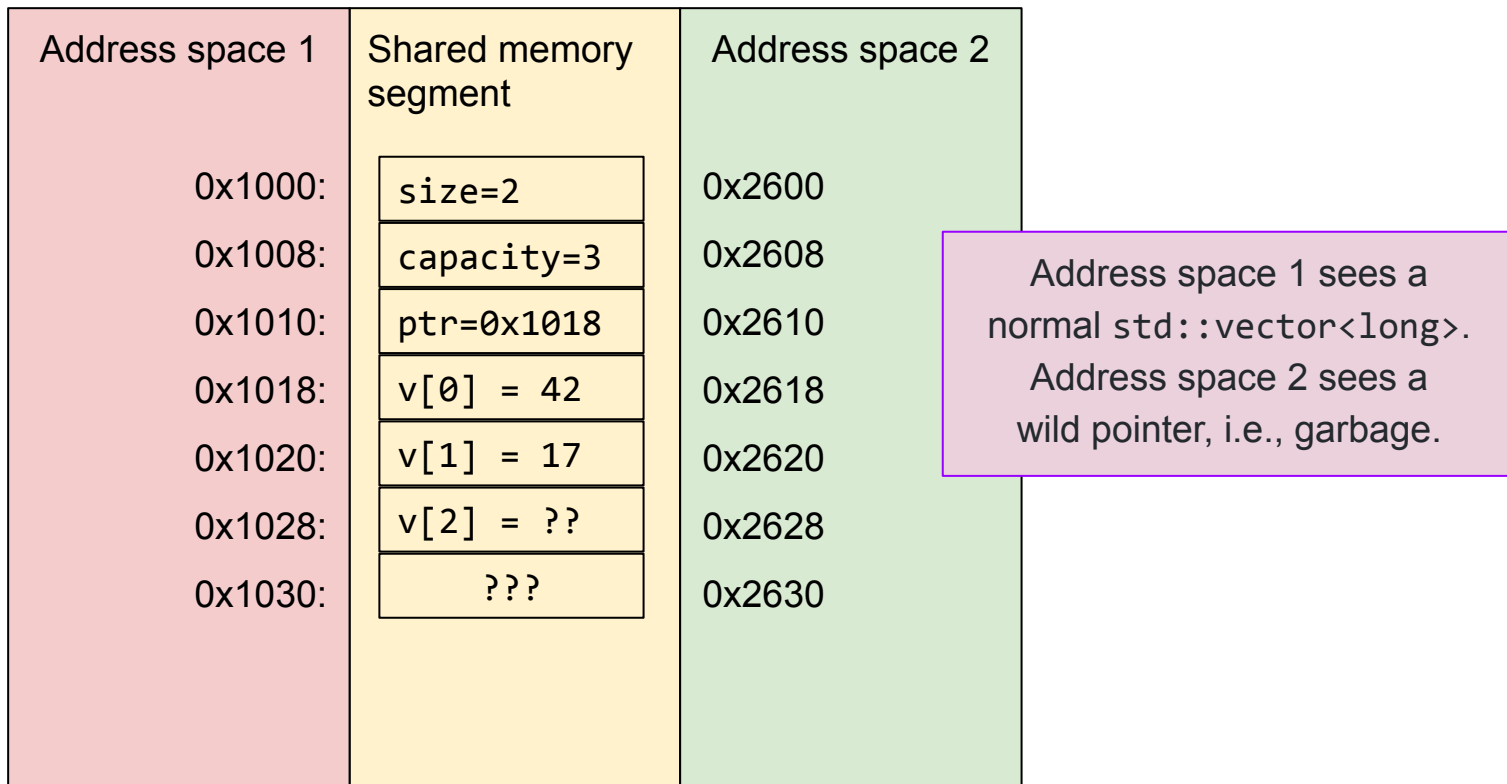
**P1144 relocation  
is not...**

# Not likely to help with “persistence”

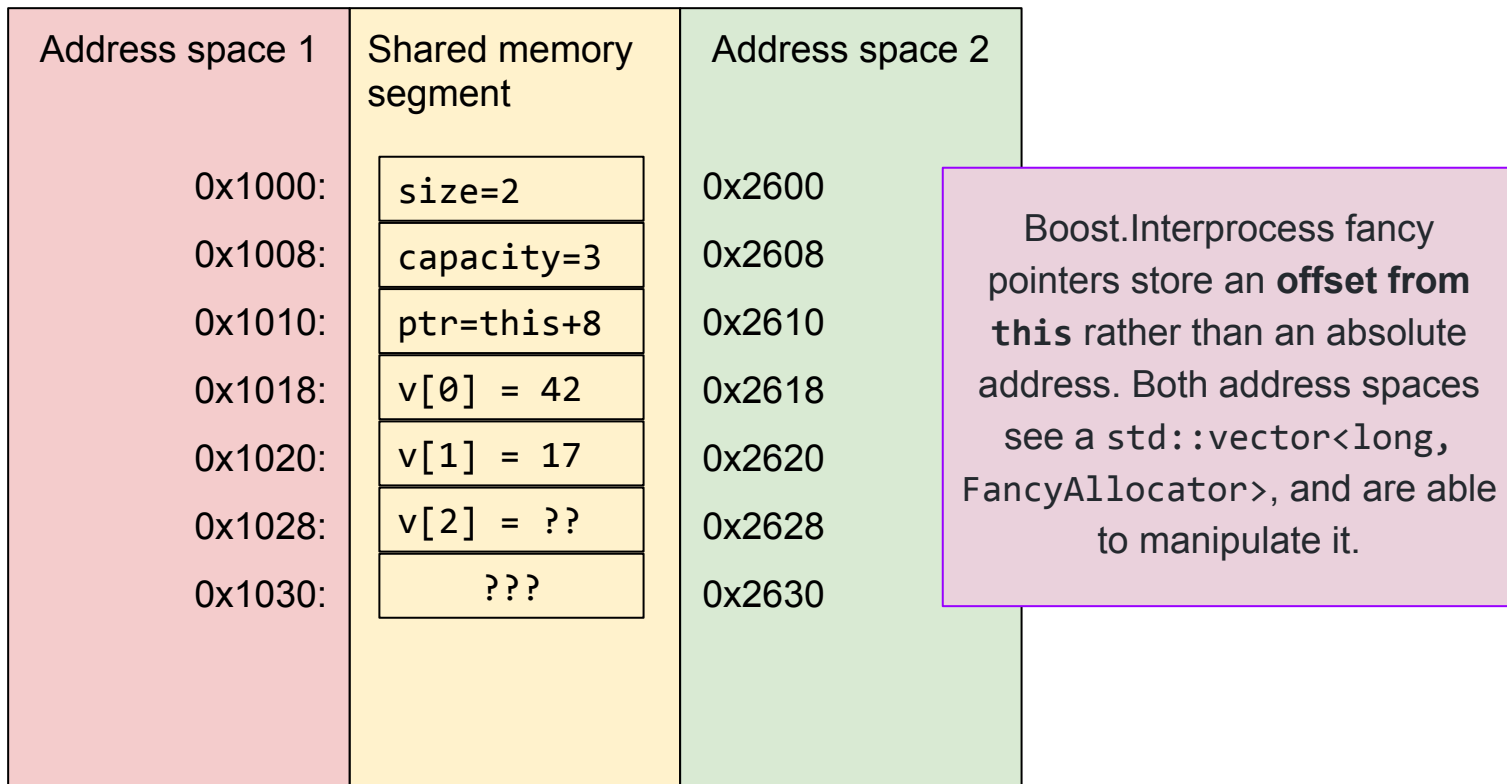
Sometimes people hear “you can relocate a whole T object with memcpy” and what they take away is “T doesn’t care what virtual address it’s at.”

Think `boost::interprocess` shared memory segments.

# Not likely to help with “persistence”

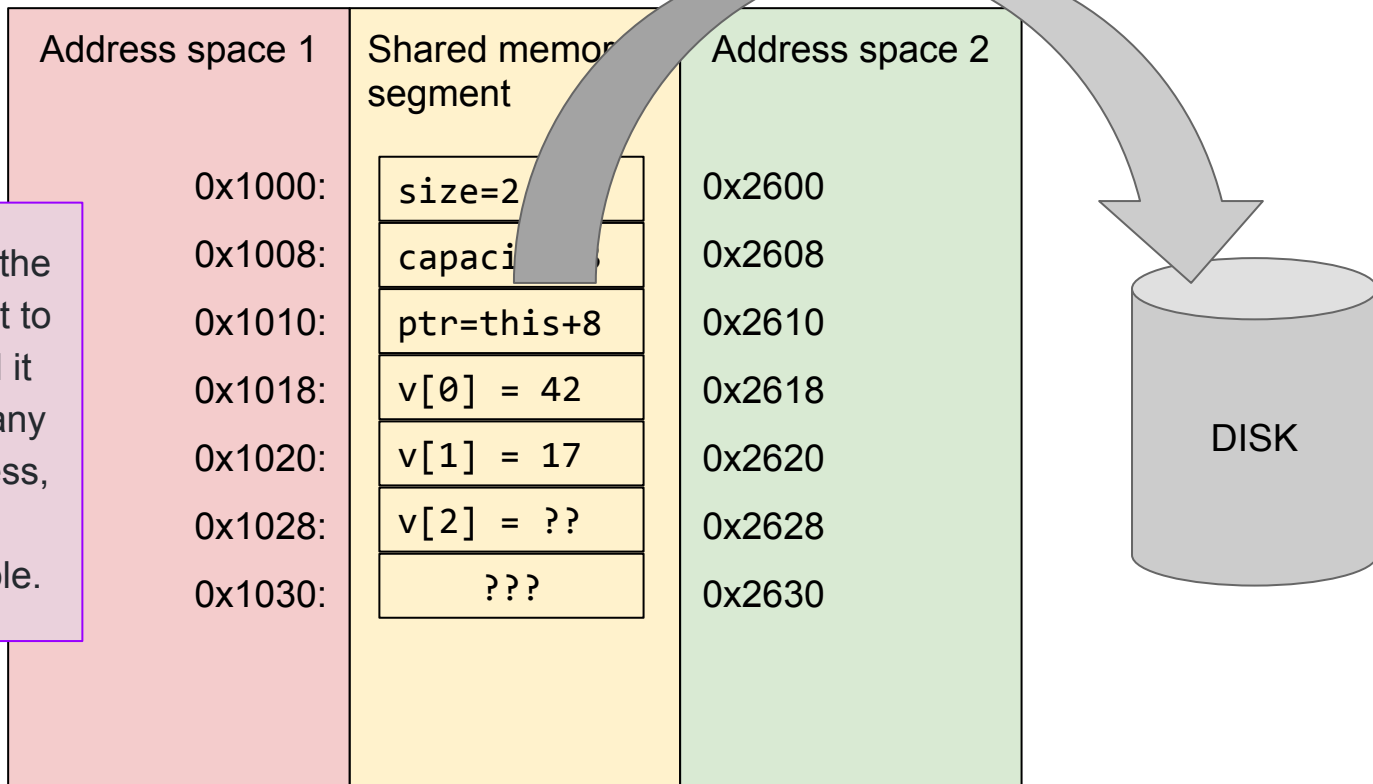


# Not likely to help with “persistence”



# Not likely to help with “persistence”

We can dump the whole segment to disk and read it back later, at any memory address, and it'll be comprehensible.

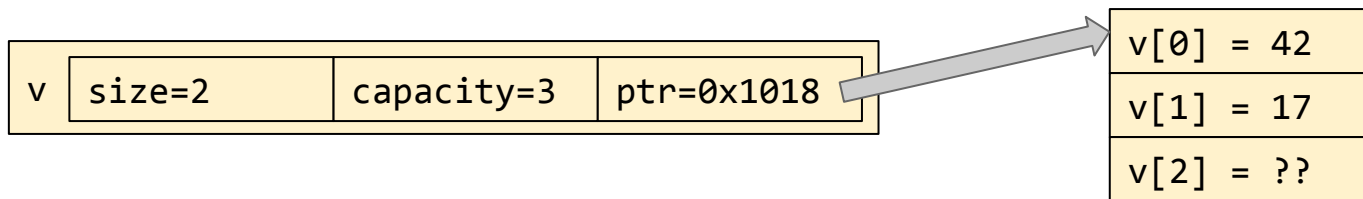


# Not likely to help with “persistence”

So `is_trivially_relocatable_v<std::vector<int, FancyAllocator>>?`

***No!***

Go back to our familiar style of diagram.





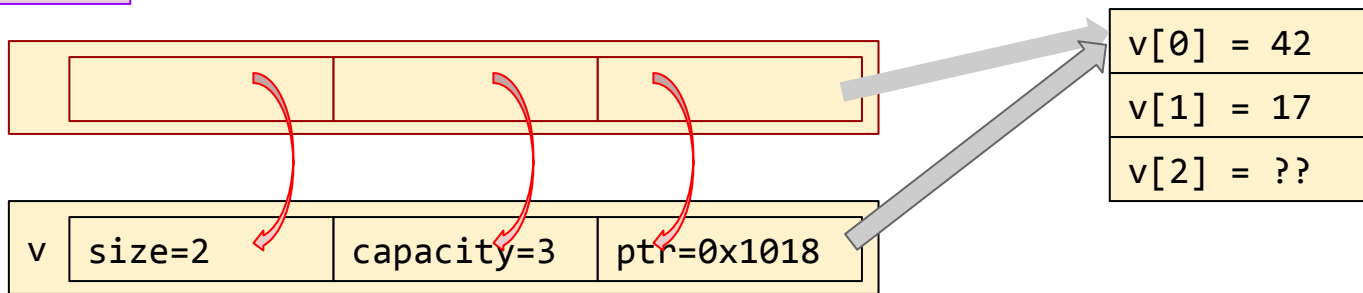
# Not likely to help with “persistence”

So `is_trivially_relocatable_v<std::vector<int, FancyAllocator>>?`

Plain vector is trivially relocatable, because plain pointers can be memcpied.

**No!**

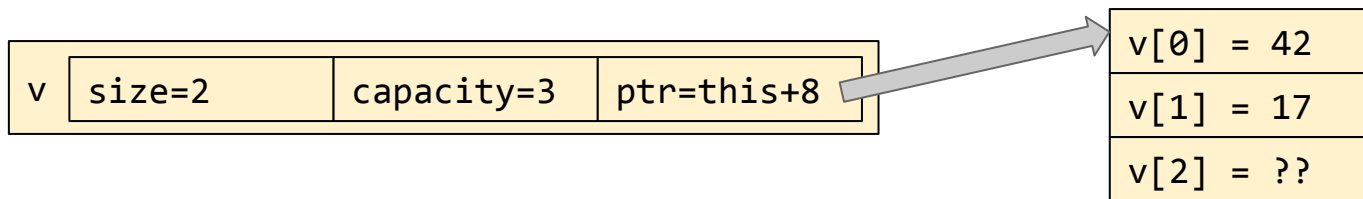
Go back to our familiar style of diagram.



# Not likely to help with “persistence”

So `is_trivially_relocatable_v<std::vector<int, FancyAllocator>>`?

***No!***

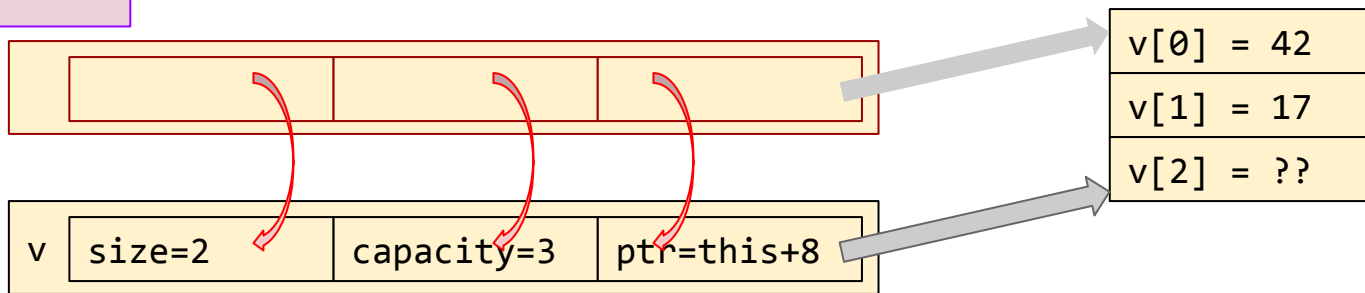


# Not likely to help with “persistence”

So `is_trivially_relocatable_v<std::vector<int, FancyAllocator>>?`

This kind of vector is **not** trivially relocatable, because `memcpy` breaks its `offset_ptrs`.

**No!**



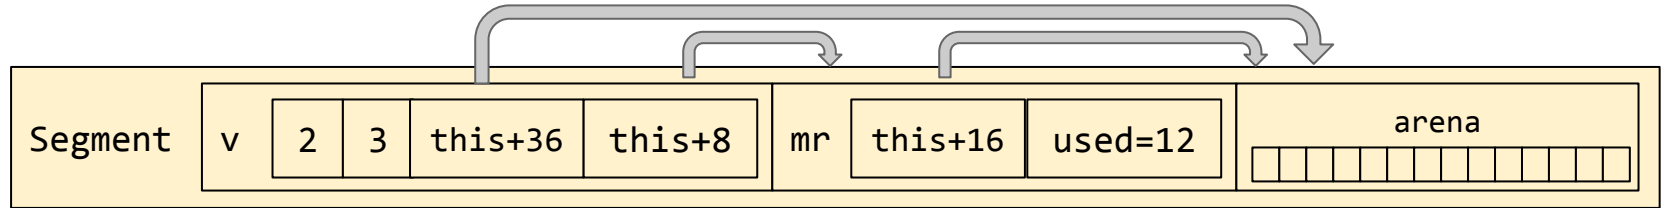
# Relocation is object-level

If you want to use trivial relocation, then you want to create some T where `is_trivially_relocatable<T>`. You must wrap the ***whole footprint*** of your relocatable arena-segment-thing into a ***single object***.

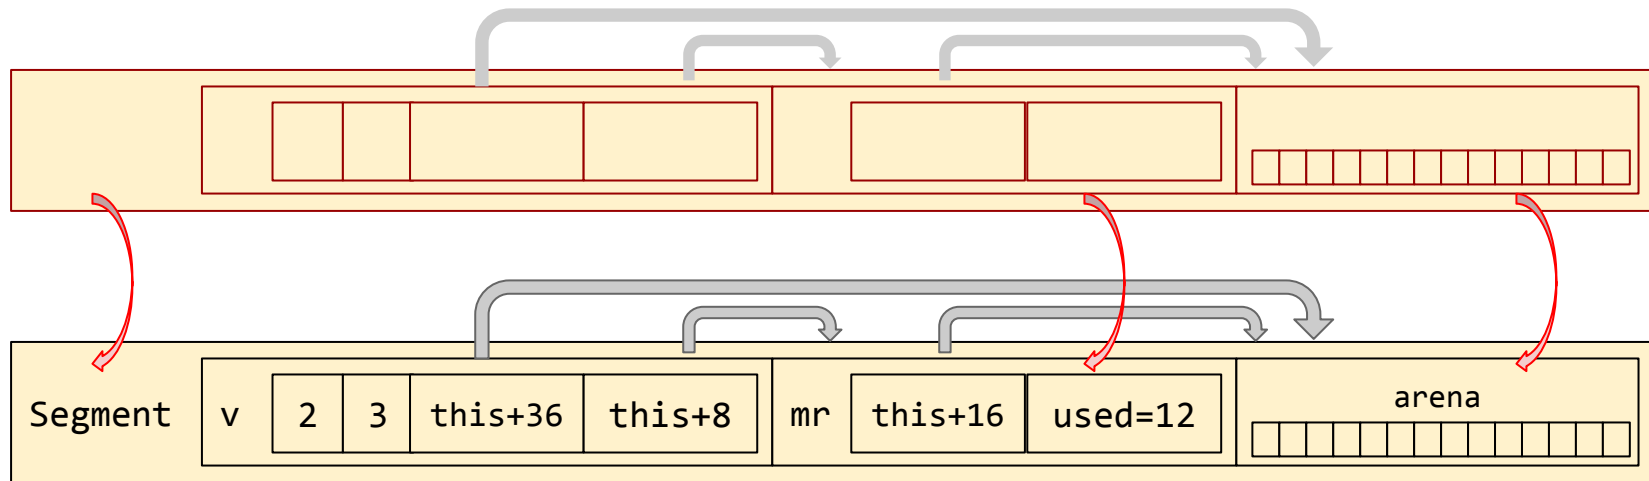
```
struct [[trivially_relocatable]] Segment {  
    std::vector<int, offset_ptr_allocator<int>> v;  
    offset_ptr_memoryresource myArena;  
    char arena[1000];  
};
```

Now you can trivially relocate a whole Segment from place to place within your address space:

# Relocation is object-level



# Relocation is object-level



Beware wrong warrants: All the stuff in the arena must **also** be trivially relocatable. (In this case, it's just longs. But what if it were `std::regexes`?)

# Relocation is object-level

That was just a proof of concept. I expect `is_trivially_relocatable` will **not** be practically useful to people who do this kind of thing.

In practice, you **never** `memcpy` a giant Segment object, because:

- it's being shared with other threads/processes and you can't move it
- `memcpy`ing a giant object is super slow

P1144's trivial relocatability is useful for **small** objects that get shuffled around in memory a lot. It's not for big objects that stay in one place.

Plus, you had to write an explicit warrant for Segment, so I bet you got it wrong. P1144 is designed for the 99% case: Rule of Zero, no warrant, free speedup.

# Open questions

- Is this what attributes are for? If not attributes, then what?
- In cases like `Segment`, `[[maybe_trivially_relocatable]]` does not suffice — we need `[[trivially_relocatable]]`. Should we provide both, anyway? Is `[[maybe_trivially_relocatable]]` “safer”?
- Technically, `swap(T& a, T& b)` cannot use “memswap” if either `a` or `b` might be a base of a non-POD class type, because there might be stuff stored in their tail padding. Every vendor gets this wrong today in `std::copy`. To what extent do we care? and if we care, to what extent does this nerf my advertised optimization of swap-based algorithms?



# Open questions

- P1144 says that “relocation” is just “move + destroy.” But our `vector::reserve` optimization is actually replacing “copy + destroy” in some cases, e.g. `libstdc++`’s motivating `vector<deque<int>>` example.
- Our `vector::insert` and `swap` optimizations are actually replacing “move-assign-from + move-assign-to” in some cases.

To what extent can the library assume that a trivially relocatable type behaves “normally”? Should P1144 consider `T` non-trivially-relocatable by default if it has *any* special members?

P1144R3 actually does make this change.

# **End of Part 2**

## **Questions?**

### **Bonus slides**

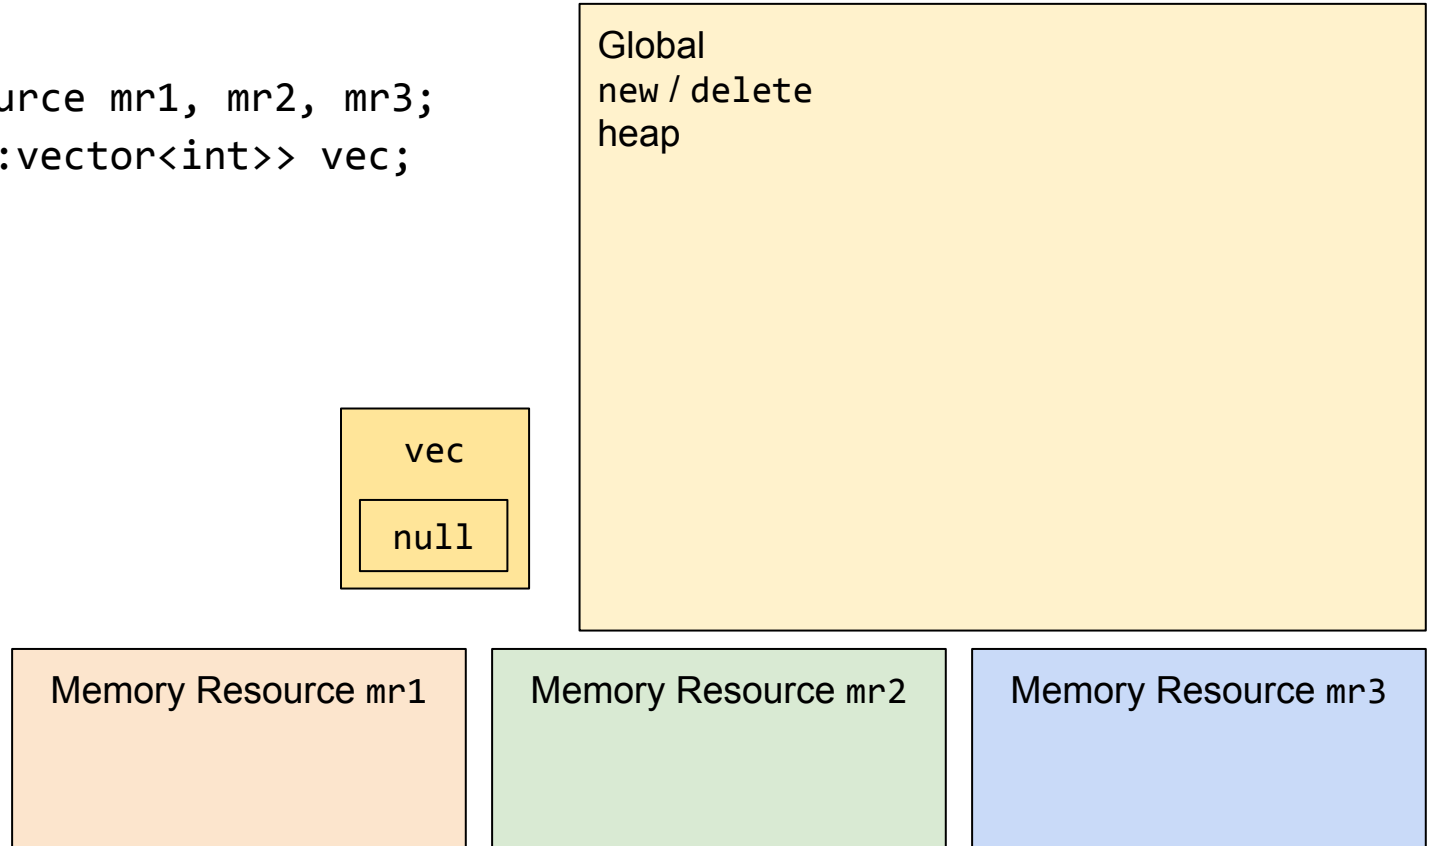
Is `pmr::vector` trivially relocatable?

# Erasing from `vector<pmr::vector<int>>`

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
vec.erase(vec.begin());
```

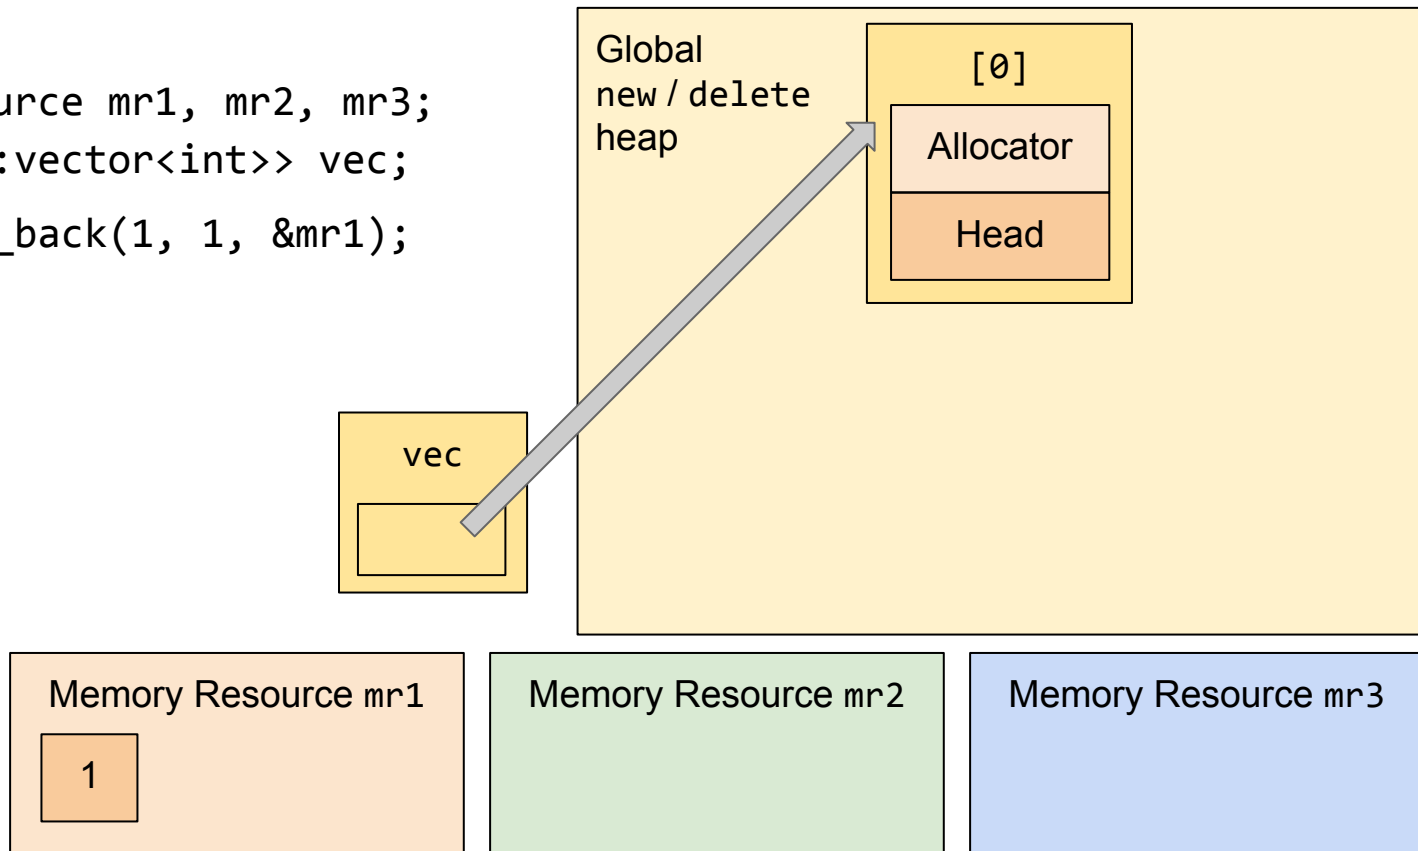
# Erasing from `vector<pmr::vector<int>>`

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;
```



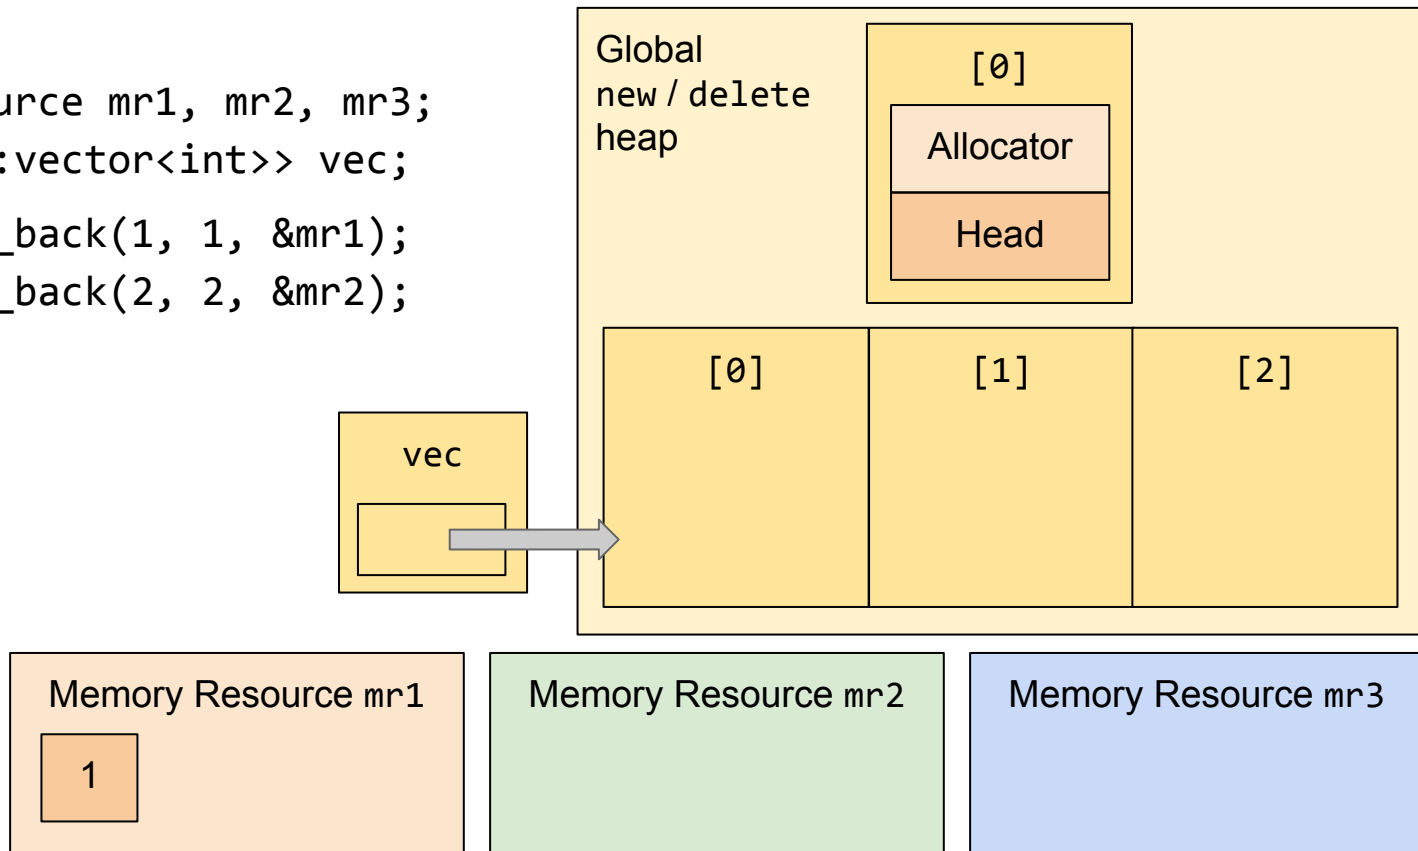
# Erasing from `vector<pmr::vector<int>>`

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
vec.emplace_back(1, 1, &mr1);
```



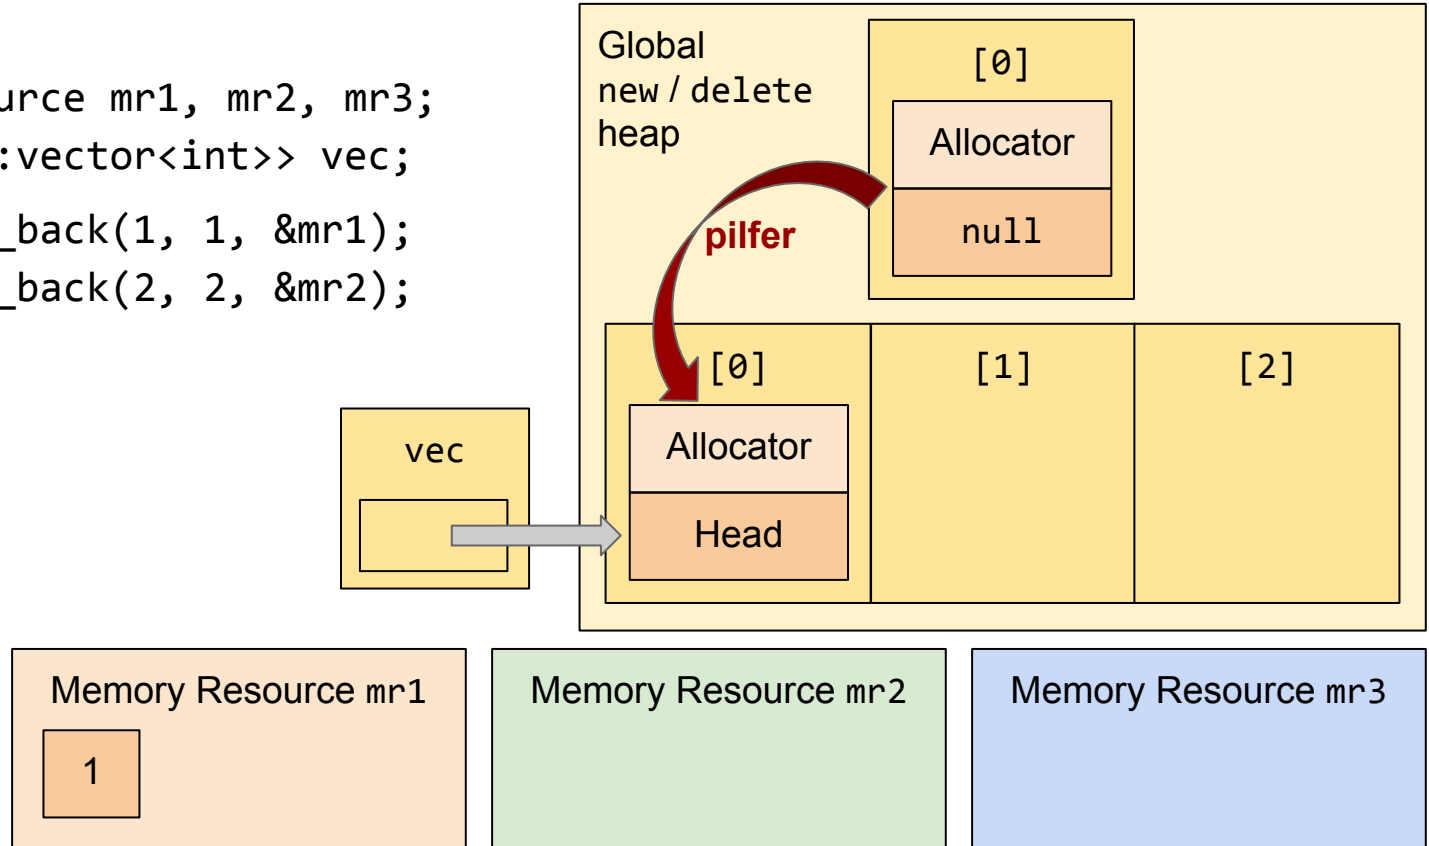
# Erasing from `vector<pmr::vector<int>>`

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);
```



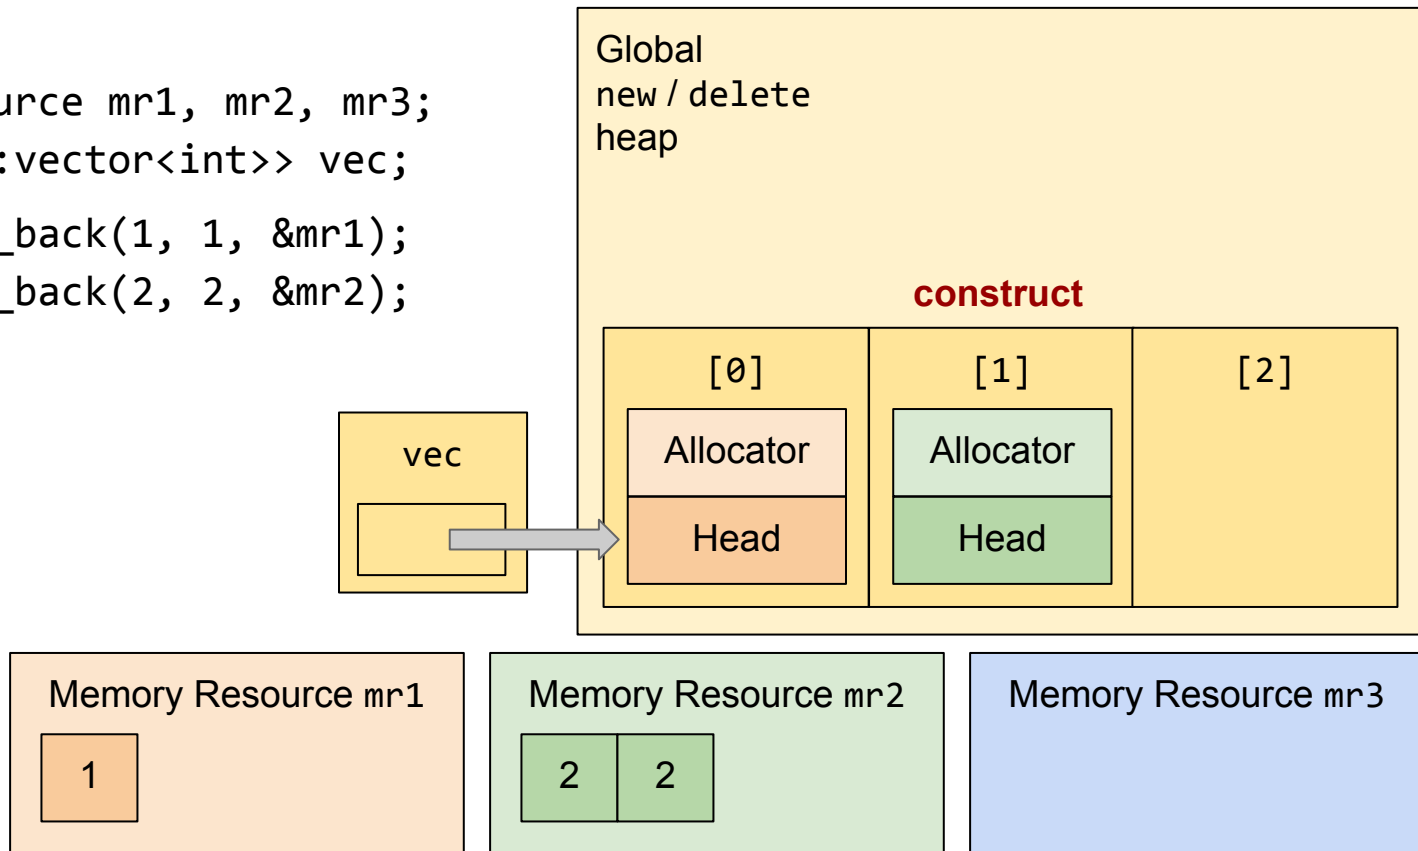
# Erasing from `vector<pmr::vector<int>>`

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);
```



# Erasing from `vector<pmr::vector<int>>`

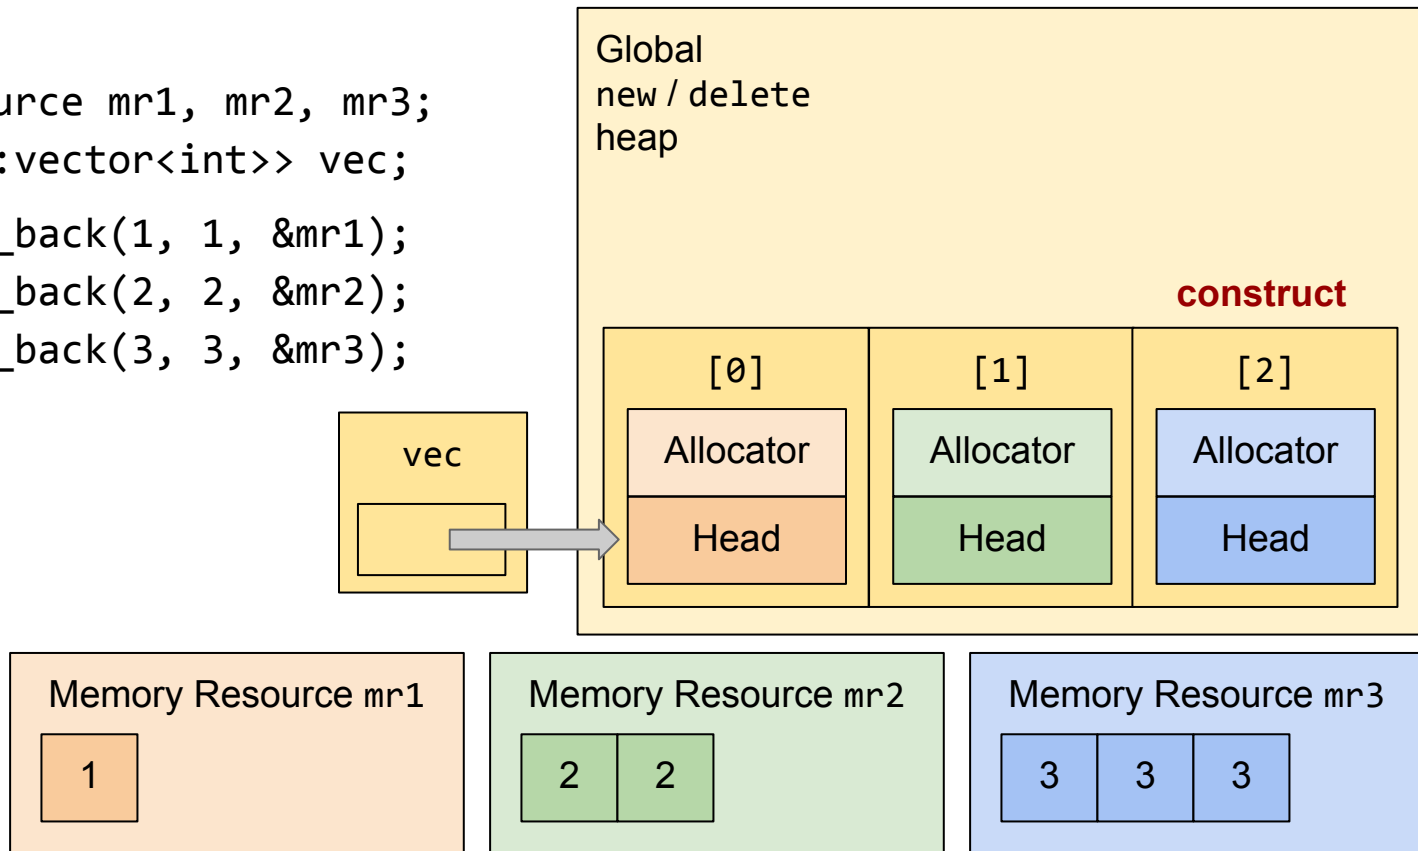
```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);
```





# Erasing from `vector<pmr::vector<int>>`

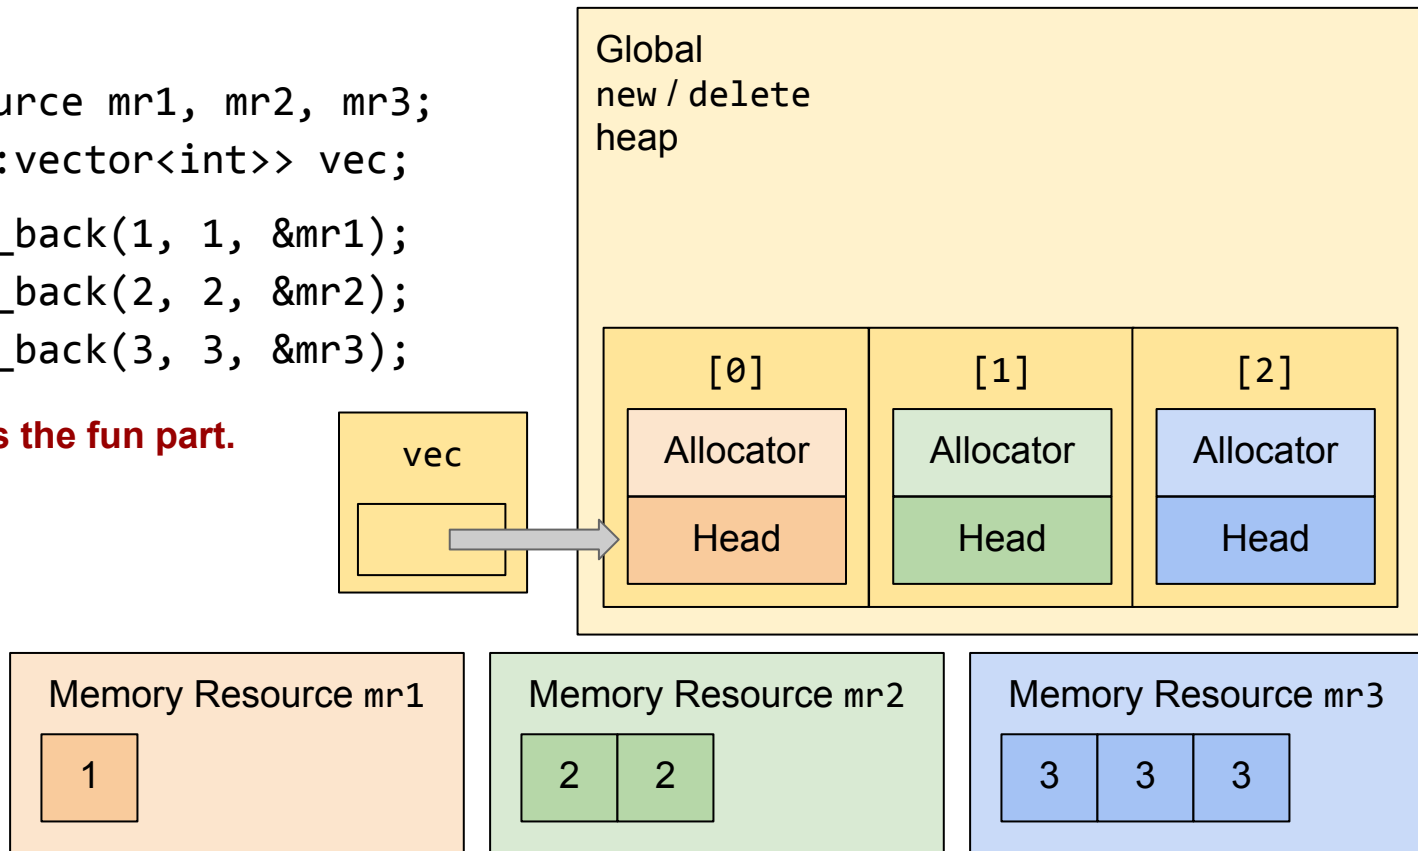
```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);
```



# Erasing from `vector<pmr::vector<int>>`

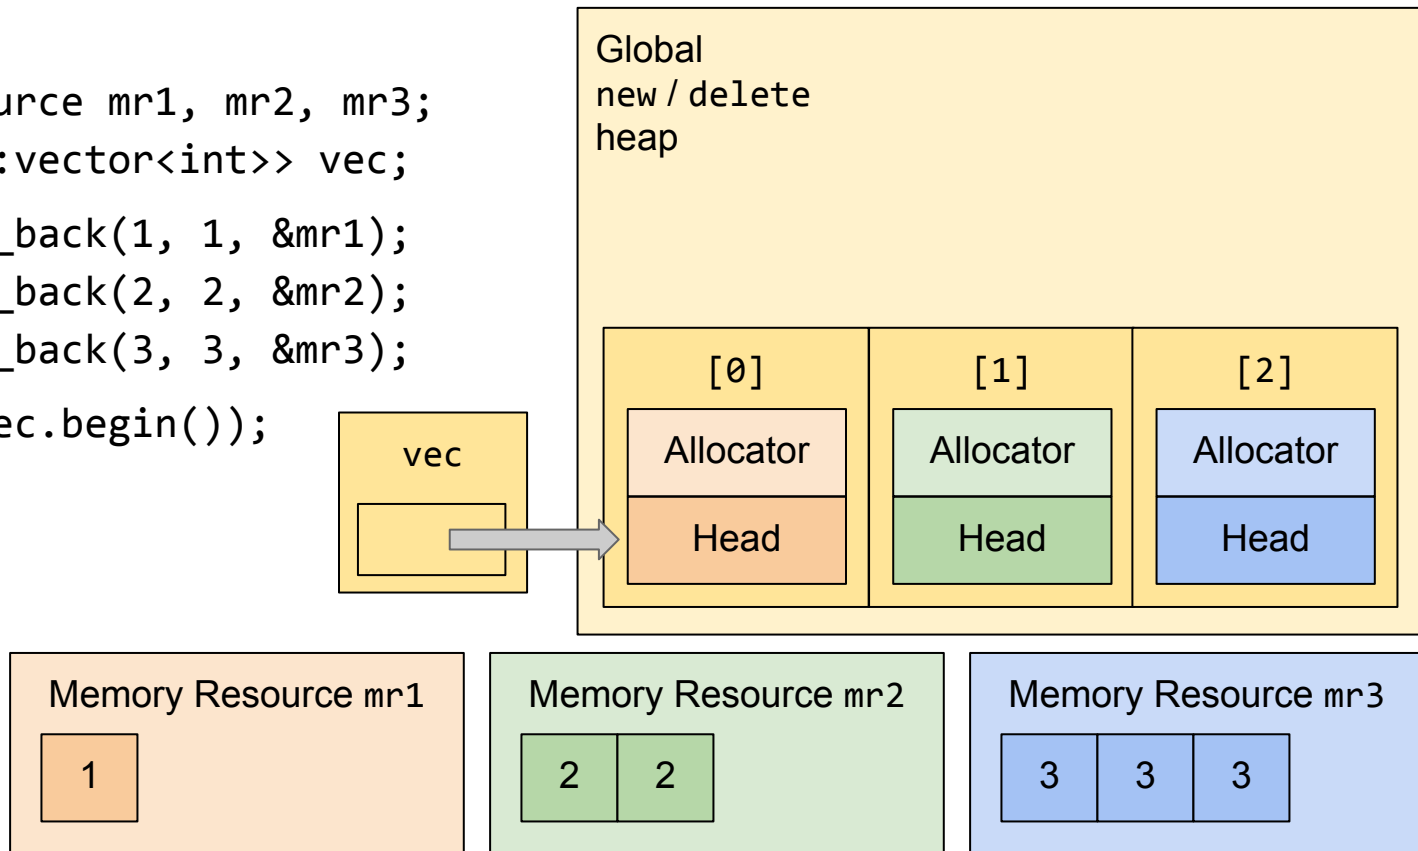
```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);
```

Now here comes the fun part.



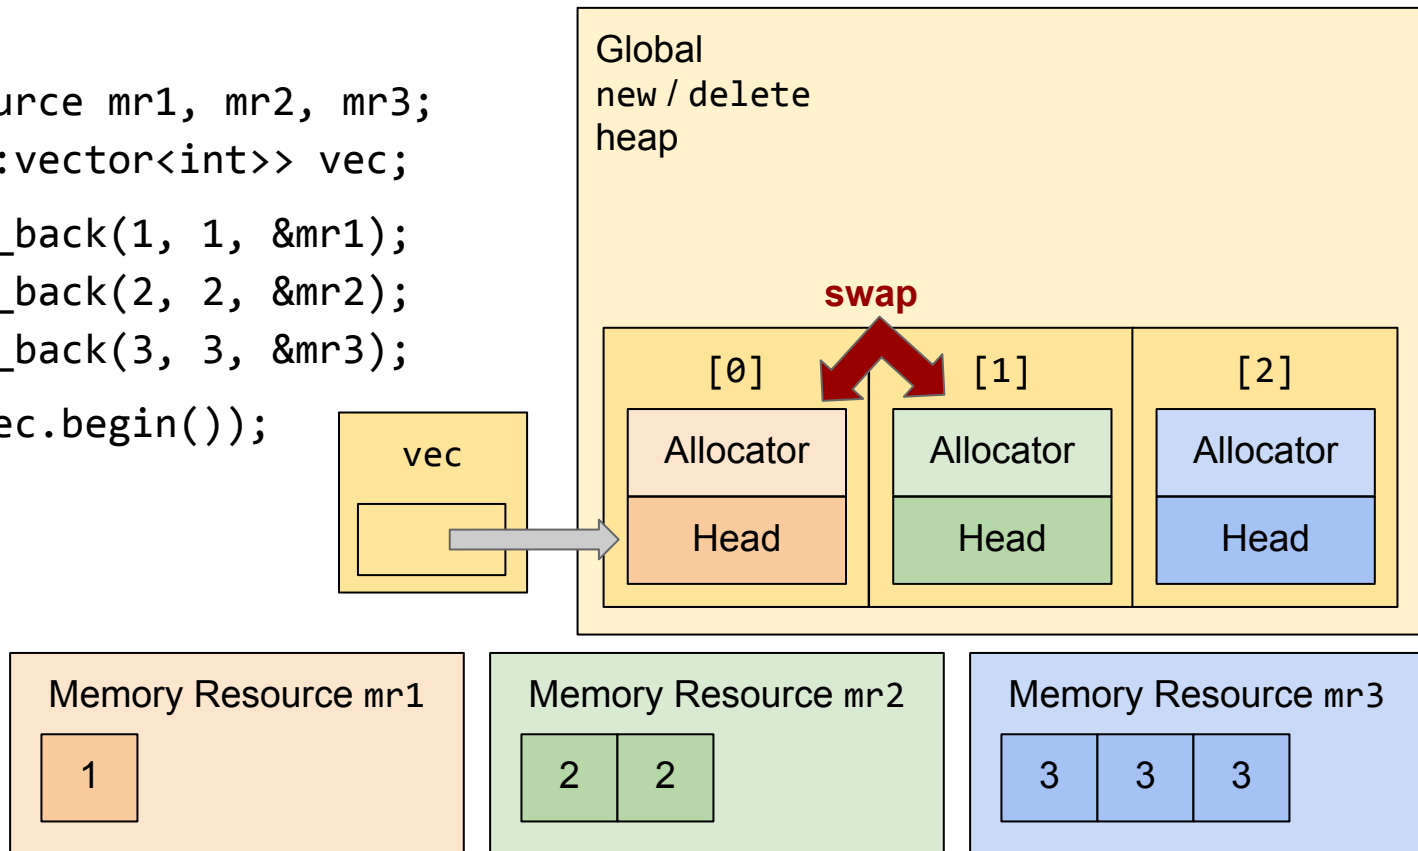
# Erasing from `vector<pmr::vector<int>>`

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
  
vec.erase(vec.begin());
```



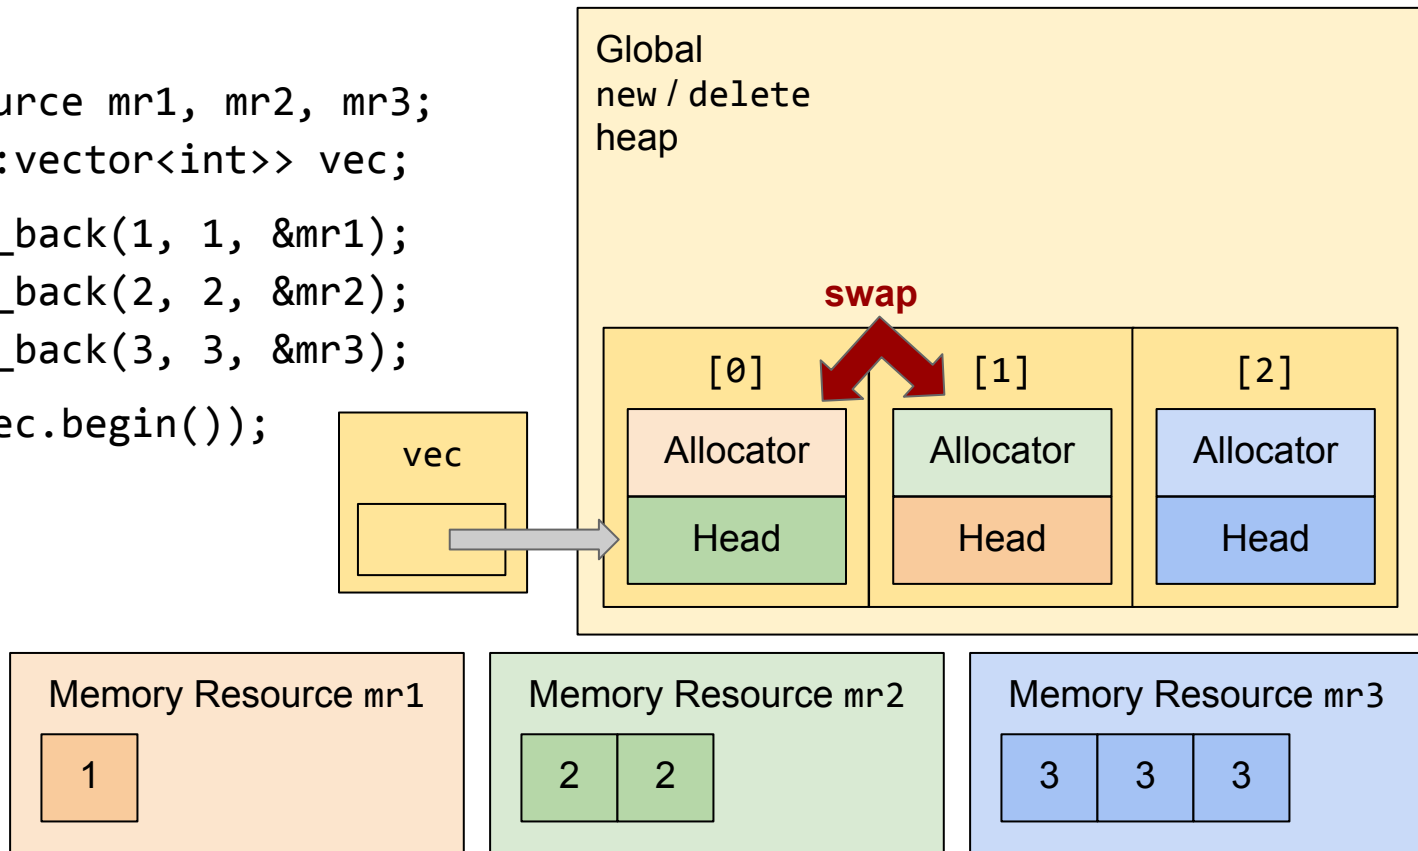
# If we implement erase via std::swap

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
  
vec.erase(vec.begin());
```



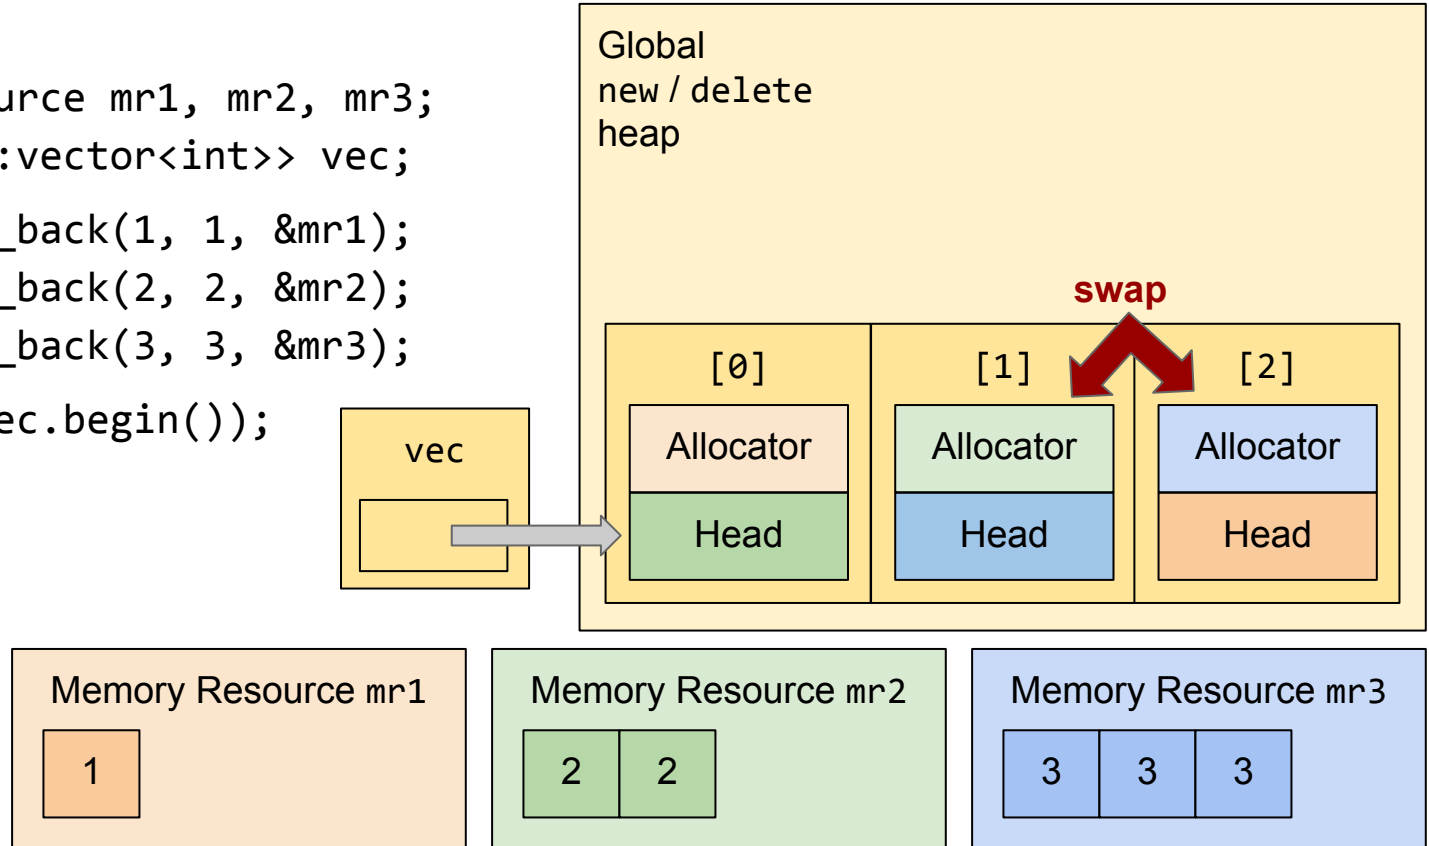
# If we implement erase via std::swap

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
  
vec.erase(vec.begin());
```



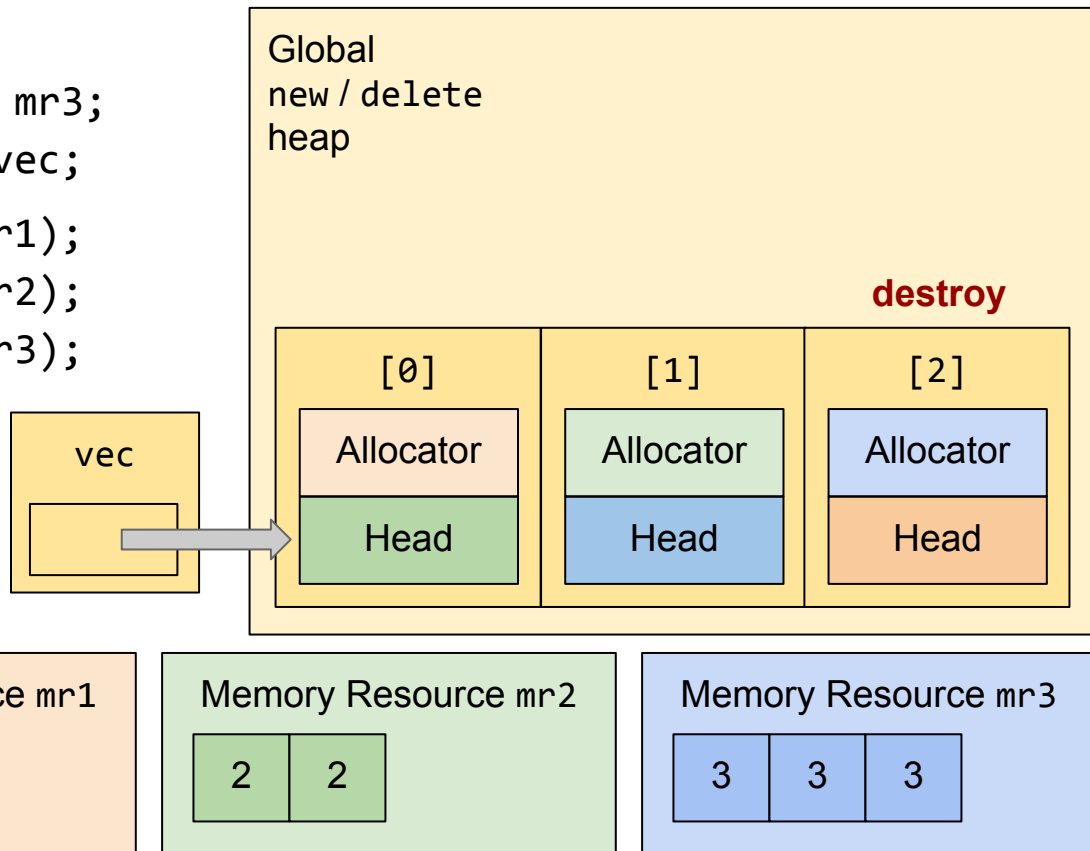
# If we implement erase via std::swap

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
  
vec.erase(vec.begin());
```



# If we implement erase via std::swap

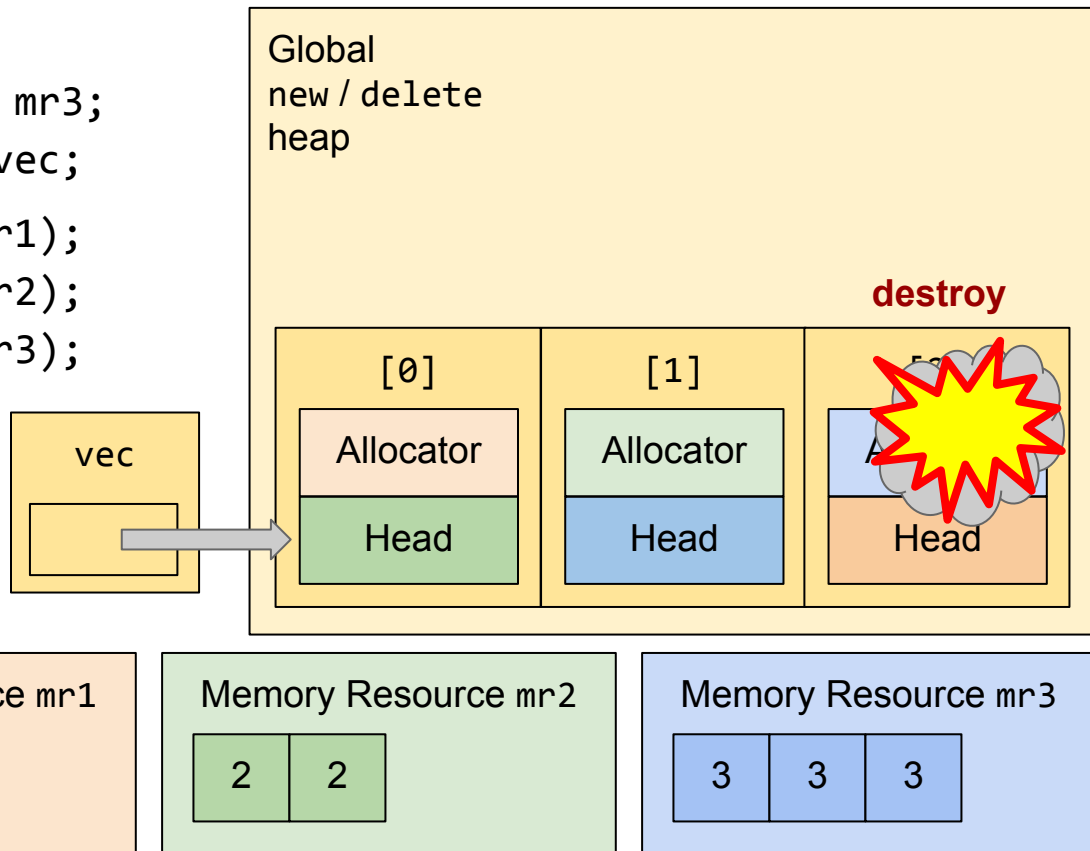
```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
vec.erase(vec.begin());
```



# If we implement erase via std::swap

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
  
vec.erase(vec.begin());
```

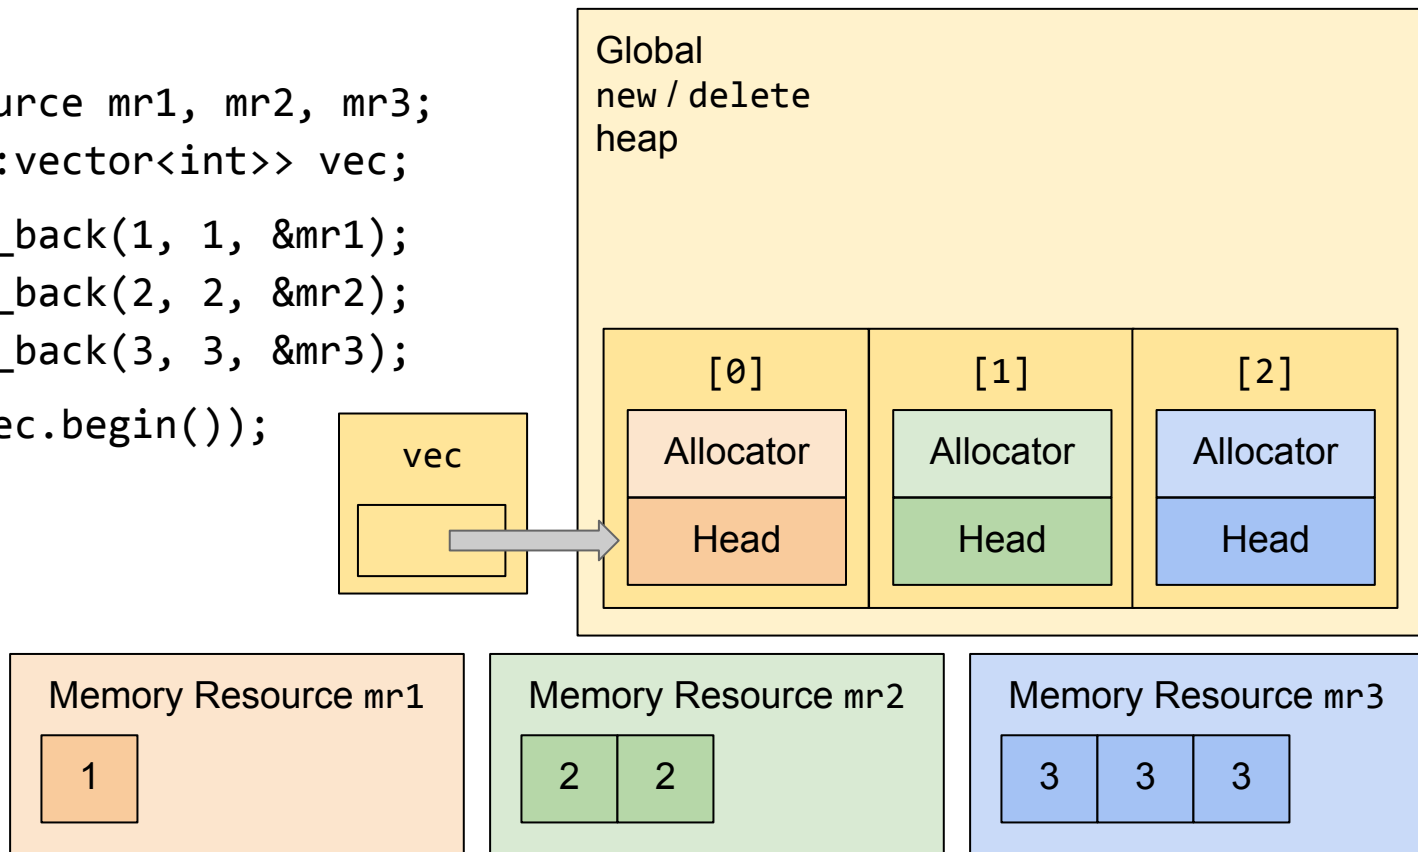
**Undefined behavior, because swapping pmr::vectors with unequal allocators is UB.**





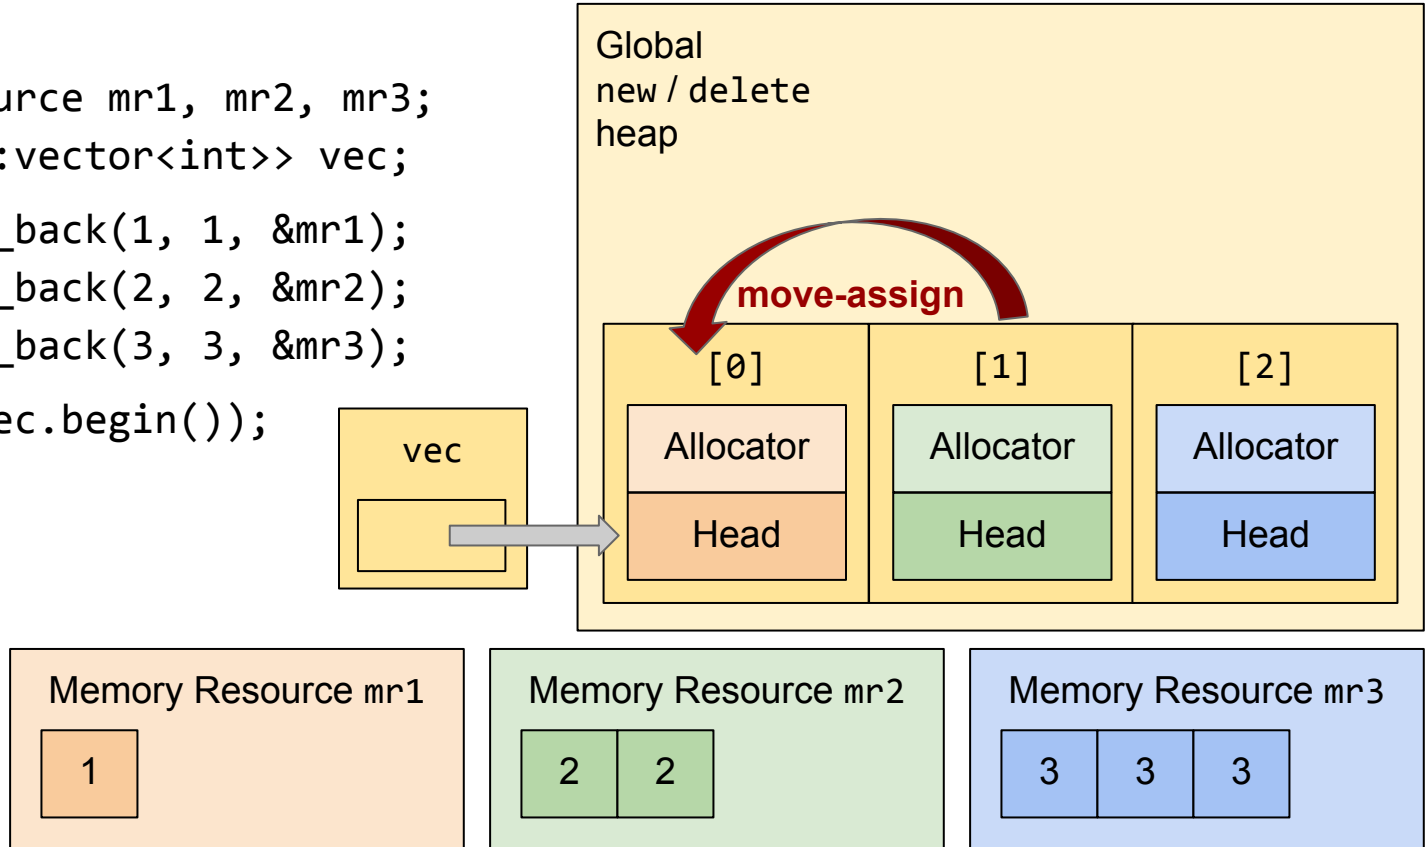
# If we implement erase via move-assign

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
  
vec.erase(vec.begin());
```



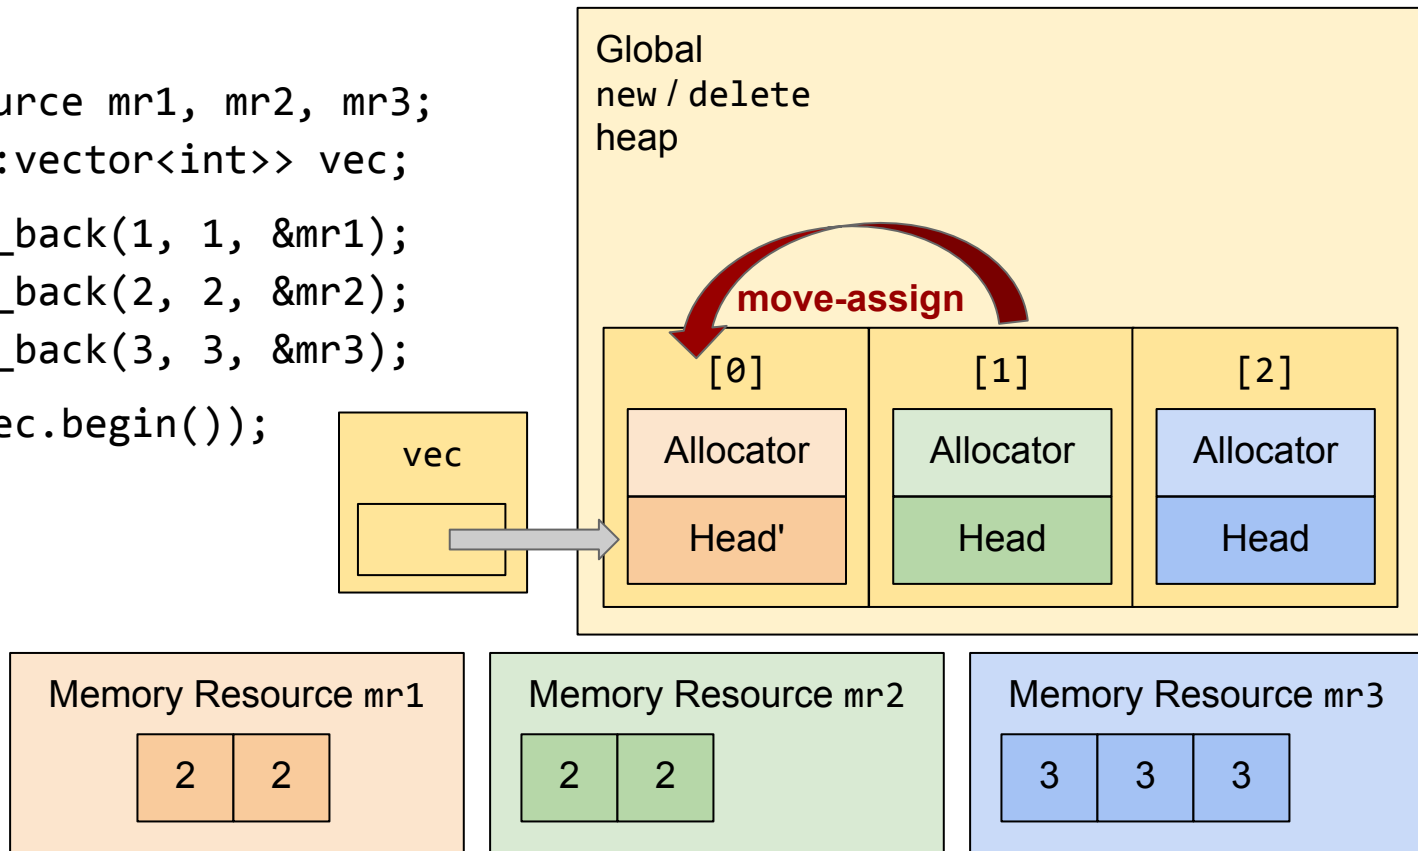
# If we implement erase via move-assign

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
  
vec.erase(vec.begin());
```



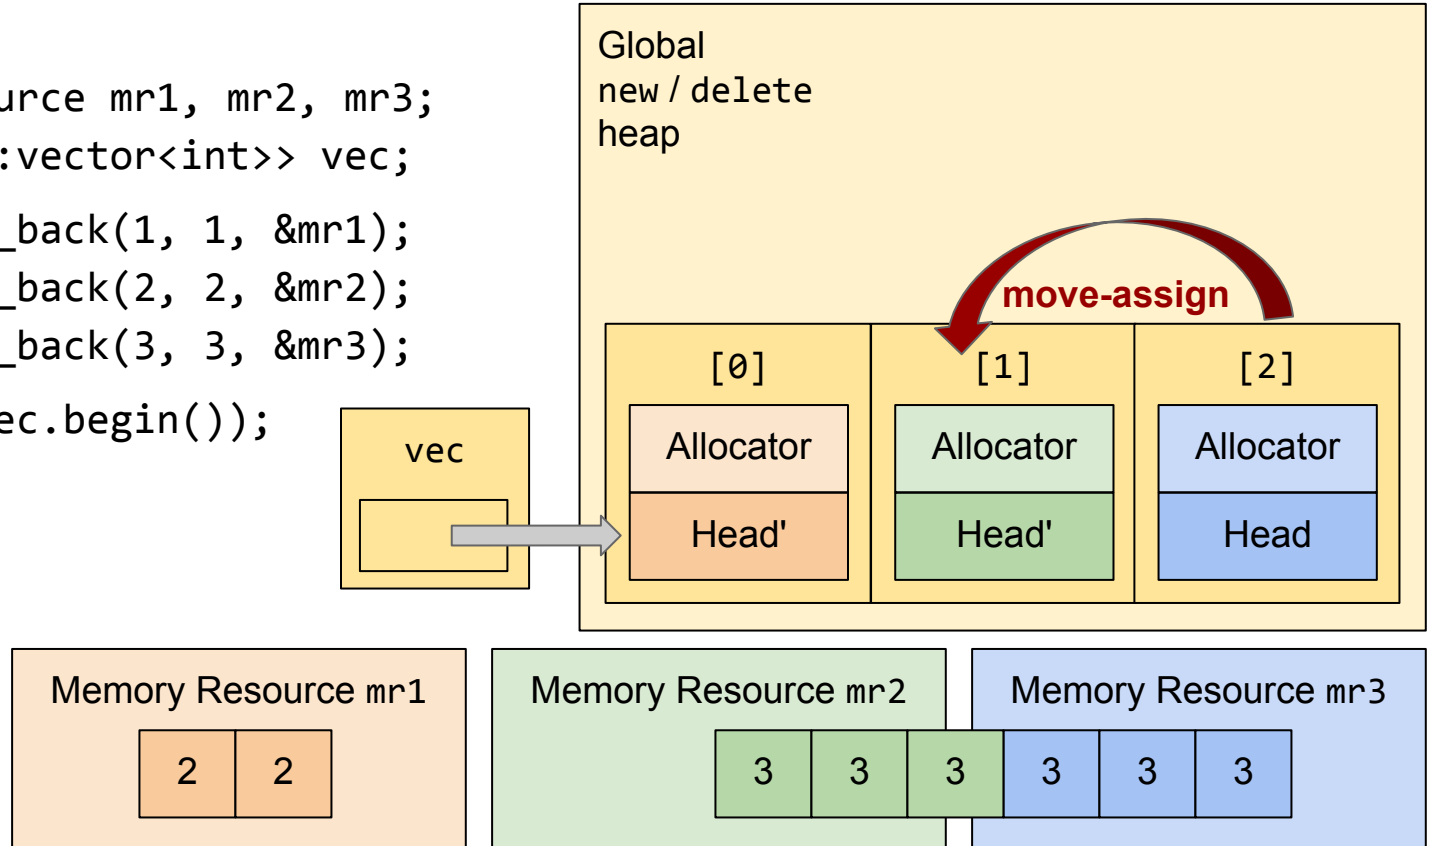
# If we implement erase via move-assign

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
  
vec.erase(vec.begin());
```



# If we implement erase via move-assign

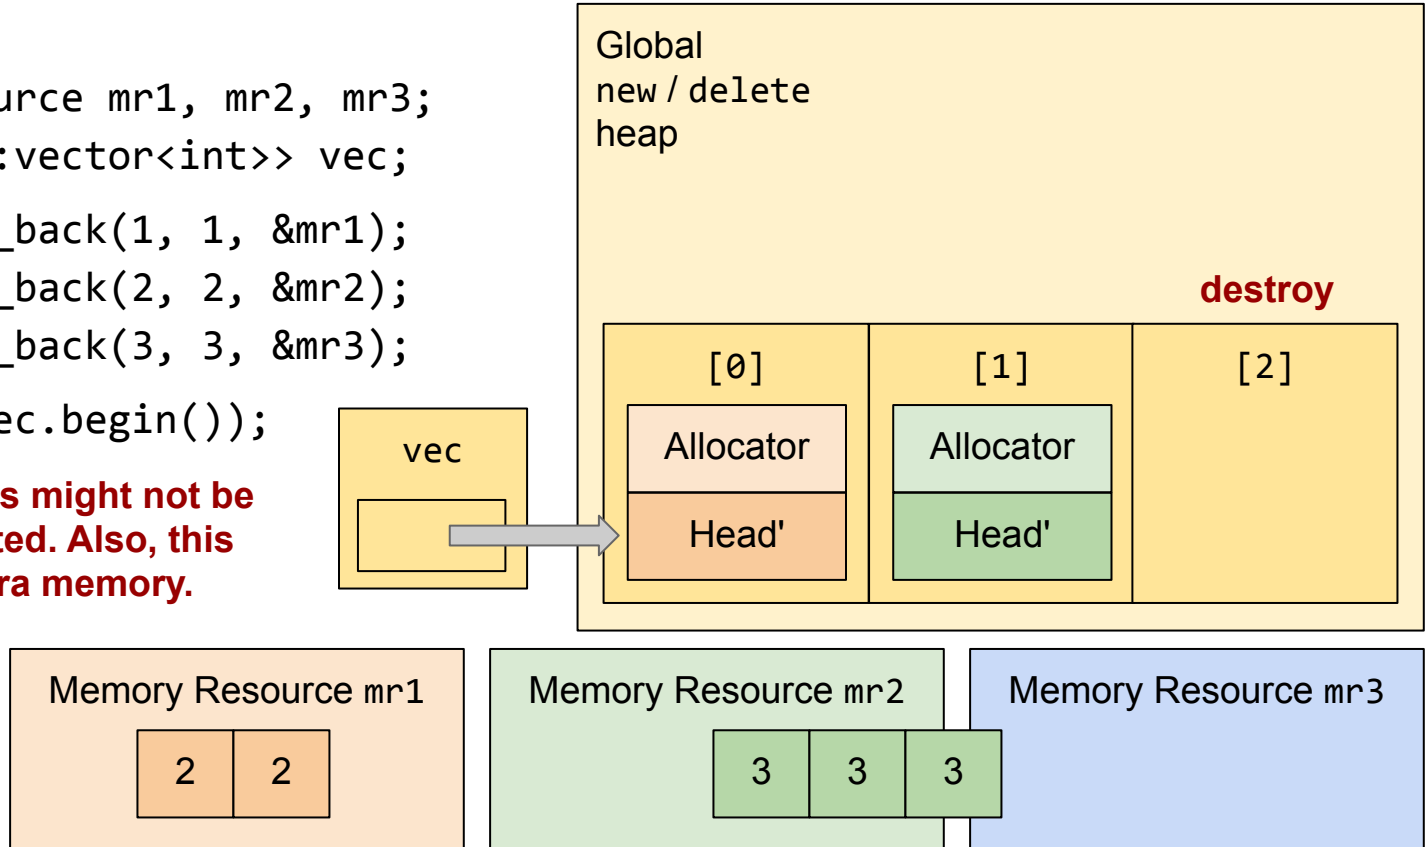
```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
  
vec.erase(vec.begin());
```



# If we implement erase via move-assign

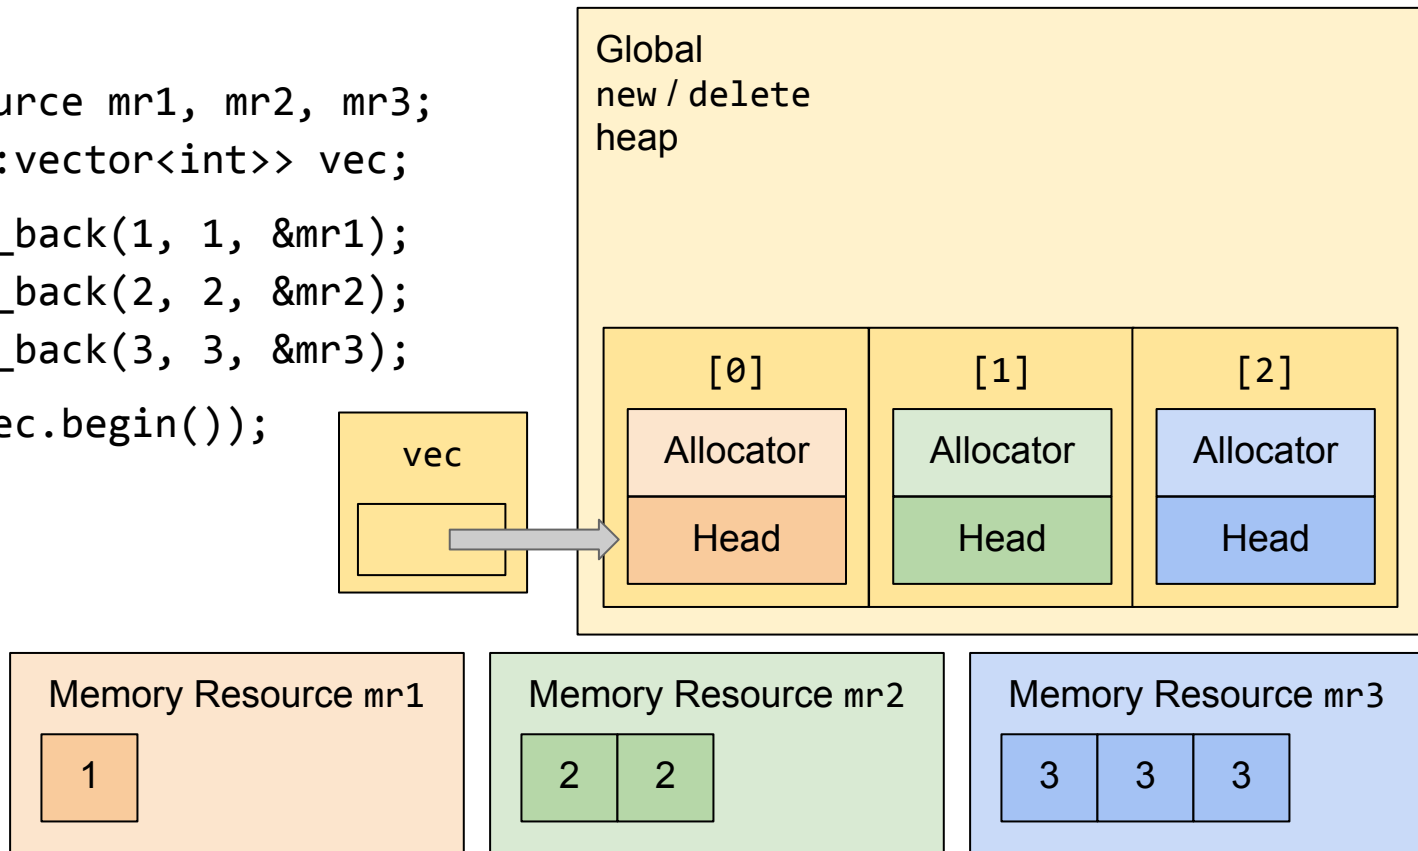
```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
  
vec.erase(vec.begin());
```

**Success, but this might not be what you expected. Also, this uses a lot of extra memory.**



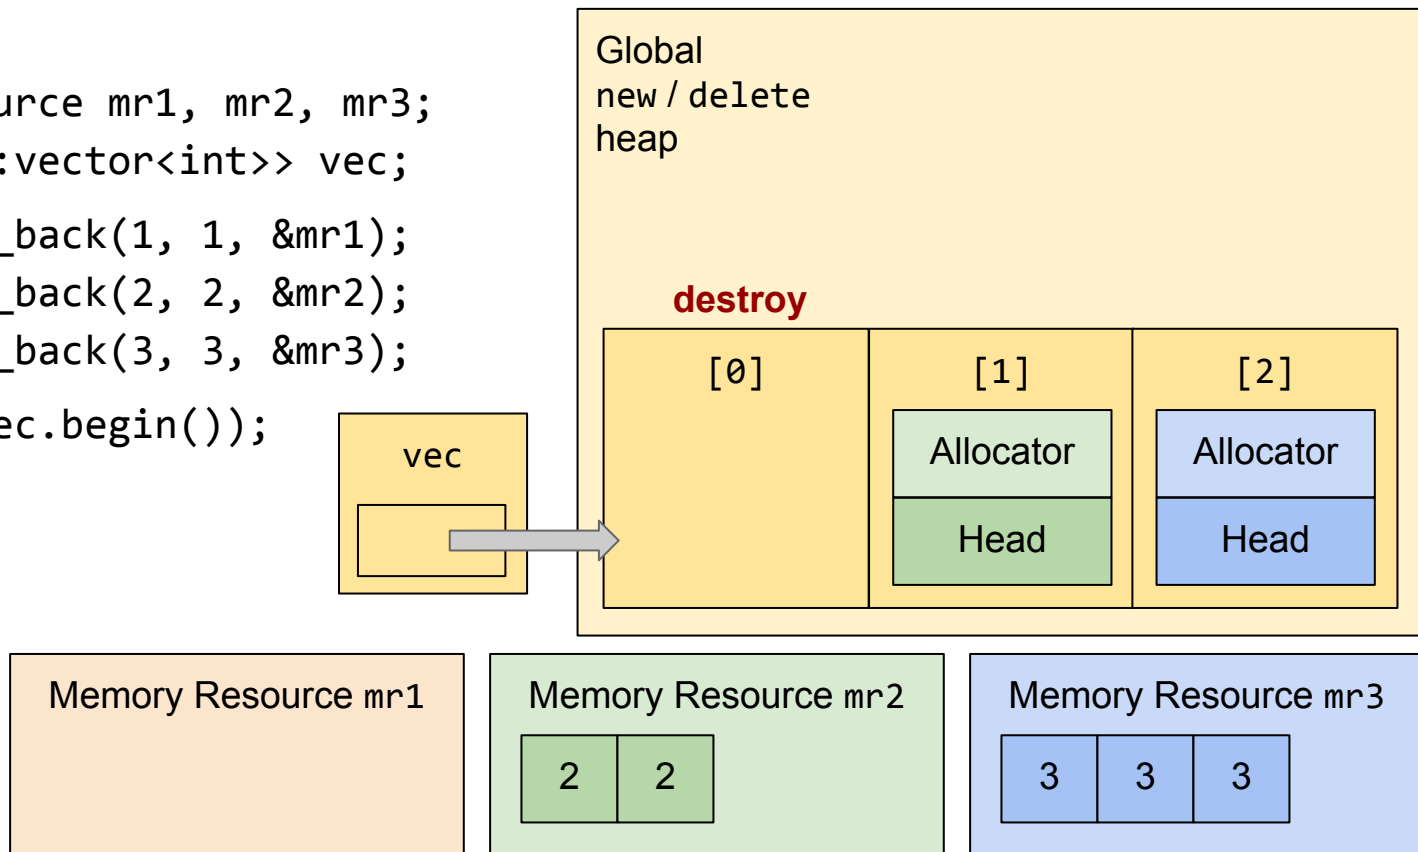
# If we implement erase via move+destroy

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
vec.erase(vec.begin());
```



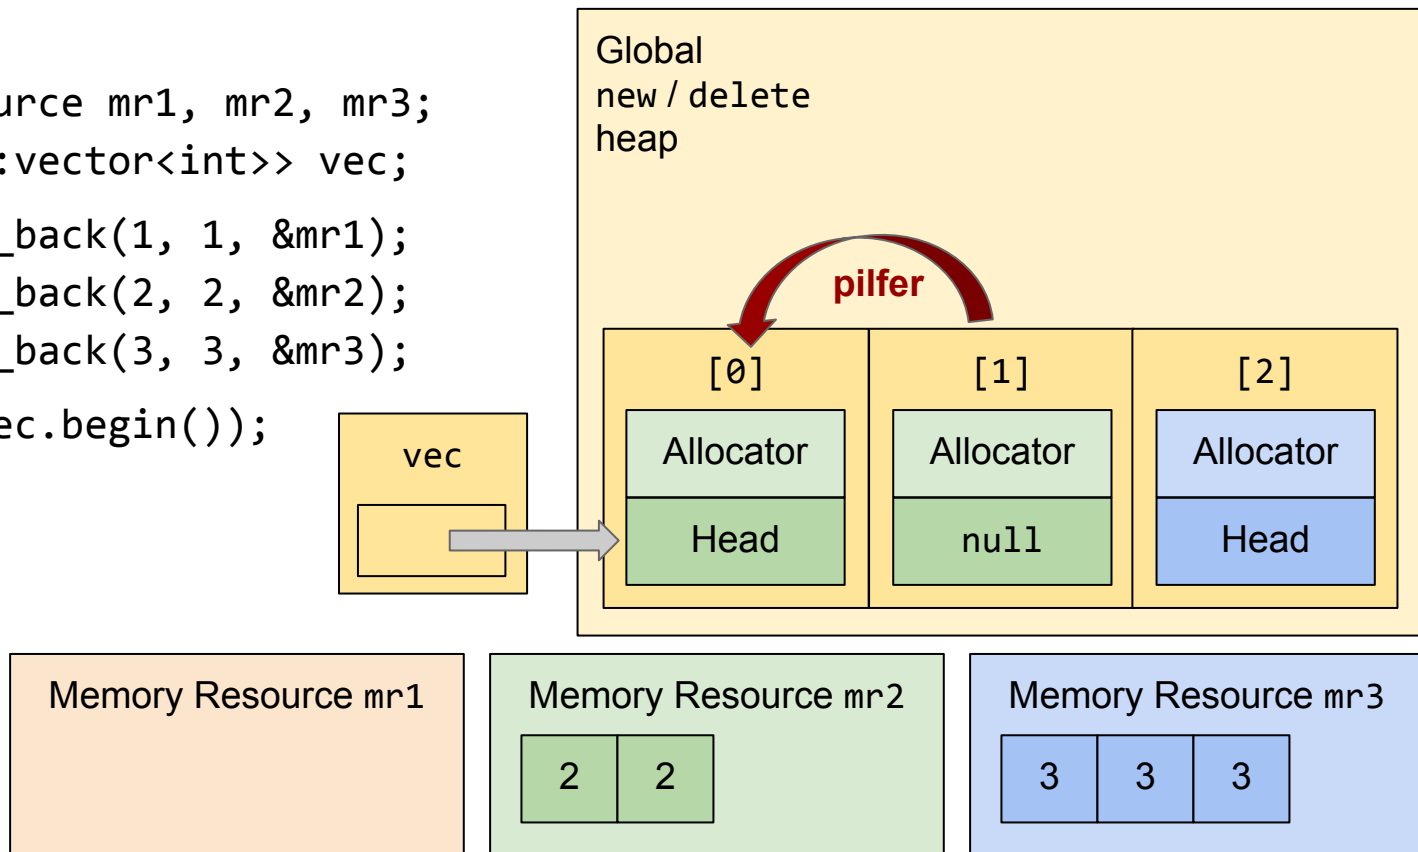
# If we implement erase via move+destroy

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
vec.erase(vec.begin());
```



# If we implement erase via move+destroy

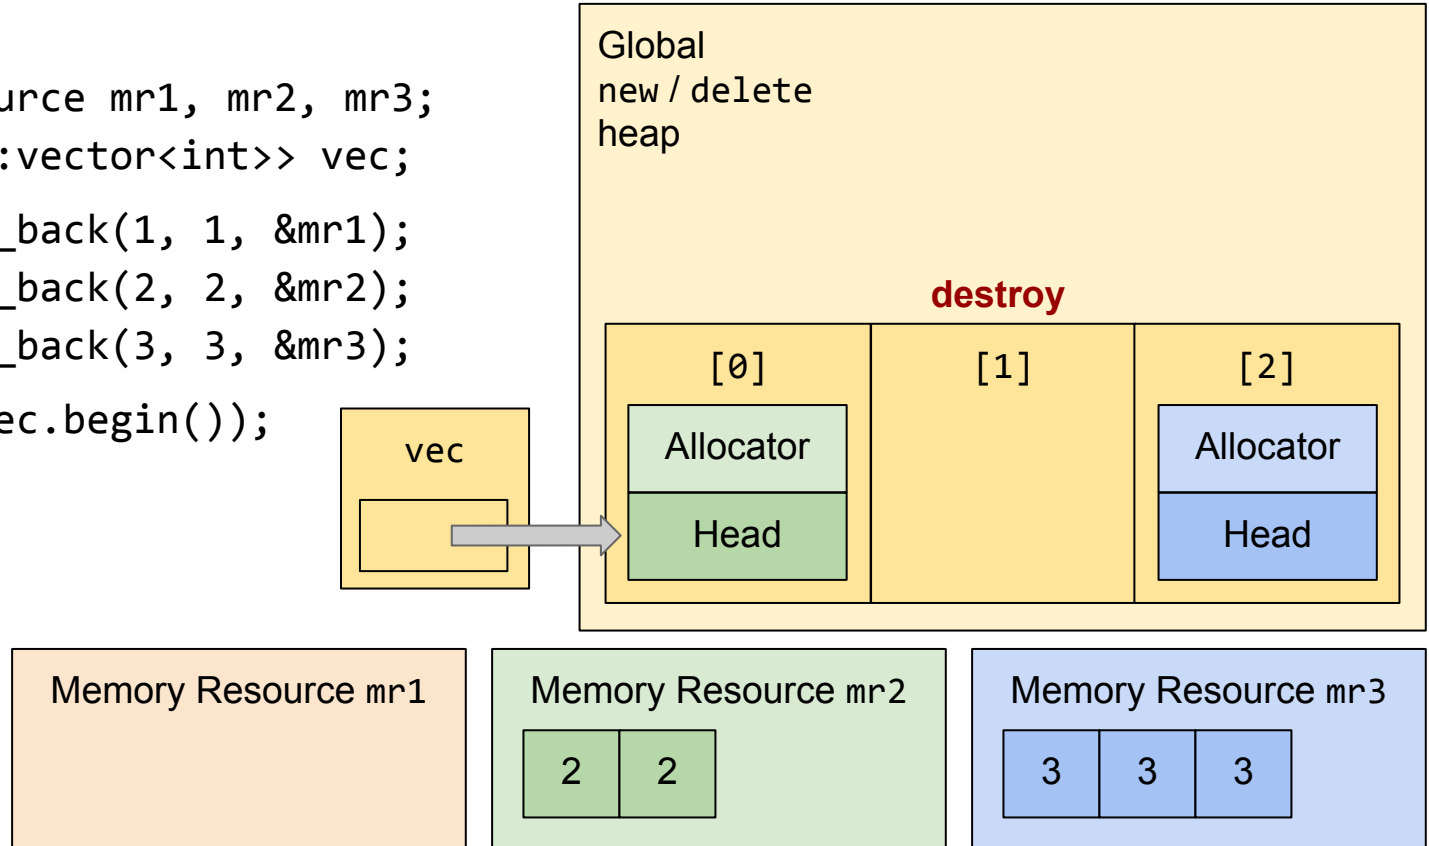
```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
  
vec.erase(vec.begin());
```





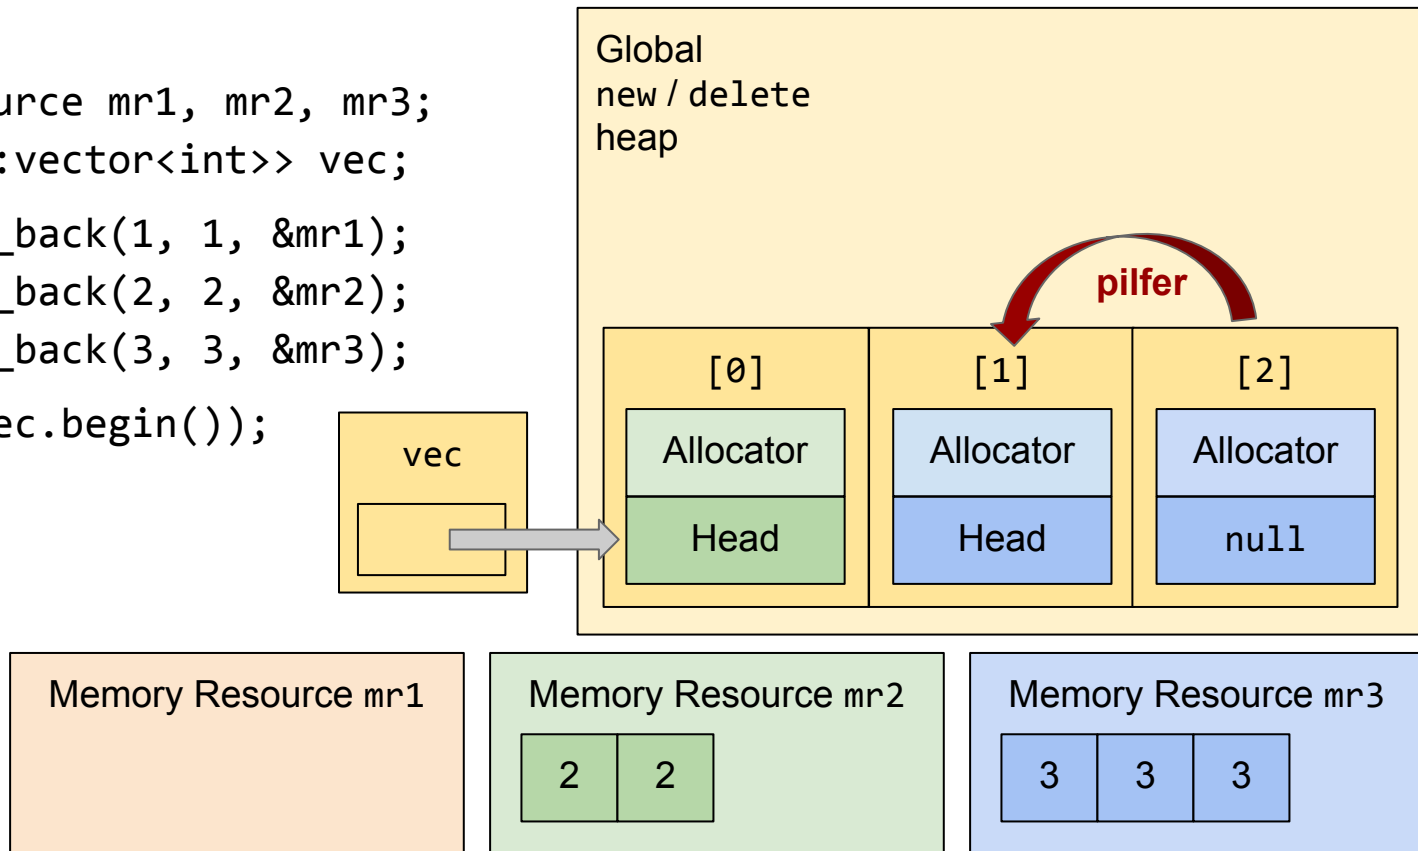
# If we implement erase via move+destroy

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
vec.erase(vec.begin());
```



# If we implement erase via move+destroy

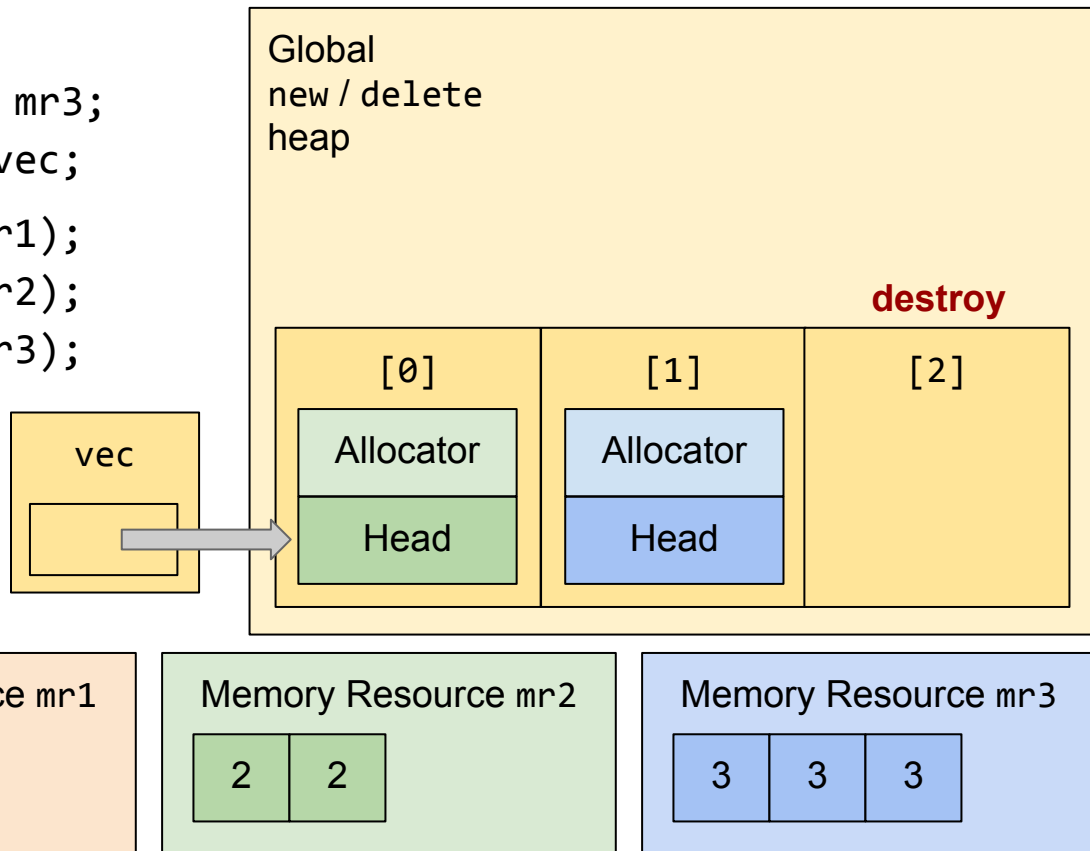
```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
vec.erase(vec.begin());
```



# If we implement erase via move+destroy

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
  
vec.erase(vec.begin());
```

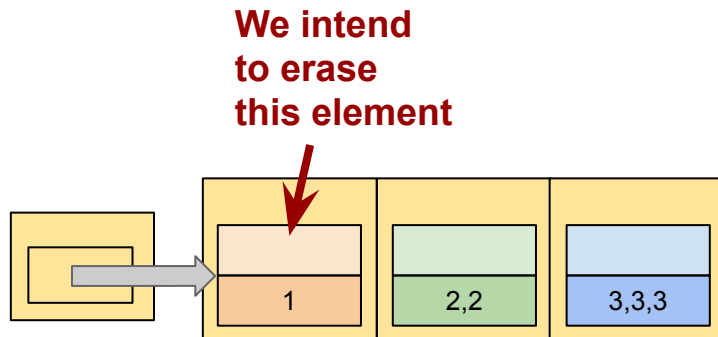
**Success, and very efficiently.  
But this might not be what you  
expected.**



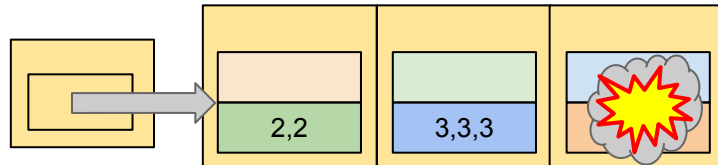
# Compare the outcomes

```
simple_resource mr1, mr2, mr3;  
vector<pmr::vector<int>> vec;  
  
vec.emplace_back(1, 1, &mr1);  
vec.emplace_back(2, 2, &mr2);  
vec.emplace_back(3, 3, &mr3);  
  
vec.erase(vec.begin());
```

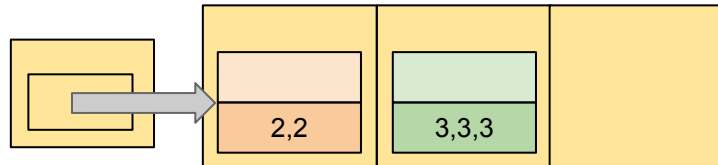
**Original:**



**Via swap:**



**Via move-assignment:**



**Via relocate:**

